

PIOTR ZWIERNIK

METHODS FOR MULTI-VARIATE DATA (STA437)

Introduction

o.o.1 About this course

These are rudimentary notes for STA437/2005 that contain more technical part of the lecture that will be presented in class on blackboard. Some multivariate methods will be presented on slides and so not treated in detail here.

Traditionally multivariate statistics is taught with a lot of focus on the multivariate normal distribution and the underlying statistical inference. I will be more brief on this part trying to get some basic overview of more modern methods for analyzing multivariate data.

My idea for this course is to mix two styles: (i) give a deeper theoretical insight into the underlying linear algebra and canonical multivariate distributions/models, (ii) give some taste for modern state-of-the art methods. For this reason, the lecture will be mostly given on blackboard but occasionally I will present some slides to give a more high-level overview of more recent material. The tutorials will also mix theory and practical considerations.

o.o.2 Prerequisites

To understand better what you are expected to know, note that one of the prerequisites for STA437 is STA302. A prerequisite for STA302 is MAT223 or MAT224 or MAT240. Students should have knowledge of linear algebra from these courses and its applications to linear models from STA302.

Multivariate statistics very heavily relies on matrix algebra and, more generally, on linear algebra. I will try to briefly recall some key facts when we need them but there will be no time to discuss these results in detail so they are assumed. Please make sure you have all the right background. There are no shortcuts here. The more you understand linear algebra, the more successful you will be in understanding this course and many other data analysis methods you will encounter in the future. The material gathered in Appendix A of ¹ should be more than enough. For deeper understanding check, for example, Chapter I in ² or Jason Siefken's course book for MAT223.

¹ Kanti V Mardia, John T Kent, and Charles C Taylor. *Multivariate analysis*, volume 88. John Wiley & Sons, 2024

² Gilbert Strang. *Linear Algebra and Learning from Data*. Wellesley-Cambridge Press, Wellesley, MA, 2019. ISBN 9780692196380

o.0.3 Complementary courses

There are several courses that are complementary to STA437/STA2005. For example, it is natural to discuss some classification algorithms in the context of multivariate statistics. However, this is covered in detail in STA314 so we will mostly omit this here. Some important examples are k-means, hierarchical clustering, and linear discriminant analysis³. Another related course, STA414/2104, focuses on unsupervised learning and we recommend students to take this course to get a full picture where the methods we learn here can be applied.

³ We will talk about Gaussian mixtures, which give access to classification methods that, in a sense, generalize both k-means and LDA.

o.0.4 Computational statistics

Although the lecture focuses on theory and methods, the computational aspect of this course is critical for building and verifying the correct intuition. Occasionally, we provide R code for your convenience. If you are not fluent in R, I recommend using AI to generate complementary code or to translate it into Python or Julia.

To report typos/mistakes, write me: piotr.zwiernik@utoronto.ca

I would like to thank Luis Sierra Muntané for his help with improving the notes.

Last update: March 4, 2025

Contents

1	<i>Preliminaries</i>	9
	1.1 <i>Basic Linear Algebra</i>	9
	1.2 <i>Random vectors and matrices</i>	13
	1.3 <i>Sample Quantities</i>	16
	1.4 <i>Exercises</i>	20
2	<i>Multivariate Normal Distribution</i>	25
	2.1 <i>Definition and basic properties</i>	25
	2.2 <i>Linear Transformation and Marginal Distribution</i>	28
	2.3 <i>Conditional Distribution and Conditional independence</i>	29
	2.4 <i>Estimation of Parameters</i>	32
	2.5 <i>Hotelling's T² Distribution*</i>	34
	2.6 <i>Gaussian Processes in Multivariate Statistics</i>	37
	2.7 <i>Exercises</i>	42
3	<i>Non-normal distributions</i>	45
	3.1 <i>Elliptical distributions</i>	45
	3.2 <i>Copula models</i>	50
	3.3 <i>Gaussian mixture models</i>	57
	3.4 <i>Higher order moments and cumulants*</i>	63
	3.5 <i>Exercises</i>	64

4	<i>Principal Component Analysis</i>	71
4.1	<i>Principal components</i>	72
4.2	<i>Principal Components for Observed Data</i>	73
4.3	<i>Simple Graphics for PCA</i>	73
4.4	<i>PCA and Affine Approximating Subspaces</i>	74
4.5	<i>Principal Component Regression (PCR) and Probabilistic PCA (PPCA)</i>	76
4.6	<i>PCA and Matrix Completion*</i>	78
4.7	<i>Appendix: Covariance matrix estimation in high dimensions</i>	79
4.8	<i>Exercises</i>	83
5	<i>Some Other Dimension Reduction Methods</i>	87
5.1	<i>Multidimensional Scaling (MDS)</i>	87
5.2	<i>Laplacian Eigenmaps (Spectral Embedding)</i>	89
5.3	<i>Uniform Manifold Approximation and Projection (UMAP)</i>	94
5.4	<i>Autoencoders</i>	97
5.5	<i>Exercises</i>	100
6	<i>Canonical Correlation Analysis</i>	103
6.1	<i>Population CCA</i>	103
6.2	<i>Sample CCA</i>	105
6.3	<i>Practical Considerations and Applications</i>	105
6.4	<i>Interpretation and Visualization of CCA Results</i>	106
6.5	<i>Example: Real Data Analysis in R</i>	107
7	<i>Factor analysis and Independent Component Analysis</i>	109
7.1	<i>Motivating examples</i>	109
7.2	<i>Definitions</i>	111
7.3	<i>Model Implications and identifiability</i>	112
7.4	<i>Fitting FA and ICA</i>	115
7.5	<i>Choosing the Number of Latent Components</i>	118
7.6	<i>Illustrative application</i>	119
7.7	<i>Exercises</i>	120

8	<i>Graphical Models</i>	123
8.1	<i>Introduction</i>	123
8.2	<i>Appendix: Conditional independence</i>	124
8.3	<i>Graphs and Conditional Independence</i>	126
8.4	<i>Gaussian Graphical Models</i>	128
8.5	<i>Log-linear Graphical Models</i>	137
8.6	<i>Exercises</i>	145

1

Preliminaries

We start by recalling basic notation and facts. This part of the lecture notes is quite fast-paced because we refresh material that should be known from earlier courses. Please fill out all the details.

1.1 Basic Linear Algebra

We use standard notations such as \mathbb{R}^m to denote the set of all m -dimensional real vectors $\mathbf{x} = (x_1, \dots, x_m)$ and $\mathbb{R}^{n \times m}$ to denote the set of all $n \times m$ real matrices $A = (A_{ij})$. All vectors in \mathbb{R}^m are treated as column vectors, that is, we take $\mathbb{R}^m = \mathbb{R}^{m \times 1}$.

Recall that a matrix $A \in \mathbb{R}^{n \times m}$ represents a linear function from \mathbb{R}^m to \mathbb{R}^n . For every $\mathbf{x} \in \mathbb{R}^m$, the vector $A\mathbf{x}$ lies in \mathbb{R}^n . Checking that the function $f(\mathbf{x}) = A\mathbf{x}$ is linear should be straightforward¹.

¹ Make sure you see why.

1.1.1 Matrix-Vector Multiplication $A\mathbf{x}$

Let $\mathbf{x} \in \mathbb{R}^m$ be a vector, with $\mathbf{x} = (x_1, \dots, x_m)$, and let $A \in \mathbb{R}^{n \times m}$ be a matrix. Recall that the matrix-vector product $A\mathbf{x}$ is defined by the rule:

$$(A\mathbf{x})_i = \sum_{j=1}^m A_{ij}x_j,$$

which states that the i -th entry of the vector $A\mathbf{x} \in \mathbb{R}^n$ is the inner product of the i -th row of A with the vector \mathbf{x} .

An important alternative interpretation of $A\mathbf{x}$ is that it is a linear combination of the columns of A with coefficients given by the entries of \mathbf{x} :

$$A\mathbf{x} = x_1\mathbf{a}_1 + \cdots + x_m\mathbf{a}_m, \quad (1.1)$$

where $\mathbf{a}_j \in \mathbb{R}^n$ denotes the j -th column of A .

← Exercise 1.4.5

Example 1.1.1. In the least squares problem from linear regression, we have a data matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$ with rows $\mathbf{x}_1, \dots, \mathbf{x}_n$ and a vector of observations $\mathbf{y} = (y_1, \dots, y_n) \in \mathbb{R}^n$. The least squares estimator $\hat{\beta} \in \mathbb{R}^p$ minimizes

Recall that $\|\mathbf{x}\| = \sqrt{\mathbf{x}^\top \mathbf{x}}$ is the Euclidean distance.

$$\|\mathbf{y} - \mathbf{X}\beta\|^2 = \sum_{i=1}^n (y_i - \mathbf{x}_i^\top \beta)^2,$$

where $\mathbf{X}\beta$ represents the prediction. The interpretation of $\mathbf{X}\beta$ as a linear combination of the columns of \mathbf{X} helps us understand that the least squares problem finds the point in the column space of \mathbf{X} that is closest to \mathbf{y} .

Example 1.1.2. Let $\mathbf{e}_1, \dots, \mathbf{e}_m \in \mathbb{R}^m$ be the canonical unit basis. We have $A\mathbf{e}_i = \mathbf{a}_i$ for $i = 1, \dots, m$ and so A maps the canonical unit basis of \mathbb{R}^m to $\mathbf{a}_1, \dots, \mathbf{a}_m$. If $m = n$ and the columns of A are linearly independent, the vectors $\mathbf{a}_1, \dots, \mathbf{a}_m$ also form a basis of \mathbb{R}^m and so A represents a change of basis.

1.1.2 Rank-One Matrix $\mathbf{x}\mathbf{y}^\top$

A **rank-one matrix** is formed by the outer product of two vectors. Specifically, if $\mathbf{x} \in \mathbb{R}^n$ and $\mathbf{y} \in \mathbb{R}^m$, the outer product $\mathbf{x}\mathbf{y}^\top \in \mathbb{R}^{n \times m}$ is defined as:

$$A = \mathbf{x}\mathbf{y}^\top \quad \text{with} \quad A_{ij} = x_i y_j \quad \text{for all } i = 1, \dots, n, j = 1, \dots, m.$$

The resulting matrix has rank² at most one because all rows (or columns) of the matrix are scalar multiples of each other.

² Recall that the rank of A is the number of linearly independent rows (or equiv. columns).

1.1.3 Matrix Multiplication AB

Matrix multiplication extends the concept of matrix-vector multiplication to two matrices. Let $A \in \mathbb{R}^{n \times m}$ and $B \in \mathbb{R}^{m \times p}$. The product $AB \in \mathbb{R}^{n \times p}$ is defined by the rule:

$$(AB)_{ik} = \sum_{j=1}^m A_{ij}B_{jk},$$

which means that the (i, k) -th entry of AB is the inner product of the i -th row of A with the k -th column of B .

Alternatively, we can view matrix multiplication as a sum of rank-one matrices. Let $\mathbf{a}_1, \dots, \mathbf{a}_m$ be the columns of A and $\mathbf{b}_1, \dots, \mathbf{b}_m$ be the rows of B . Then:

$$AB = \sum_{j=1}^m \mathbf{a}_j \mathbf{b}_j^\top, \tag{1.2}$$

which expresses AB as the sum of rank-one matrices formed from the columns of A and the rows of B .

← Exercise 1.4.8

Example 1.1.3. Suppose $\mathbf{X} \in \mathbb{R}^{n \times p}$ is a data matrix with rows $\mathbf{x}_1, \dots, \mathbf{x}_n$. Then the matrix $\frac{1}{n}\mathbf{X}^\top \mathbf{X}$ can be written as:

$$\frac{1}{n}\mathbf{X}^\top \mathbf{X} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top.$$

This matrix will play an important role in this course. Note that if the rows of \mathbf{X} are independent and identically distributed (i.i.d.) samples from some distribution with mean zero and covariance matrix $\Sigma = \mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top]$, then by the law of large numbers, we have:

$$\frac{1}{n} \mathbf{X}^\top \mathbf{X} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top \xrightarrow{p} \Sigma \quad \text{as } n \rightarrow \infty.$$

1.1.4 Singular Value Decomposition (SVD)

The **singular value decomposition (SVD)** is one of the most important linear algebra tools in this course. It decomposes any matrix into a product of three matrices: two orthogonal matrices and a diagonal matrix.

Denote by $O(m)$ the set of orthogonal $m \times m$ matrices:

$$O(m) = \{U \in \mathbb{R}^{m \times m} : UU^\top = I_m\}.$$

An orthogonal matrix has the property that its rows (and columns) form an orthonormal basis for \mathbb{R}^m . Note that if $U \in O(m)$ then $\det(U) = \pm 1$, which follows from multiplicativity of the determinant and the fact that $\det(I_m) = 1$. An orthogonal matrix U is called a rotation matrix if $\det(U) = 1$ but this terminology is often used for all orthogonal matrices.

Theorem 1.1.4 (Singular Value Decomposition). *Let $A \in \mathbb{R}^{n \times m}$ with $n \geq m$.³ Then there exist orthogonal matrices $U \in O(n)$ and $V \in O(m)$ such that:*

$$A = UDV^\top,$$

where $D \in \mathbb{R}^{n \times m}$ is a “diagonal” matrix with non-negative entries on the diagonal:

$$D = \begin{bmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_m \\ \mathbf{0}_{n-m} & \cdots & \mathbf{0}_{n-m} \end{bmatrix}, \quad \sigma_1 \geq \sigma_2 \geq \cdots \geq \sigma_m \geq 0.$$

The values $\sigma_1, \dots, \sigma_m$ are called the **singular values** of A . The singular values are always uniquely defined.

Denote the columns of U by u_1, \dots, u_n and the columns of V by v_1, \dots, v_m . Note that, by any $i = 1, \dots, m$,

$$Av_i = UDV^\top v_i = \sigma_i Ue_i = \sigma_i u_i$$

and so each v_i is transformed into a scaled version of u_i , $i = 1, \dots, m$. We refer to v_i 's as right singular vectors and to u_i 's as the left singular vectors.

Here \xrightarrow{p} denotes convergence in distribution. You know the definition for scalar random variables: $X_n \xrightarrow{p} X$ if, for all $\epsilon > 0$, $\mathbb{P}(|X_n - X| > \epsilon) \rightarrow 0$. In other words, $|X_n - X| \xrightarrow{p} 0$. Now, given a sequence of random vectors (X_n) we say it converges in probability if $\|X_n - X\| \xrightarrow{p} 0$, where $\|\cdot\|$ is the Euclidean distance. Similarly, for a sequence of random matrices (S_n) , we have $S_n \xrightarrow{p} S$ if $\|S_n - S\|_F \xrightarrow{p} 0$, where $\|\cdot\|_F$ is the Frobenius norm.

Defining equation $UU^\top = I_m$ implies that $U^{-1} = U^\top$ and so $U^\top U = I_m$ too.

³ The case $n \leq m$ can be treated analogously.

$\mathbf{0}_n \in \mathbb{R}^n$ is the zero vector and $\mathbf{1}_n \in \mathbb{R}^n$ is the vector of ones.

The singular value decomposition expresses A as a sum of rank-one matrices:

$$A = \sigma_1 \mathbf{u}_1 \mathbf{v}_1^\top + \cdots + \sigma_m \mathbf{u}_m \mathbf{v}_m^\top.$$

Note that $\mathbf{u}_{m+1}, \dots, \mathbf{u}_n$ do not feature here so they are pretty much irrelevant.

1.1.5 Spectrum of a symmetric matrix

Let $A \in \mathbb{R}^{m \times m}$ be a square matrix. A non-zero vector $\mathbf{v} \in \mathbb{R}^m$ is called an **eigenvector** of A if there exists a scalar $\lambda \in \mathbb{R}$ such that:

$$A\mathbf{v} = \lambda\mathbf{v}. \quad (1.3)$$

Think about a geometric meaning of this equation.

In this case, λ is called an **eigenvalue** of A corresponding to the eigenvector \mathbf{v} .

The matrix A has at most m eigenvalues (counting multiplicities). The eigenvalues are the real roots of the characteristic polynomial $\det(A - \lambda I_m)$.⁴

Denote by S^m the set of all $m \times m$ symmetric matrices.

Lemma 1.1.5. *If $A \in S^m$, then all its eigenvalues are always real, and the eigenvectors corresponding to distinct eigenvalues are orthogonal.*

Proof. Since $A \in S^m$, it is symmetric, so $A = A^\top$. Suppose $A\mathbf{v} = \lambda\mathbf{v}$, where $\lambda = a + ib$ with $a, b \in \mathbb{R}$ (i.e., assume for the sake of contradiction that λ could be complex). Recall that the standard inner product in \mathbb{C}^m is $\langle \mathbf{x}, \mathbf{y} \rangle = \sum_{i=1}^m x_i \bar{y}_i$, where \bar{y} is the complex conjugate of y (if $y = a + bi$ then $\bar{y} = a - bi$). If $\mathbf{x}, \mathbf{y} \in \mathbb{R}^m$ then $\langle \mathbf{x}, \mathbf{y} \rangle = \mathbf{x}^\top \mathbf{y}$ and so

$$\mathbf{v}^\top A\mathbf{v} = \langle \mathbf{v}, A\mathbf{v} \rangle = \langle \mathbf{v}, \lambda\mathbf{v} \rangle = \bar{\lambda}\|\mathbf{v}\|^2.$$

But, by the symmetry of A ,

$$\mathbf{v}^\top A\mathbf{v} = (A\mathbf{v})^\top \mathbf{v} = \langle A\mathbf{v}, \mathbf{v} \rangle = \langle \lambda\mathbf{v}, \mathbf{v} \rangle = \lambda\|\mathbf{v}\|^2.$$

This implies

$$\lambda\|\mathbf{v}\|^2 = \bar{\lambda}\|\mathbf{v}\|^2.$$

Since $\|\mathbf{v}\|^2 \neq 0$, it follows that $\lambda = \bar{\lambda}$, which implies that $\lambda \in \mathbb{R}$. Therefore, all eigenvalues of A are real.

Now, for the orthogonality of eigenvectors corresponding to distinct eigenvalues, let \mathbf{v}_1 and \mathbf{v}_2 be eigenvectors corresponding to distinct eigenvalues λ_1 and λ_2 of A . Then

$$A\mathbf{v}_1 = \lambda_1 \mathbf{v}_1 \quad \text{and} \quad A\mathbf{v}_2 = \lambda_2 \mathbf{v}_2.$$

Taking the inner product of the first equation with \mathbf{v}_2 and the second with \mathbf{v}_1 , we get

$$\mathbf{v}_2^\top A\mathbf{v}_1 = \lambda_1 \mathbf{v}_2^\top \mathbf{v}_1 \quad \text{and} \quad \mathbf{v}_1^\top A\mathbf{v}_2 = \lambda_2 \mathbf{v}_1^\top \mathbf{v}_2.$$

← Exercise 1.4.7

⁴ We can extend this definition to complex eigenvalues and eigenvectors. The fundamental theorem of algebra tells us that the characteristic polynomial always has m complex roots (counting with multiplicities). This is not so important for our course.

$$S^m = \{A \in \mathbb{R}^{m \times m} : A = A^\top\}$$

Since A is symmetric, $\mathbf{v}_2^\top A \mathbf{v}_1 = \mathbf{v}_1^\top A \mathbf{v}_2$. Therefore,

$$\lambda_1 \mathbf{v}_2^\top \mathbf{v}_1 = \lambda_2 \mathbf{v}_1^\top \mathbf{v}_2.$$

Since $\lambda_1 \neq \lambda_2$, it must be that $\mathbf{v}_1^\top \mathbf{v}_2 = 0$, so \mathbf{v}_1 and \mathbf{v}_2 are orthogonal.
This completes the proof. \square

This directly gives the following fundamental result.

Theorem 1.1.6 (Spectral Theorem for Symmetric Matrices). *If $A \in \mathbb{S}^m$, then there exists an orthogonal matrix $U \in O(m)$ and a diagonal matrix $\Lambda \in \mathbb{R}^{m \times m}$ such that:*

$$A = U \Lambda U^\top,$$

where the diagonal entries of Λ , denoted by $\lambda_1, \dots, \lambda_m$, are the eigenvalues of A . The columns of U , denoted by $\mathbf{u}_1, \dots, \mathbf{u}_m$, are the corresponding eigenvectors of A , i.e.,

$$A \mathbf{u}_i = \lambda_i \mathbf{u}_i \quad \text{for all } i = 1, \dots, m.$$

The spectral theorem provides a full diagonalization of a symmetric matrix, allowing us to express it as a sum of projections onto its eigenvectors:

$$A = \sum_{i=1}^m \lambda_i \mathbf{u}_i \mathbf{u}_i^\top.$$

Each term in this sum is a rank-one matrix corresponding to an eigenvalue and its associated eigenvector.

Question: In what sense is the matrix $\mathbf{u}_i \mathbf{u}_i^\top$ a projection matrix?

1.1.6 Positive definite matrices and quadratic forms

Every symmetric matrix $A \in \mathbb{S}^p$ defines a quadratic form

$$q(\mathbf{x}) = \mathbf{x}^\top A \mathbf{x} = \sum_{i,j=1}^p A_{ij} x_i x_j.$$

By definition, A is called **positive semi-definite (PSD)** if $q(\mathbf{x}) \geq 0$ for all \mathbf{x} . If the inequality is strict for all non-zero \mathbf{x} , then A is **positive definite (PD)**. The set of PD matrices is denoted by \mathbb{S}_+^m and the set of PSD matrices is denoted by $\bar{\mathbb{S}}_+^m$.

Notation $\bar{\mathbb{S}}_+^m$ suggests that the set of PSD matrices is the Euclidean closure of the set of PD matrices, which is indeed true.

← Exercise 1.4.15

1.2 Random vectors and matrices

This section covers the fundamental concepts of random vectors, random matrices, and their associated mean and covariance structures.

1.2.1 Mean and Covariance

Let $X = (X_1, \dots, X_m) \in \mathbb{R}^m$ be a random vector with a probability density function $f(\mathbf{x})$. The expectation (or mean) of a scalar-valued function $g(X)$ is defined as:

$$\mathbb{E}[g(X)] = \int_{\mathbb{R}^m} g(\mathbf{x})f(\mathbf{x}) d\mathbf{x}.$$

We assume that all necessary integrals converge, ensuring that the expectation is finite.

Now, consider a random vector $X = (X_1, \dots, X_m) \in \mathbb{R}^m$ or a random matrix $S = (S_{ij}) \in \mathbb{R}^{n \times m}$. When we refer to $\mathbb{E}[X]$ or $\mathbb{E}[S]$, we mean that the expectation is applied element-wise:

$$\mathbb{E}[X] = (\mathbb{E}[X_i])_{i=1}^m, \quad \mathbb{E}[S] = (\mathbb{E}[S_{ij}])_{i,j}.$$

We denote the mean vector of $X \in \mathbb{R}^m$ as:

$$\mu = \mathbb{E}[X] = (\mathbb{E}[X_1], \dots, \mathbb{E}[X_m]) = (\mu_1, \dots, \mu_m).$$

The covariance matrix of X , denoted by Σ , is defined as:

$$\Sigma = \text{var}(X) = \mathbb{E}[(X - \mu)(X - \mu)^\top].$$

The diagonal elements of Σ represent the variances of the individual components of X :

$$\Sigma_{jj} = \mathbb{E}[(X_j - \mu_j)^2] = \text{var}(X_j),$$

while the off-diagonal elements represent the covariances between the components:

$$\Sigma_{jk} = \mathbb{E}[(X_j - \mu_j)(X_k - \mu_k)] = \text{cov}(X_j, X_k).$$

Recall from Section 1.1.6 the definition of PSD matrices. The covariance matrix Σ is always positive semi-definite. Indeed, for any vector $\mathbf{u} \in \mathbb{R}^m$, we have:

$$\mathbf{u}^\top \text{var}(X) \mathbf{u} = \sum_{i,j} u_i u_j \mathbb{E}(X_i - \mu_i)(X_j - \mu_j) = \mathbb{E}[\sum_i u_i (X_i - \mu_i) \sum_j u_j (X_j - \mu_j)] = \mathbb{E}(\mathbf{u}^\top (X - \mu))^2,$$

which is always nonnegative. The covariance matrix Σ is positive definite unless there is a linear combination of the components of X that is degenerate.

Example 1.2.1. Consider a covariance matrix $\Sigma \in \mathbb{S}^m$. Since covariance matrices are symmetric and positive semi-definite, they admit a spectral decomposition:

$$\Sigma = U \Lambda U^\top,$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_m)$ contains the eigenvalues (which are non-negative). This decomposition is useful in multivariate statistics, for example, when performing **Principal Component Analysis (PCA)**.

For discrete distributions we can take f to be the probability mass function and replace integration with summation.

More generally, for two random vectors $X \in \mathbb{R}^p$ and $Y \in \mathbb{R}^q$, the covariance matrix between X and Y is given by:

$$\text{cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])^\top] \in \mathbb{R}^{p \times q}.$$

In particular, the variance of X is just the covariance of X with itself: $\text{var}(X) = \text{cov}(X, X)$.

1.2.2 Linearity of Expectation and Bilinearity of Covariance

The following properties of expectation are straightforward but important. They can be verified by checking each coordinate:

Lemma 1.2.2. *Let $X \in \mathbb{R}^m$ be a random vector, and $S \in \mathbb{R}^{p \times q}$ be a random matrix. Let $\mathbf{b} \in \mathbb{R}^n$ be a vector and let $A \in \mathbb{R}^{n \times p}$, $B \in \mathbb{R}^{q \times m}$ be matrices.*

Then

$$\mathbb{E}[AX + \mathbf{b}] = A\mathbb{E}[X] + \mathbf{b}.$$

Moreover,

$$\mathbb{E}[ASB] = A\mathbb{E}[S]B.$$

Proof. For the first equality, we check it on each coordinate. The i -th entry on the left is

$$\mathbb{E}\left(\sum_{j=1}^m A_{ij}X_j + b_i\right) = \sum_{j=1}^m A_{ij}\mathbb{E}(X_j) + b_i,$$

which is the same as the i -th entry of $A\mathbb{E}[X] + \mathbf{b}$ on the right. For the second equality, we see that the (i, j) -th entry of the matrix on the left is

$$\mathbb{E}\left[\sum_{k,l} A_{ik}S_{k,l}B_{l,j}\right] = \sum_{k,l} A_{ik}\mathbb{E}[S_{k,l}]B_{l,j}$$

with the same conclusion as before. \square

Lemma 1.2.3. *Consider random vectors $X \in \mathbb{R}^p$, $Y \in \mathbb{R}^q$. Then*

1. $\text{cov}(X, Y) = \text{cov}(Y, X)^\top$ (symmetry).
2. $\text{cov}(X + X', Y) = \text{cov}(X, Y) + \text{cov}(X', Y)$ (additivity).

Moreover, for any matrices $A \in \mathbb{R}^{n \times p}$, $B \in \mathbb{R}^{m \times q}$

$$\text{cov}(AX, BY) = A \text{cov}(X, Y) B^\top$$

and

$$\text{var}(AX) = A \text{var}(X) A^\top.$$

Exercise 1.2.4. *Prove this lemma. Hint: Use the definition of the covariance and Lemma 1.2.2.*

1.3 Sample Quantities

We now turn to the sample analogs of the mean and covariance matrix, which are important when working with observed data. Suppose we observe an i.i.d. sample $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^m$. We collect these observed vectors into a data matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$, where each row corresponds to one observation.

Two sample statistics are particularly important: the sample mean and the sample covariance matrix.

1.3.1 Sample Mean

The sample mean vector is defined as:

$$\bar{\mathbf{x}}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i = \frac{1}{n} \mathbf{X}^\top \mathbf{1}_n,$$

where $\mathbf{1}_n$ is the vector of all ones in \mathbb{R}^n . This vector summarizes the central location of the sample.

1.3.2 Sample Covariance Matrix

The sample covariance matrix is an estimator of the population covariance matrix. It measures the variability and linear relationships between the components of the random vector across the sample. It is defined as:

$$S_n = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}}_n)(\mathbf{x}_i - \bar{\mathbf{x}}_n)^\top. \quad (1.4)$$

We define the **centering matrix** as the matrix

$$H = I_n - \frac{1}{n} \mathbf{1}_n \mathbf{1}_n^\top.$$

The matrix H is symmetric, idempotent ($H^2 = H$), and projects any vector in \mathbb{R}^n onto the space orthogonal to the vector $\mathbf{1}_n$. This matrix is used to remove the effect of the sample mean from the data and it gives us a useful alternative formulation of S_n .

← Exercise 1.4.24

Lemma 1.3.1. *The sample covariance matrix can be written as:*

$$S_n = \frac{1}{n} \mathbf{X}^\top H \mathbf{X}.$$

Proof. The rows of the matrix

$$\mathbf{X} - \mathbf{1}_n \bar{\mathbf{x}}_n^\top = \mathbf{X} - \frac{1}{n} \mathbf{1}_n \mathbf{1}_n^\top \mathbf{X} = H \mathbf{X}$$

are precisely $\mathbf{x}_i - \bar{\mathbf{x}}_n$. Using Subsection 1.1.3, we see that S_n defined in (1.4) has a matrix representation $\frac{1}{n} \mathbf{X}^\top H^\top H \mathbf{X}$, which concludes the proof by the fact that H is symmetric and idempotent. □

Exercise 1.3.2. Show that the kernel of H is the span of $\mathbf{1}_n$. Conclude that $\text{rank}(H) = n - 1$.

In statistical software like R, the function `scale(X, center = TRUE, scale = FALSE)` can be used to center the data matrix X by subtracting the column means.

Recall that the rank of a matrix is the maximal number of linearly independent row/columns.

1.3.3 Properties of Sample Statistics

Now, suppose that the sample $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^m$ comes from a population with the same distribution as the random vector $X \in \mathbb{R}^m$. As before, we collect the sample as rows in the data matrix $\mathbf{X} \in \mathbb{R}^{n \times m}$. Suppose that $\mathbb{E}[X] = \mu$ and $\text{var}(X) = \Sigma$ are the population mean and covariance, respectively.

We now examine some properties of the sample statistics.

Lemma 1.3.3. The sample mean $\bar{\mathbf{x}}_n$ is an unbiased estimator of the population mean:

$$\mathbb{E}[\bar{\mathbf{x}}_n] = \mu.$$

Proof. By the linearity of expectation:

$$\mathbb{E}[\bar{\mathbf{x}}_n] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[\mathbf{x}_i] = \frac{1}{n} \sum_{i=1}^n \mu = \mu.$$

□

Next, we compute the variance of the sample mean.

Lemma 1.3.4. The variance of the sample mean is:

$$\text{var}(\bar{\mathbf{x}}_n) = \frac{1}{n} \Sigma.$$

Proof. Since the components of the sample are i.i.d., we can apply the bilinearity of covariance:

$$\begin{aligned} \text{var}\left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i\right) &= \text{cov}\left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i, \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i\right) = \frac{1}{n^2} \sum_{i,j} \text{cov}(\mathbf{x}_i, \mathbf{x}_j) \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{cov}(\mathbf{x}_i, \mathbf{x}_i) = \frac{1}{n^2} \sum_{i=1}^n \text{var}(\mathbf{x}_i) = \frac{1}{n^2} \cdot n\Sigma = \frac{1}{n} \Sigma \end{aligned}$$

Note that $\text{cov}(\mathbf{x}_i, \mathbf{x}_j) = 0$ by independence.

□

These two lemmas show that $\bar{\mathbf{x}}_n$ is a decent estimator of μ and its distribution quickly concentrates around μ . To show one version of this concentration phenomenon, recall the most basic probabilistic inequality.

← Exercise 1.4.25

Lemma 1.3.5 (Markov inequality). *If Z is a non-negative random variable, then*

$$\mathbb{P}(Z \geq t) \leq \frac{\mathbb{E}Z}{t}.$$

Proof. We have⁵

$$t\mathbb{P}(Z \geq t) = \mathbb{E}[t\mathbf{1}(Z \geq t)] \leq \mathbb{E}[Z\mathbf{1}(Z \geq t)] \leq \mathbb{E}Z,$$

where the last inequality follows because $Z \geq 0$. \square

Proposition 1.3.6. *The sample mean satisfies the following concentration inequality*

$$\mathbb{P}(\|\bar{\mathbf{x}}_n - \mu\| \geq t) \leq \frac{\text{tr}(\Sigma)}{nt^2}.$$

Proof. Consider $Y = \|\bar{\mathbf{x}}_n - \mu\|$ so that

$$Y^2 = \|\bar{\mathbf{x}}_n - \mu\|^2 = (\bar{\mathbf{x}}_n - \mu)^\top (\bar{\mathbf{x}}_n - \mu) = \text{tr}((\bar{\mathbf{x}}_n - \mu)(\bar{\mathbf{x}}_n - \mu)^\top).$$

We have

$$\mathbb{E}Y^2 = \text{tr}(\mathbb{E}((\bar{\mathbf{x}}_n - \mu)(\bar{\mathbf{x}}_n - \mu)^\top)) = \text{tr}(\text{var}(\bar{\mathbf{x}}_n)) = \frac{1}{n}\text{tr}(\Sigma).$$

Thus, using the Markov inequality, we get

$$\mathbb{P}(\|\bar{\mathbf{x}}_n - \mu\| \geq t) = \mathbb{P}(Y^2 \geq t^2) \leq \frac{\text{tr}(\Sigma)}{nt^2}.$$

⁵ We use the fact that $\mathbb{E}\mathbf{1}\{X \in A\} = \mathbb{P}(X \in A)$

We use the fact that $\mathbb{R} = \mathbb{R}^{1 \times 1}$ and so $(\bar{\mathbf{x}}_n - \mu)^\top (\bar{\mathbf{x}}_n - \mu) = \text{tr}((\bar{\mathbf{x}}_n - \mu)^\top (\bar{\mathbf{x}}_n - \mu))$. Moreover $\text{tr}(AB) = \text{tr}(BA)$ whenever AB and BA are both well defined.

This is only a very simple concentration bound which works under minimal assumptions. Much sharper bounds can be found if the distribution of the sample is more structured.

Finally, we consider the bias of the sample covariance matrix.

Lemma 1.3.7. *The sample covariance matrix S_n is a biased estimator of the population covariance matrix Σ . Its expectation is:*

$$\mathbb{E}[S_n] = \frac{n-1}{n}\Sigma.$$

Proof. Expanding S_n in terms of $\mathbf{x}_i - \mu$ and $\bar{\mathbf{x}}_n - \mu$, we get

$$S_n = \frac{1}{n} \sum_{i=1}^n ((\mathbf{x}_i - \mu) - (\bar{\mathbf{x}}_n - \mu)) ((\mathbf{x}_i - \mu) - (\bar{\mathbf{x}}_n - \mu))^\top.$$

Expanding this expression, we find:

$$S_n = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \mu)(\mathbf{x}_i - \mu)^\top - (\bar{\mathbf{x}}_n - \mu)(\bar{\mathbf{x}}_n - \mu)^\top.$$

Taking expectations, we get

$$\mathbb{E}[S_n] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[(\mathbf{x}_i - \mu)(\mathbf{x}_i - \mu)^\top] - \mathbb{E}[(\bar{\mathbf{x}}_n - \mu)(\bar{\mathbf{x}}_n - \mu)^\top].$$

Since $\mathbb{E}[(\mathbf{x}_i - \mu)(\mathbf{x}_i - \mu)^\top] = \Sigma$ and $\mathbb{E}[(\bar{\mathbf{x}}_n - \mu)(\bar{\mathbf{x}}_n - \mu)^\top] = \frac{\Sigma}{n}$, we have

$$\mathbb{E}[S_n] = \Sigma - \frac{1}{n}\Sigma = \frac{n-1}{n}\Sigma.$$

Simplifying, we find

$$\mathbb{E}[S_n] = \frac{n}{n-1}\Sigma - \frac{\Sigma}{n-1} = \frac{n-1}{n}\Sigma.$$

□

This result shows that S_n is a biased estimator of Σ with $\mathbb{E}(S_n - \Sigma) = -\frac{1}{n}\Sigma$, which is asymptotically negligible.

1.3.4 Correlation and Sample Correlation

In many cases, it is useful to measure the linear relationship between the components of a random vector using the **correlation** rather than the covariance. The correlation matrix normalizes the covariance matrix so that all diagonal elements are equal to 1.

If Σ is a covariance matrix, we define the diagonal matrix $D_\Sigma = \text{diag}(\Sigma_{11}, \dots, \Sigma_{mm})$. The **correlation matrix** R is given by:

$$R = D_\Sigma^{-1/2} \Sigma D_\Sigma^{-1/2}.$$

The elements of R are:

$$R_{ij} = \begin{cases} 1 & \text{if } i = j, \\ \frac{\Sigma_{ij}}{\sqrt{\Sigma_{ii}\Sigma_{jj}}} & \text{if } i \neq j. \end{cases}$$

The correlation matrix measures the strength of the linear relationship between pairs of variables, with R_{ij} taking values in the range $[-1, 1]$. The value $R_{ij} = 1$ indicates perfect positive correlation, while $R_{ij} = -1$ indicates perfect negative correlation.

← Exercise 1.4.26

The **sample correlation matrix** is defined in the same way as the population correlation matrix, but using the sample covariance matrix S_n instead of Σ . That is, the sample correlation matrix \hat{R} is given by:

$$\hat{R} = D_{S_n}^{-1/2} S_n D_{S_n}^{-1/2},$$

where D_{S_n} is the diagonal matrix of the variances of the components in the sample.

The correlation matrix is often preferred over the covariance matrix when the variables have different units or scales, as it provides a dimensionless measure of linear dependence.

1.4 Exercises

In this section we provide some additional exercises that will help you to refresh the basic concepts and practice the new ones.

Exercise 1.4.1. Consider a matrix $A \in \mathbb{R}^{3 \times 3}$ and a vector $\mathbf{x} = (1, 2, 3) \in \mathbb{R}^3$. Let

$$A = \begin{pmatrix} 1 & 2 & 3 \\ 0 & 1 & 4 \\ 5 & 6 & 0 \end{pmatrix}.$$

Compute $A\mathbf{x}$ using both interpretations of matrix-vector multiplication: (i) by taking inner products of rows with \mathbf{x} , and (ii) as a linear combination of the columns of A .

Exercise 1.4.2. Let $\mathbf{x} = (1, 2, 3) \in \mathbb{R}^3$ and $\mathbf{y} = (4, 5, 6) \in \mathbb{R}^3$. Compute the matrix $\mathbf{x}\mathbf{y}^\top$. Argue why this matrix has rank one. Compute $\text{tr}(\mathbf{x}\mathbf{y}^\top)$ and $\mathbf{x}^\top \mathbf{y}$. Is it a coincidence that these two numbers are equal?

Exercise 1.4.3. Suppose $A = \mathbf{u}\mathbf{v}^\top$ is a rank-one matrix where $\mathbf{u} \in \mathbb{R}^n$ and $\mathbf{v} \in \mathbb{R}^m$. Show that for any vector $\mathbf{x} \in \mathbb{R}^m$, the vector $A\mathbf{x}$ is a scalar multiple of \mathbf{u} . What does this imply about the image of A ?

Exercise 1.4.4. Let $A \in \mathbb{R}^{3 \times 2}$ be the matrix:

$$A = \begin{pmatrix} 1 & 2 \\ 0 & 3 \\ 4 & 0 \end{pmatrix}.$$

Compute the singular value decomposition (SVD) of A by finding orthogonal matrices U and V , and a diagonal matrix D such that $A = UDV^\top$. (Also try to do it in R/Python.)

Exercise 1.4.5. Prove (1.1) by showing that the i -th entry of both expressions for $A\mathbf{x}$ is the same. Specifically, verify that the equality holds for each coordinate $i = 1, \dots, n$.

Exercise 1.4.6. Let $A \in \mathbb{R}^{n \times m}$ and $\mathbf{x} \in \mathbb{R}^m$. Suppose A represents a linear transformation, and \mathbf{x} is a unit vector ($\|\mathbf{x}\| = 1$). Explain the geometric meaning of the vector $A\mathbf{x}$.

a) What can you say about $A\mathbf{x}$ if A is orthogonal? (here $m = n$)

b) What if A is a diagonal matrix with positive entries?

Exercise 1.4.7. Show that eigenvalues defined in (1.3) can be equivalently characterized as the roots of the characteristic polynomial $\det(A - \lambda I_m)$.

Exercise 1.4.8. Prove (1.2) by expanding both sides of the equation and verifying that the (i, k) -th entries match for each i and k .

Exercise 1.4.9. Let $A \in \mathbb{S}^2$ be the matrix

$$A = \begin{pmatrix} 3 & 2 \\ 2 & 6 \end{pmatrix}.$$

Compute the eigenvalues and eigenvectors of A by solving the characteristic equation and verify the spectral theorem by expressing A as $A = U\Lambda U^\top$, where U is an orthogonal matrix of eigenvectors and Λ is the diagonal matrix of eigenvalues.

Exercise 1.4.10. Let $A \in \mathbb{R}^{n \times m}$. Express the singular values of AA^\top and $A^\top A$ in terms of the singular values of A .

Exercise 1.4.11. Let $A \in \mathbb{S}^m$ be a symmetric matrix with eigenvalues $\lambda_1, \dots, \lambda_m$.

a) Show that the largest eigenvalue of A can be characterized as $\lambda_{\max} = \max_{\mathbf{x}: \|\mathbf{x}\|=1} \mathbf{x}^\top A \mathbf{x}$.

b) Using this, prove that $\lambda_{\max} \geq \frac{1}{m} \text{tr}(A)$.

Exercise 1.4.12. Prove that the covariance matrix Σ is always positive semi-definite. Specifically, show that for any non-zero vector $\mathbf{u} \in \mathbb{R}^m$, $\mathbf{u}^\top \Sigma \mathbf{u} \geq 0$. Explain why this result implies that all eigenvalues of a covariance matrix are non-negative.

Exercise 1.4.13. Suppose you have the following data points in \mathbb{R}^2 : $(1, 2), (2, 4), (3, 6)$. Compute the sample mean vector and the sample covariance matrix for this dataset.

Exercise 1.4.14. Consider the matrix

$$A \in \mathbb{R}^{5 \times 5} : A = \begin{pmatrix} 1 & 2 & 3 & 4 & 5 \\ 2 & 3 & 4 & 5 & 6 \\ 3 & 4 & 5 & 6 & 7 \\ 4 & 5 & 6 & 7 & 8 \\ 5 & 6 & 7 & 8 & 9 \end{pmatrix}.$$

Use the singular value decomposition (SVD) to find a rank-2 approximation of A . Express this approximation as a sum of two rank-one matrices.

Exercise 1.4.15. Show that $A \in \mathbb{S}^m$ is positive definite if and only if all its eigenvalues are strictly positive. Hint: You can use Theorem 1.1.6.

Exercise 1.4.16. Let $A \in \mathbb{S}_+^m$ be a symmetric positive definite matrix with eigenvalues $\lambda_1, \dots, \lambda_m$.

a) Prove that adding a small perturbation ϵI (where I is the identity matrix and $\epsilon > 0$) to A preserves positive definiteness.

b) How do the eigenvalues of $A + \epsilon I$ relate to those of A ?

Exercise 1.4.17. An antisymmetric matrix is a square matrix whose transpose equals its negative (i.e., $A^\top = -A$). Let $A \in \mathbb{R}^{m \times m}$ be an antisymmetric matrix, show that

- $I_m - A$ and $I_m + A$ are nonsingular matrices
- $(I_m - A)^{-1}(I_m + A)$ is an orthogonal matrix.

Exercise 1.4.18. Suppose $A, B \in \mathbb{S}_+^m$. Show that (i) $A + B \in \mathbb{S}_+^m$, (ii) for every $a, b > 0$, $B(aA + bB)^{-1}A \in \mathbb{S}_+^m$.

Exercise 1.4.19. Prove that the eigenvalues of an idempotent matrix are either 0 or 1.

Exercise 1.4.20. Show that $\mathbf{X} \in \mathbb{R}^{n \times m}$ is such that $\mathbf{X}^\top \mathbf{X}$ is invertible, then the eigenvalues of $\mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top$ are either 0 or 1.

Exercise 1.4.21. Consider the power series for the real exponential function, given by

$$e^x = \sum_{k \geq 0} \frac{x^k}{k!}.$$

For a symmetric matrix $A \in \mathbb{S}^m$, we want to apply the same idea to compute “ e^A ” as

$$e^A := \sum_{k \geq 0} \frac{A^k}{k!} = I_m + A + \frac{A^2}{2!} + \dots$$

1. Use the spectral decomposition of A to find a simple expression for A^k in terms of the spectral decomposition of A .
2. Use the previous point to come up with a simple simple expression for e^A in terms of the spectral decomposition of A .
3. Compute e^A for

$$A = \begin{bmatrix} 2.2 & 0.4 \\ 0.4 & 2.8 \end{bmatrix}.$$

Exercise 1.4.22. Show that if the eigenvalues of a symmetric matrix $A \in \mathbb{S}^m$ are $\lambda_1, \dots, \lambda_m$, then for any $k \in \mathbb{N}$, the eigenvalues of A^k are $\lambda_1^k, \dots, \lambda_m^k$. If A is invertible (i.e., all $\lambda_i \neq 0$), the same result holds for all integer powers $k \in \mathbb{Z}$.

Exercise 1.4.23. Let $\mathbf{X} \in \mathbb{R}^{n \times m}$. Show that the matrix $\mathbf{X}^\top \mathbf{X}$ has rank less than or equal to n . Conclude that it cannot be PD if $n < m$.

Exercise 1.4.24. Show that the matrix $H = I_n - \frac{1}{n}\mathbf{1}_n\mathbf{1}_n^\top$ is symmetric, idempotent ($H^2 = H$), and projects any vector in \mathbb{R}^n onto the space orthogonal to the vector $\mathbf{1}_n$.

Exercise 1.4.25. The mean squared error (MSE) of the sample mean is defined as $\text{MSE} = \mathbb{E}\|\bar{\mathbf{x}}_n - \mu\|^2$. Show that if we allow the underlying dimension m grow together with n then the MSE may not be negligible.

Exercise 1.4.26. Prove that the correlation matrix R satisfies $R_{ij} \in [-1, 1]$ for all i, j . Hint: One idea is to use the fact that R is positive semidefinite.

Exercise 1.4.27. Show that if $Y = aX + b$ then $\text{corr}(X, Y) = \pm 1$.

Exercise 1.4.28. Consider the random vector $X = (X_1, X_2, X_3, X_4)$ with the following covariance matrix

$$\Sigma = \begin{bmatrix} 3 & 1 & 1 & 1 \\ 1 & 3 & 1 & 1 \\ 1 & 1 & 3 & 1 \\ 1 & 1 & 1 & 3 \end{bmatrix}.$$

Let the random vector $Y = (Y_1, Y_2, Y_3)$ be

$$Y_1 = X_1 - X_2, \quad Y_2 = X_1 + X_2 - 2X_3, \quad Y_3 = X_1 + X_2 + X_3 - 3X_4.$$

(a) Compute $\text{var}(Y)$.

(b) Which components of Y are uncorrelated?

Exercise 1.4.29. Let $X \in \mathbb{R}^m$ be a random vector with mean vector μ and covariance matrix Σ , also let $A \in \mathbb{R}^{m \times m}$ be a fixed matrix. Prove that

$$\mathbb{E}(X^\top AX) = \text{tr}(A\Sigma) + \mu^\top A\mu.$$

2

Multivariate Normal Distribution

Multivariate normal distribution is the most important distribution in multivariate statistics. Here we give a relatively streamlined treatment. If you are interested in more details, see Chapter 4 in ¹.

2.1 Definition and basic properties

The multivariate normal distribution is fundamental in multivariate statistics due to several key reasons:

1. It is highly tractable: closed under linear transformations, conditioning, and marginalization.
2. It provides a good approximation for sample distributions via the Central Limit Theorem.
3. It frequently arises naturally in the modelling of physical phenomena. See, for example, the work of Gauß in the context of measurement errors, the work of Einstein on Brownian motion, Galton board, and many other examples.

¹ Kanti V Mardia, John T Kent, and Charles C Taylor. *Multivariate analysis*, volume 88. John Wiley & Sons, 2024

As a result, the multivariate Gaussian distribution appears in many generative models, such as (Bayesian) linear regression, Gaussian Mixture Models, Probabilistic Principal Component Analysis (PPCA), Factor Analysis, Linear Dynamical Systems, Gaussian Processes, Gaussian Conditional Random Fields, autoencoders, etc.

2.1.1 Recall Univariate Gaussian

Recall that $X \sim N(\mu, \sigma^2)$ has the density

$$f_X(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}}, \quad x \in \mathbb{R}.$$

If $X \sim N(\mu, \sigma^2)$ then $\mathbb{E}X = \mu$ and $\text{var}(X) = \sigma^2$.

If $X_j \sim N(\mu_j, \sigma_j^2)$ for $j = 1, \dots, m$ are independent, then the vector $X = (X_1, \dots, X_m)$ has the density² for $x = (x_1, \dots, x_m)$

² Recall that if X has independent entries then its density is the product of densities of the individual entries.

$$f_X(\mathbf{x}) = \prod_{j=1}^m \frac{1}{\sqrt{2\pi}\sigma_j} e^{-\frac{1}{2}\frac{(x_j - \mu_j)^2}{\sigma_j^2}} = \frac{1}{(2\pi)^{m/2}} (\det(\Sigma))^{-1/2} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{x}-\boldsymbol{\mu})}, \quad \mathbf{x} \in \mathbb{R}^m,$$

where $\Sigma = \text{diag}(\sigma_1^2, \dots, \sigma_m^2)$. Here we used the fact that the determinant of a diagonal matrix is simply the products of its diagonal entries.

2.1.2 Definition of Multivariate Normal

The vector $\mathbf{X} = (X_1, \dots, X_m)$ is said to have a multivariate normal distribution with mean vector $\boldsymbol{\mu} \in \mathbb{R}^m$ and covariance matrix $\Sigma \in \mathbb{S}_+^m$ if it has the density

$$f_X(\mathbf{x}) = \frac{1}{(2\pi)^{m/2}} (\det(\Sigma))^{-1/2} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{x}-\boldsymbol{\mu})}. \quad (2.1)$$

We denote this distribution by $N_m(\boldsymbol{\mu}, \Sigma)$. The random vector $Z \sim N_m(\mathbf{0}, I_m)$ is called standard normal and its density if

$$\phi(\mathbf{z}) = \frac{1}{(2\pi)^{m/2}} e^{-\frac{1}{2}\|\mathbf{z}\|^2}.$$

Note that $\|\mathbf{z}\|^2 = \mathbf{z}^\top \mathbf{z}$.

It is extremely useful to think about any multivariate normal distribution as an affine transformation of the standard normal distribution.

Proposition 2.1.1. *If $Z \sim N_m(\mathbf{0}, I_m)$ then $\mathbf{X} = \boldsymbol{\mu} + \Sigma^{1/2}Z \sim N_m(\boldsymbol{\mu}, \Sigma)$.*

Proof. Recall the formula for the change of variables: If Z is a random vector with density p_z and $X = G(Z)$, with G bijective and differentiable, then the density of X is

$$p_x(\mathbf{x}) = p_z(G^{-1}(\mathbf{x})) |\det \nabla G^{-1}|.$$

We apply this in our situation, where $G(\mathbf{z}) = \boldsymbol{\mu} + \Sigma^{1/2}\mathbf{z}$ and $G^{-1}(\mathbf{x}) = \Sigma^{-1/2}(\mathbf{x} - \boldsymbol{\mu})$. In this case $\nabla G^{-1} = \Sigma^{-1/2}$ and so

$$p_x(\mathbf{x}) = \phi(\Sigma^{-1/2}(\mathbf{x} - \boldsymbol{\mu})) (\det \Sigma)^{-1/2},$$

which gives exactly the formula in (2.1). □

It is easy to check that if $\mathbf{X} \sim N_m(\boldsymbol{\mu}, \Sigma)$ then

$$\mathbb{E}\mathbf{X} = \boldsymbol{\mu} \quad \text{and} \quad \text{var}(\mathbf{X}) = \Sigma.$$

This could be derived by computing the corresponding integrals, e.g. $\mathbb{E}\mathbf{X} = \int_{\mathbb{R}^m} \mathbf{x} f_X(\mathbf{x}) d\mathbf{x}$. However, working with Gaussian distributions we should get accustomed to more algebraic proofs. For example, show first that if $Z \sim N_m(\mathbf{0}, I_m)$ then $\mathbb{E}Z = \mathbf{0}$ and $\text{var}(Z) = I_m$, which can be directly inferred from the univariate result and by independence of the components of Z . To conclude the general result, we can then use Lemma 1.2.2, Lemma 1.2.3, and Proposition 2.1.1.




Figure 2.1: Contour plot of the two-dimensional density of the Gaussian distribution with mean $(0,0)$ and correlation $\rho = 0.6$.

2.1.3 Geometric Interpretation

The contours of the density f_X in (2.1) are ellipsoids³ (c.f. Figure 2.1.2) defined by the set of points \mathbf{x} such that

$$(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) = c^2.$$

To see this explicitly, consider the spectral decomposition $\Sigma = U\Lambda U^\top$ so that $\Sigma^{-1} = U\Lambda^{-1}U^\top$. Define a new variable $\mathbf{y} = U^\top(\mathbf{x} - \boldsymbol{\mu})$ (an affine change of coordinates by translation and rotation), and then:

$$(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) = \mathbf{y}^\top \Lambda^{-1} \mathbf{y} = \sum_{j=1}^m \frac{y_j^2}{\lambda_j}.$$

Hence, the contours are ellipses aligned with the eigenvectors of Σ , with axis lengths proportional to the square roots of the eigenvalues of Σ .

2.1.4 Mahalanobis Distance

Fix a covariance matrix Σ . An important concept in multivariate normal theory is the **Mahalanobis distance**

$$\|\mathbf{x} - \mathbf{y}\|_\Sigma = \sqrt{(\mathbf{x} - \mathbf{y})^\top \Sigma^{-1} (\mathbf{x} - \mathbf{y})}.$$

This distance takes into account the covariance structure of the data.

Proposition 2.1.2. If $X \sim N_m(\boldsymbol{\mu}, \Sigma)$, then $(X - \boldsymbol{\mu})^\top \Sigma^{-1} (X - \boldsymbol{\mu}) \sim \chi_m^2$.

We leave the proof as an exercise.

³ Recall that an ellipsoid in \mathbb{R}^m is a quadratic surface which, for fixed $\lambda_1, \dots, \lambda_m, c$, is defined by the equation

$$\frac{y_1^2}{\lambda_1} + \dots + \frac{y_m^2}{\lambda_m} = c^2.$$

If $\Sigma = I_m$, we recover the Euclidean distance.

Recall that χ_m^2 has stochastic representation $Z_1^2 + \dots + Z_m^2$, where Z_i are i.i.d. $N(0, 1)$.

← Exercise 2.7.13

2.1.5 Characteristic Function

If X is a random variable then the characteristic function of X is a function defined by $\phi_X(s) = \mathbb{E}e^{isX}$, where i is the imaginary unit ($i^2 = -1$). For a random vector $X = (X_1, \dots, X_m)$ in \mathbb{R}^m we define the characteristic function as

$$\phi_X(\mathbf{t}) = \mathbb{E}e^{i\mathbf{t}^\top X}, \quad \text{where } \mathbf{t} = (t_1, \dots, t_m).$$

One of the important properties of characteristic functions is that they define the underlying probability distribution uniquely.

Recall that the characteristic function of $Z \sim N(0, 1)$ is:

$$\phi_Z(s) := \mathbb{E}[e^{isZ}] = e^{-s^2/2}.$$

From this, we can derive the characteristic function for $X = \boldsymbol{\mu} + \sigma Z \sim N(\boldsymbol{\mu}, \sigma^2)$:

$$\phi_X(s) = \mathbb{E}[e^{isX}] = e^{is\boldsymbol{\mu} - \sigma^2 s^2/2}. \quad (2.2)$$

For the multivariate case, let $Z = (Z_1, \dots, Z_m) \sim N_m(\mathbf{0}, I_m)$. Then:

$$\phi_Z(\mathbf{t}) = \mathbb{E}[e^{i\mathbf{t}^\top Z}] = \prod_{j=1}^m \mathbb{E}e^{it_j Z_j} = \prod_{j=1}^m e^{-\frac{1}{2}t_j^2} = e^{-\frac{1}{2}\|\mathbf{t}\|^2}.$$

Thus, for $X = \mu + \Sigma^{1/2}Z$, we get:

$$\phi_X(\mathbf{t}) = \mathbb{E}e^{i\mathbf{t}^\top (\mu + \Sigma^{1/2}Z)} = e^{i\mathbf{t}^\top \mu} \mathbb{E}e^{i(\Sigma^{1/2}\mathbf{t})^\top Z} = e^{i\mathbf{t}^\top \mu} \phi_Z(\Sigma^{1/2}\mathbf{t}) = e^{i\mathbf{t}^\top \mu - \frac{1}{2}\mathbf{t}^\top \Sigma \mathbf{t}}, \quad (2.3)$$

which gives the formula for the characteristic function of the multivariate normal distribution for general μ and Σ .

2.1.6 Alternative Characterization of the Gaussian distribution

The characteristic function provides an alternative characterization of the multivariate normal distribution. Specifically, we say that $X \sim N_m(\mu, \Sigma)$ if, for every $\mathbf{u} \in \mathbb{R}^m$, the linear combination $\mathbf{u}^\top X \sim N(\mathbf{u}^\top \mu, \mathbf{u}^\top \Sigma \mathbf{u})$.

To prove that this defines the same distribution, let $\phi_X(\mathbf{t})$ denote the characteristic function of the vector X . Since $\mathbf{u}^\top X \sim N(\mathbf{u}^\top \mu, \mathbf{u}^\top \Sigma \mathbf{u})$, by (2.2), for any $\mathbf{u} \in \mathbb{R}^m$, we have:

$$\phi_{\mathbf{u}^\top X}(s) = e^{is(\mathbf{u}^\top \mu) - \frac{1}{2}(\mathbf{u}^\top \Sigma \mathbf{u})s^2}.$$

On the other hand,

This is the crucial observation here.

$$\phi_{\mathbf{u}^\top X}(s) = \mathbb{E}[e^{is\mathbf{u}^\top X}] = \phi_X(s\mathbf{u}).$$

Comparing these two expressions, we conclude that $\phi_X(\mathbf{t}) = e^{i\mathbf{t}^\top \mu - \frac{1}{2}\mathbf{t}^\top \Sigma \mathbf{t}}$ as in (2.3), which proves the claim.

2.2 Linear Transformation and Marginal Distribution

By Proposition 2.1.1, if $Z \sim N_m(\mathbf{0}, I_m)$ then $\mu + \Sigma^{1/2}Z \sim N_m(\mu, \Sigma)$. In this section, we generalize this result. More generally, for any matrix $A \in \mathbb{R}^{p \times m}$ and vector $b \in \mathbb{R}^p$, if $X \sim N_m(\mu, \Sigma)$, then the linear transformation $AX + b$ is distributed as:

$$AX + b \sim N_p(A\mu + b, A\Sigma A^\top). \quad (2.4)$$

This again follows easily by using the characteristic function. Let $Y = AX + b$ then

$$\phi_Y(\mathbf{s}) = \mathbb{E}e^{i\mathbf{s}^\top (AX+b)} = e^{i\mathbf{s}^\top b} \phi_X(A^\top \mathbf{s}) = e^{i\mathbf{s}^\top (A\mu+b) - \frac{1}{2}\mathbf{s}^\top A\Sigma A^\top \mathbf{s}},$$

which is the characteristic function of the Gaussian distribution with mean $A\mu + b$ and covariance $A\Sigma A^\top$.

Consider now a subvector $X' = (X_1, \dots, X_p)$ of the vector $X = (X_1, \dots, X_p, \dots, X_m)$ with $p < m$. We have $X' = AX$ where $A = [I_p | \mathbf{0}_{p \times (m-p)}] \in \mathbb{R}^{p \times m}$, where $\mathbf{0}_{p \times (m-p)}$ denotes the matrix of zeros with p rows and $m - p$ columns. Using (2.4), we get that X' has Gaussian distribution with mean $\mu' = A\mu$ and covariance $\Sigma' = A\Sigma A^\top$. Note that $\mu' = (\mu_1, \dots, \mu_p)$ and Σ' is the submatrix of Σ corresponding to rows/columns in $\{1, \dots, p\}$.

More generally, for any subset $I \subseteq \{1, \dots, p\}$ denote by $X_I = (X_i)_{i \in I}$ the subvector of X with coordinates indexed by I . Similarly, by $\Sigma_{I,I}$ denote the submatrix of Σ with rows/columns in I . We get the following result.

Proposition 2.2.1. Suppose $X \sim N_m(\mu, \Sigma)$ and let $I \subseteq \{1, \dots, m\}$. Then the distribution of X_I is Gaussian with mean μ_I and covariance $\Sigma_{I,I}$.

← Exercise 2.7.14

We now discuss an important result regarding independence. Recall that if $X_i \perp\!\!\!\perp X_j$ then $\text{cov}(X_i, X_j) = 0$. The reverse conclusion is typically not true but it is true for multivariate Gaussian distributions. Indeed, suppose $X \sim N_m(\mu, \Sigma)$ and suppose that $\Sigma_{ij} = 0$ for some i, j . By Proposition 2.2.1, the marginal distribution of (X_i, X_j) is Gaussian with mean (μ_i, μ_j) and covariance

$$\Sigma_{I,I} = \begin{bmatrix} \Sigma_{ii} & 0 \\ 0 & \Sigma_{jj} \end{bmatrix}.$$

Since $\Sigma_{I,I}$ is diagonal, so is its inverse, and so the joint density of (X_i, X_j) factorizes, showing that $X_i \perp\!\!\!\perp X_j$. Indeed, this density is simply

$$f_{ij}(x_i, x_j) = \frac{1}{2\pi\sqrt{\Sigma_{ii}\Sigma_{jj}}} e^{-\frac{1}{2\Sigma_{ii}}(x_i-\mu_i)^2 - \frac{1}{2\Sigma_{jj}}(x_j-\mu_j)^2} = \left(\frac{1}{\sqrt{2\pi\Sigma_{ii}}} e^{-\frac{1}{2\Sigma_{ii}}(x_i-\mu_i)^2} \right) \left(\frac{1}{\sqrt{2\pi\Sigma_{jj}}} e^{-\frac{1}{2\Sigma_{jj}}(x_j-\mu_j)^2} \right).$$

More generally, we have the following result:

Lemma 2.2.2. Suppose $X \sim N_m(\mu, \Sigma)$. Then $AX \perp\!\!\!\perp BX$ if and only if $A\Sigma B^\top = 0$.

The proof is left as an exercise.

← Exercise 2.7.15

2.3 Conditional Distribution and Conditional independence

In the previous section we saw that Gaussian distributions are closed under taking margins. In this section we show that they are also closed under conditioning. We then discuss conditional independence in Gaussian distributions.

2.3.1 Conditional distribution

Suppose $X \sim N_m(\mu, \Sigma)$. Split $X = (X_A, X_B)$ into a p -dimensional subvector X_A and a $(m - p)$ -dimensional subvector X_B . This induces corresponding splits in the mean vector $\mu = (\mu_A, \mu_B)$ and the covariance matrix:

$$\Sigma = \begin{bmatrix} \Sigma_{A,A} & \Sigma_{A,B} \\ \Sigma_{B,A} & \Sigma_{B,B} \end{bmatrix}.$$

Theorem 2.3.1. *The marginal distribution of X_A is Gaussian with mean μ_A and covariance $\Sigma_{A,A}$. The conditional distribution of $X_B | X_A = x_A$ is Gaussian with mean:*

$$\mathbb{E}(X_B | X_A = x_A) = \mu_B + \Sigma_{B,A}\Sigma_{A,A}^{-1}(x_A - \mu_A),$$

and covariance:

$$\text{var}(X_B | X_A = x_A) = \Sigma_{B,B} - \Sigma_{B,A}\Sigma_{A,A}^{-1}\Sigma_{A,B}.$$

Proof. The first part was proved in Proposition 2.2.1. For the second part, define the new variable $X_{B \cdot A} := X_B - \Sigma_{B,A}\Sigma_{A,A}^{-1}X_A$, which is Gaussian as an affine function of a Gaussian vector. Computing the expectation and the covariance matrix, it can be shown that

$$X_{B \cdot A} \sim N_{m-p}(\mu_B - \Sigma_{B,A}\Sigma_{A,A}^{-1}\mu_A, \Sigma_{B,B} - \Sigma_{B,A}\Sigma_{A,A}^{-1}\Sigma_{A,B}).$$

Note that

$$\text{cov}(X_{B \cdot A}, X_A) = \text{cov}(X_B, X_A) - \Sigma_{B,A}\Sigma_{A,A}^{-1}\Sigma_{A,A} = \Sigma_{B,A} - \Sigma_{B,A} = 0$$

and so $X_{B \cdot A} \perp\!\!\!\perp X_A$. Thus, we can decompose X_B as a sum of two independent terms

$$X_B = X_{B \cdot A} + \Sigma_{B,A}\Sigma_{A,A}^{-1}X_A.$$

In particular, given $X_A = x_A$, the conditional distribution of X_B is Gaussian with mean and covariance as stated in the theorem. \square

2.3.2 Conditional Independence

Let $A = \{i, j\}$ and consider splitting the vector X into X_A and X_B , where $B = \{1, \dots, m\} \setminus A$. By Theorem 2.3.1, the conditional covariance matrix of X_A given X_B is:

$$\Sigma_{AA} - \Sigma_{AB}\Sigma_{BB}^{-1}\Sigma_{BA} \in \mathbb{S}_+^2.$$

The variables X_i and X_j are conditionally independent given the variables in X_B if and only if the conditional covariance matrix has a zero off-diagonal entry. Equivalently, we have $\Sigma_{ij} = \Sigma_{iB}\Sigma_{BB}^{-1}\Sigma_{Bj}$. This leads to several useful equivalent conditions:

For example, $X = (X_1, X_2, X_3, X_4)$, $A = \{1, 3\}$, $B = \{2, 4\}$ so that $X_A = (X_1, X_3)$ and $X_B = (X_2, X_4)$.

Proposition 2.3.2. Suppose $X \sim N(0, \Sigma)$. Fix $i, j \in \{1, \dots, m\}$ and let $B = \{1, \dots, m\} \setminus \{i, j\}$. The following conditions are equivalent:

- (i) $X_i \perp\!\!\!\perp X_j | X_B$,
- (ii) $\Sigma_{ij} = \Sigma_{iB}\Sigma_{BB}^{-1}\Sigma_{Bj}$,
- (iii) $K_{ij} = 0$, where $K = \Sigma^{-1}$ is the precision matrix.
- (iv) The linear regression coefficient vector $\beta_{i \cdot B \cup \{j\}} = \Sigma_{i, B \cup \{j\}}\Sigma_{B \cup \{j\}, B \cup \{j\}}^{-1}$ has a zero entry in the j -th index.

We are not going to prove this in class but the proof is relatively elementary with a bit of annoying bookkeeping. We provide it for your convenience.

Proof. The equivalence of (i) and (ii) is clear from the discussion above. For the other implications we rely on standard formulas for block matrices. To show equivalence with (iii) we use the formula for the inverse of a block matrix that we develop independently in Section 2.3.3. From the formula it follows that

$$K_{A,A} = (\Sigma_{A,A} - \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A})^{-1}. \quad (2.5)$$

We have (ii) if and only if $\Sigma_{A,A} - \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A}$ is diagonal. By (2.5), this is equivalent to $K_{A,A}$ being diagonal, which is equivalent to (iii). For the equivalence between (iii) and (iv) use Proposition 2.3.3 again to conclude that

$$K_{i, B \cup \{j\}} = -K_{ii}\Sigma_{i, B \cup \{j\}}(\Sigma_{B \cup \{j\}, B \cup \{j\}})^{-1}.$$

Since $K_{ii} > 0$, then $K_{ij} = 0$ if and only if the j -th entry of $\Sigma_{i, B \cup \{j\}}\Sigma_{B \cup \{j\}, B \cup \{j\}}^{-1}$ is zero. \square

The relationship between conditional independence and zeros in the precision matrix K is fundamental in **Gaussian graphical models**, which we will cover later in this course; see Chapter 8.

2.3.3 Partitioned Inverse*

Let $X \sim N(0, \Sigma)$ and consider the partition of X into two subvectors X_A and X_B . Set

$$K = \Sigma^{-1} = \begin{bmatrix} K_{A,A} & K_{A,B} \\ K_{B,A} & K_{B,B} \end{bmatrix},$$

where K is the precision matrix.

Proposition 2.3.3. *We have the following formulas for the blocks of the precision matrix in terms of the blocks of Σ :*

$$\begin{aligned} K_{A,A} &= (\Sigma_{A,A} - \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A})^{-1}, \\ K_{A,B} &= -K_{A,A}\Sigma_{A,B}\Sigma_{B,B}^{-1}, \\ K_{B,B} &= \Sigma_{B,B}^{-1} + \Sigma_{B,B}^{-1}\Sigma_{B,A}K_{A,A}\Sigma_{A,B}\Sigma_{B,B}^{-1} \end{aligned}$$

Proof. We verify that these formulas satisfy $\Sigma K = I_m$. For instance, consider multiplying the first block row of Σ by the first block column of K :

$$\Sigma_{A,A}K_{A,A} + \Sigma_{A,B}K_{B,A} = \Sigma_{A,A}K_{A,A} - \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A}K_{A,A} = (\Sigma_{A,A} - \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A})K_{A,A} = I.$$

Similarly, multiplying the first block row of Σ by the second block column of K gives:

$$\begin{aligned} \Sigma_{A,A}K_{A,B} + \Sigma_{A,B}K_{B,B} &= -\Sigma_{A,A}K_{A,A}\Sigma_{A,B}\Sigma_{B,B}^{-1} + \Sigma_{A,B}\Sigma_{B,B}^{-1} + \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A}K_{A,A}\Sigma_{A,B}\Sigma_{B,B}^{-1} \\ &= -(\Sigma_{A,A} - \Sigma_{A,B}\Sigma_{B,B}^{-1}\Sigma_{B,A})K_{A,A}\Sigma_{A,B}\Sigma_{B,B}^{-1} + \Sigma_{A,B}\Sigma_{B,B}^{-1} \\ &= 0. \end{aligned}$$

The remaining cases can be checked similarly. \square

This partitioned inverse result is useful in various scenarios, such as Gaussian graphical models and Bayesian networks. We already saw an application in the previous subsection.

2.4 Estimation of Parameters

Suppose we have data $\mathbf{X} \in \mathbb{R}^{n \times m}$, where the rows $\mathbf{x}_i \in \mathbb{R}^m$ are i.i.d. samples from $N_m(\mu, \Sigma)$. A natural way to estimate μ and Σ is through the Maximum Likelihood Estimator (MLE). In this section, we show that the MLEs $\hat{\mu}$ and $\hat{\Sigma}$ satisfy:

$$\hat{\mu} = \bar{\mathbf{x}}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i,$$

and

$$\hat{\Sigma} = S_n = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}}_n)(\mathbf{x}_i - \bar{\mathbf{x}}_n)^\top.$$

2.4.1 The Likelihood Function

Let $f(\mathbf{x})$ denote the density of $N_m(\mu, \Sigma)$ as given in (2.1). After taking logs, we get

$$\log f(\mathbf{x}) = -\frac{m}{2} \log(2\pi) + \frac{1}{2} \log \det(\Sigma^{-1}) - \frac{1}{2} (\mathbf{x} - \mu)^\top \Sigma^{-1} (\mathbf{x} - \mu).$$

The log-likelihood of the data $\mathbf{x}_1, \dots, \mathbf{x}_n$ is:

$$\ell(\mu, \Sigma) = \sum_{i=1}^n \log f(\mathbf{x}_i) = \text{const} + \frac{n}{2} \log \det(\Sigma^{-1}) - \frac{1}{2} \sum_{i=1}^n (\mathbf{x}_i - \mu)^\top \Sigma^{-1} (\mathbf{x}_i - \mu).$$

2.4.2 Optimize over μ

Using basic calculus, the derivative of the log-likelihood with respect to μ is:⁴

$$\nabla_{\mu} \ell(\mu, \Sigma) = \Sigma^{-1} \left(\sum_{i=1}^n \mathbf{x}_i - n\mu \right).$$

Setting this equal to zero gives the MLE for μ :

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i = \bar{\mathbf{x}}_n.$$

We get this result irrespective of what Σ is as the only thing we used is that Σ is invertible⁵.

2.4.3 Optimize over Σ

Knowing that $\hat{\mu} = \bar{\mathbf{x}}_n$, we now optimize $\ell(\bar{\mathbf{x}}_n, \Sigma)$ over Σ . For this, we first rewrite $\ell(\bar{\mathbf{x}}_n, \Sigma)$ in a suitable form. It is convenient to list the following useful observation.

Useful observation: We know that the trace satisfies $\text{tr}(AB) = \text{tr}(BA)$ whenever both multiplications are well defined. In particular, for any two $\mathbf{x}, \mathbf{y} \in \mathbb{R}^m$ we have

$$\mathbf{x}^\top \mathbf{y} = \text{tr}(\mathbf{x}^\top \mathbf{y}) = \text{tr}(\mathbf{y} \mathbf{x}^\top),$$

where $\mathbf{y} \mathbf{x}^\top \in \mathbb{R}^{m \times m}$. Using this, we get

$$\begin{aligned} \ell(\bar{\mathbf{x}}_n, \Sigma) &= \text{const} + \frac{n}{2} \log \det(\Sigma^{-1}) - \frac{1}{2} \sum_{i=1}^n \text{tr}((\mathbf{x}_i - \bar{\mathbf{x}}_n)(\mathbf{x}_i - \bar{\mathbf{x}}_n)^\top \Sigma^{-1}) \\ &= \text{const} + \frac{n}{2} \log \det(\Sigma^{-1}) - \frac{n}{2} \text{tr}\left(\frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}}_n)(\mathbf{x}_i - \bar{\mathbf{x}}_n)^\top \Sigma^{-1}\right) \\ &= \text{const} + \frac{n}{2} \log \det(\Sigma^{-1}) - \frac{n}{2} \text{tr}(S_n \Sigma^{-1}), \end{aligned}$$

where $S_n = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}}_n)(\mathbf{x}_i - \bar{\mathbf{x}}_n)^\top$ is the sample covariance matrix.

To optimize $\ell(\hat{\mu}, \Sigma)$ over Σ , we rewrite the log-likelihood in terms of $K = \Sigma^{-1}$:

$$\ell(\hat{\mu}, K) = \text{const} + \frac{n}{2} \log \det(K) - \frac{n}{2} \text{tr}(S_n K). \quad (2.6)$$

We are not going to show this formally here but this function is a strictly concave function of $K = \Sigma^{-1}$. In particular, every local maximum is a global maximum and if a global optimum exists, it is unique.

Useful observation: Another useful observation is that for any two $A, B \in \mathbb{S}^m$

$$\text{tr}(AB) = \sum_{i,j=1}^m A_{ij} B_{ij}.$$

⁴ Verify that $\nabla_{\mathbf{y}}(\mathbf{b}^\top \mathbf{y}) = \mathbf{b}$ and $\nabla_{\mathbf{y}} \mathbf{y}^\top K \mathbf{y} = 2K \mathbf{y}$.
⁵ Recall that $\ker(A) = \{\mathbf{x} : A\mathbf{x} = \mathbf{0}\}$ and A is invertible if and only if $\ker(A) = \{\mathbf{0}\}$.

This can be checked easily directly as $\text{tr}(\mathbf{x} \mathbf{y}^\top) = \sum_i (\mathbf{x} \mathbf{y}^\top)_{ii} = \sum_i x_i y_i = \mathbf{x}^\top \mathbf{y}$.

Note that the trace is a linear function: $\text{tr}(aA + bB) = a \text{tr}(A) + b \text{tr}(B)$.

In particular, this defines an inner product on \mathbb{S}^m , which is called the trace inner product⁶.

Using this, we easily check that $\nabla_K \text{tr}(S_n K) = S_n$. It can be also shown that $\nabla_K \log \det(K) = K^{-1} = \Sigma$; for a formal derivation check p. 641 [here](#). Thus, taking the derivative of $\ell(\bar{\mathbf{x}}_n, K)$ with respect to K and setting it equal to zero gives the MLE for Σ : $\hat{\Sigma} = S_n$ as long as S_n is positive definite. In the case when $n < m$, the matrix S_n cannot be positive definite⁷ and so the MLE does not exist.

2.4.4 Independence of $\bar{\mathbf{x}}_n$ and S_n

If $\mathbf{x}_1, \dots, \mathbf{x}_n$ is a sample from $N(\mu, \Sigma)$ then $\bar{\mathbf{x}}_n \sim N(\mu, \frac{1}{n}\Sigma)$ and the sample covariance matrix S_n is another random quantity of interest.

Proposition 2.4.1. *In the Gaussian sample, the sample mean $\bar{\mathbf{x}}_n = \frac{1}{n}\mathbf{X}^\top \mathbf{1}_n$ and the sample covariance $S_n = \frac{1}{n}\mathbf{X}^\top H\mathbf{X}$ are independent.*

Proof. Since $S_n = \frac{1}{n}\mathbf{X}^\top H\mathbf{X} = \frac{1}{n}(\mathbf{H}\mathbf{X})^\top (\mathbf{H}\mathbf{X})$, it is enough to show that $\bar{\mathbf{x}}_n$ is independent of $H\mathbf{X}$. Since both $\bar{\mathbf{x}}_n$ and $H\mathbf{X}$ have Gaussian distributions, independence can be checked by computing the covariance. Since the k -th row of $H\mathbf{X}$ is $\mathbf{x}_k - \bar{\mathbf{x}}_n$ we check

$$\text{cov}(\bar{\mathbf{x}}_n, \mathbf{x}_k - \bar{\mathbf{x}}_n) = \text{cov}(\bar{\mathbf{x}}_n, \mathbf{x}_k) - \text{var}(\bar{\mathbf{x}}_n) = \frac{1}{n} \text{var}(\mathbf{x}_k) - \frac{1}{n}\Sigma = 0.$$

□

2.5 Hotelling's T^2 Distribution*

Hotelling's T^2 distribution is a multivariate extension of the univariate Student's t-distribution. It is particularly useful in testing hypotheses about the mean in a single population or equality of means in two multivariate populations.

2.5.1 Testing equality of means in two univariate populations

Recall that if you have a pair of samples $x_i, i = 1, \dots, n$ from $N(\mu_x, \sigma^2)$ and $y_i, i = 1, \dots, m$ from $N(\mu_y, \sigma^2)$ then we can test $H_0 : \mu_x = \mu_y$ using the two-sample t-test. To see this, define

$$S = \frac{1}{n+m-2} \left[\sum_{i=1}^n (x_i - \bar{x})^2 + \sum_{j=1}^m (y_j - \bar{y})^2 \right] = \frac{n-1}{n+m-2} S_1 + \frac{m-1}{n+m-2} S_2$$

to be the pooled sample variance estimator of σ^2 . Define the test statistic

$$t = \frac{\bar{x} - \bar{y}}{\sqrt{S(\frac{1}{n} + \frac{1}{m})}} \tag{2.7}$$

To analyze its distribution, we first formulate the following lemma.

⁶ In the literature it is common to denote $\langle A, B \rangle := \text{tr}(AB)$. More generally, if $A, B \in \mathbb{R}^{n \times m}$ then $\text{tr}(AB^\top)$ also defines an inner product.

⁷ $S_n = \frac{1}{n}\mathbf{X}^\top H\mathbf{X}$ has rank at most n and $n < m$.

Recall $H = I_n - \frac{1}{n}\mathbf{1}_n\mathbf{1}_n^\top \in \mathbb{S}_+^m$ and $H^2 = H$

Note that we assume that the variance in both distributions is the same.

Lemma 2.5.1. If x_1, \dots, x_n is a random sample from $N(\mu, \sigma^2)$ then

$$\frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \bar{x})^2 \sim \chi_{n-1}^2.$$

Proof. Recall that $H = I_n - \frac{1}{n}\mathbf{1}_n\mathbf{1}_n^\top$ is a symmetric PSD matrix satisfying $H^2 = H$. Let $\mathbf{x} = (x_1, \dots, x_n)$ and recall that $H\mathbf{x}$ has entries $x_i - \bar{x}$. It follows that $\sum_{i=1}^n (x_i - \bar{x})^2 = \mathbf{x}^\top H\mathbf{x}$. Let $H = U\Lambda U^\top$. Since H is a projection matrix, $\Lambda = \Lambda^2$ and so the diagonal entries of Λ are either 0 or 1. Since $\text{rank}(H) = n - 1$ it follows that $\Lambda_{11} = 0$ and $\Lambda_{ii} = 1$ for $i = 2, \dots, n - 1$. Define $\mathbf{z} = U^\top \mathbf{x}$. Since $\mathbf{x} \sim N_n(\mu\mathbf{1}, \sigma^2 I_n)$, it follows that $\mathbf{z} \sim N_n(\mu U^\top \mathbf{1}, \sigma^2 I_n)$. Moreover, $U^\top \mathbf{1} = \sqrt{n}\mathbf{e}_1$, where $\mathbf{e}_1 = (1, 0, \dots, 0)$. Denote $\mathbf{z}_{n-1} = (z_2, \dots, z_n)$ obtained from \mathbf{z} by dropping the first coordinate. Then, in particular,

$$\frac{1}{\sigma} \mathbf{z}_{n-1} \sim N_{n-1}(\mathbf{0}, I_{n-1}).$$

We have

$$\frac{1}{\sigma^2} \mathbf{x}^\top H \mathbf{x} = \frac{1}{\sigma^2} \mathbf{z}^\top \Lambda \mathbf{z} = \frac{1}{\sigma^2} \mathbf{z}_{n-1}^\top \mathbf{z}_{n-1}$$

and we already noted that $\mathbf{z}_{n-1}^\top \mathbf{z}_{n-1} \sim \chi_{n-1}^2$. \square

Proposition 2.5.2. Under H_0 , the statistic t , defined in (2.7), has t-Student distribution with $n + m - 2$ degrees of freedom.

Proof. Recall that t-student distribution with r degrees of freedom has stochastic representation

$$t_r = \frac{U}{\sqrt{V/r}} \quad \text{with } U \sim N(0, 1), \quad V \sim \chi_r^2, \quad U \perp\!\!\!\perp V.$$

$U \perp\!\!\!\perp V$ denotes independence of U, V

Recall that $\bar{x} \sim N(\mu_x, \sigma^2/n)$, $\bar{y} \sim N(\mu_y, \sigma^2/m)$ and so, by independence, under H_0

$$\bar{x} - \bar{y} \sim N(0, \sigma^2(\frac{1}{n} + \frac{1}{m})) \quad \text{and so} \quad U := \frac{\bar{x} - \bar{y}}{\sigma \sqrt{\frac{1}{n} + \frac{1}{m}}} \sim N(0, 1).$$

Define $V = \frac{n+m-2}{\sigma^2} S$ and note that

$$t = \frac{\frac{\bar{x} - \bar{y}}{\sigma \sqrt{\frac{1}{n} + \frac{1}{m}}}}{\sqrt{\frac{1}{\sigma^2} S}} = \frac{U}{\sqrt{\frac{V}{m+n-2}}}.$$

We have

$$V = \frac{n+m-2}{\sigma^2} S = \frac{n-1}{\sigma^2} S_1 + \frac{m-1}{\sigma^2} S_2.$$

By Lemma 2.5.1, $\frac{n-1}{\sigma^2} S_1 \sim \chi_{n-1}^2$ and $\frac{m-1}{\sigma^2} S_2 \sim \chi_{m-1}^2$. Moreover, both these random variables are independent. It follows that $V \sim \chi_{m+n-2}^2$, which concludes the proof. \square

As many proofs in multivariate statistics, this one combines some basic results in linear algebra with properties of the Gaussian distribution.

2.5.2 Hotelling distribution

The Hotelling distribution describes the behavior of Hotelling's T^2 -statistic, a key tool in multivariate hypothesis testing. It generalizes the Student's t-distribution to the multivariate case and is commonly used to test hypotheses about the mean vector of a multivariate normal population.

Although we do not derive the Hotelling's distributions from the first principles here, we link it to another standard distribution called the **F-distribution**. Recall that if $X_1 \sim \chi_{d_1}^2$ and $X_2 \sim \chi_{d_2}^2$ are two independent chi-squared random variables, then the random variable: $F_{d_1, d_2} = \frac{(X_1/d_1)}{(X_2/d_2)}$ follows an F-distribution with d_1 and d_2 degrees of freedom.

The Hotelling's $T^2(d_1, d_2)$ is a simple rescaling of the F-distribution.

Definition 2.5.3. *The Hotelling's distribution $T^2(d_1, d_2)$ satisfies*

$$T^2(d_1, d_2) = \frac{d_1 d_2}{d_2 - d_1 + 1} F_{d_1, d_2 - d_1 + 1}.$$

Given a sample $\mathbf{x}_1, \dots, \mathbf{x}_n$ from $N_m(\mu, \Sigma)$ with $n \geq m$ (so that S_n is invertible) define the following statistic:⁸

$$D^2 = (n-1)(\bar{\mathbf{x}}_n - \mu)^\top S_n^{-1}(\bar{\mathbf{x}}_n - \mu),$$

Remark 2.5.4. Note that D^2 is invariant under nonsingular linear transformations $\mathbf{x} \mapsto A\mathbf{x} + \mathbf{b}$.

Theorem 2.5.5. We have $D^2 \sim T^2(p, n-1)$.

We omit the proof. If you are interested in details, check Corollary 4.5.6 in ⁹.

2.5.3 Multivariate Test for Equality of Means

Consider now a sample $\mathbf{x}_1, \dots, \mathbf{x}_n$ from $N_p(\mu_x, \Sigma)$ and a sample $\mathbf{y}_1, \dots, \mathbf{y}_m$ from $N_p(\mu_y, \Sigma)$. We can extend the ideas from the univariate case to define a multivariate test. Let

$$S = \frac{1}{n+m-2} [nS_x + mS_y],$$

where S_x and S_y are the sample covariance matrices for the two samples¹⁰. Define the test statistic:

$$D^2 = (\bar{\mathbf{x}}_n - \bar{\mathbf{y}}_m)^\top S^{-1}(\bar{\mathbf{x}}_n - \bar{\mathbf{y}}_m).$$

Theorem 2.5.6. Under $H_0 : \mu_x = \mu_y$, we have

$$\frac{nm}{n+m} D^2 \sim T^2(p, n+m-2).$$




Figure 2.2: The density of F_{d_1, d_2} for different combinations of d_1, d_2 .

⁸ Recall that $(\bar{\mathbf{x}}_n - \mu) \perp\!\!\!\perp S_n$.

⁹ Kanti V Mardia, John T Kent, and Charles C Taylor. *Multivariate analysis*, volume 88. John Wiley & Sons, 2024

¹⁰ Note that S is an unbiased estimate of Σ .

For details, if you are interested, see Theorem 4.6.1 in ¹¹.

In practice, we can implement this test in R using the function `HotellingsT2(X, Y)` in the package `ICSNP`. Note that the assumption of equal covariance matrices for the two populations can sometimes be problematic, but when the assumption approximately holds, this multivariate test is often more powerful than performing a series of univariate tests.

¹¹ Kanti V Mardia, John T Kent, and Charles C Taylor. *Multivariate analysis*, volume 88. John Wiley & Sons, 2024

2.6 Gaussian Processes in Multivariate Statistics

Gaussian Processes (GPs) are a powerful tool in statistics and machine learning for modeling continuous, multivariate data. A Gaussian Process is a generalization of the Gaussian distribution over a finite set of points to a distribution over functions. In this section we give only basic definitions and some practical examples. We omit many important computational aspects. Our goal is simply to illustrate that very basic results on the multivariate normal distribution, lead to powerful techniques.

2.6.1 Definition of a Gaussian Process

A Gaussian Process on the set T can be defined as follows:

Definition 2.6.1. *A Gaussian Process on a set T is a collection of random variables $\mathbb{X}_T := (X_t)_{t \in T}$ such that for any $n \in \mathbb{N}$ and any collection of points t_1, \dots, t_n in T , the vector $(X_{t_1}, \dots, X_{t_n})$ has a multivariate normal distribution.*

Effectively a Gaussian process \mathbb{X}_T is defined by specifying the **mean function** $m : T \rightarrow \mathbb{R}$ such that $m(t) = \mathbb{E}X_t$ and the kernel function $k : T \times T \rightarrow \mathbb{R}$ such that $k(t, t') = \text{cov}(X_t, X_{t'})$.

Although T can be pretty much arbitrary, we will assume that it is a metric space; often $T = \mathbb{R}^m$.

The kernel function must be symmetric ($k(t, t') = k(t', t)$) and it satisfies that for all $n \in \mathbb{N}$ and all $t_1, \dots, t_n \in T$ the matrix $\Sigma \in \mathbb{S}^n$ defined by $\Sigma_{ij} = k(t_i, t_j)$ is positive definite. Such functions are very well studied in mathematics and their theory is very rich.

The choice of the kernel function $k(t, t')$ plays a crucial role in determining the properties of functions sampled from the GP, such as smoothness and periodicity. Typically the kernel functions $k(\mathbf{x}, \mathbf{x}')$ considered are decreasing functions of the distance $\|\mathbf{x} - \mathbf{x}'\|$ between \mathbf{x} and \mathbf{x}' . Roughly speaking, if the kernel function is such that the correlations between points drops slowly with the distance, the realizations of the corresponding Gaussian process will be relatively smooth.

Popular choices include the squared exponential kernel and the Matérn kernel. If $T \subseteq \mathbb{R}^m$ then the exponential kernel is defined as

$$k_E(\mathbf{x}, \mathbf{x}') = \sigma^2 \exp\left(-\frac{1}{\ell} \|\mathbf{x} - \mathbf{x}'\|\right),$$

The exponential kernel generates functions that are not differentiable. This makes it suitable for modeling rough or jagged processes.

Matérn kernel is defined as

$$k_M(\mathbf{x}, \mathbf{x}') = \sigma^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\sqrt{2\nu} \frac{\|\mathbf{x} - \mathbf{x}'\|}{\ell} \right)^\nu K_\nu \left(\sqrt{2\nu} \frac{\|\mathbf{x} - \mathbf{x}'\|}{\ell} \right),$$

where $\nu > 0$ controls the smoothness of the kernel, $\ell > 0$ is the length scale, controlling the distance over which points are correlated, $\sigma^2 > 0$ is the signal variance, K_ν is the modified Bessel function of the second kind of order ν , and $\Gamma(\cdot)$ is the gamma function.

2.6.2 Example: Gaussian Processes for Spatial Data in R

Gaussian Processes are commonly applied in spatial statistics to model observations that are correlated across geographic locations. In this example, we use Gaussian Processes to model temperature data observed at various locations, capturing spatial dependencies across the region.

Suppose we observe temperatures across a grid of spatial locations. We can model the temperature at any given location as a sample from a Gaussian Process with a spatial covariance structure, such as the exponential or Matérn kernel, which are popular choices for spatial modeling.

```
# Load required packages
library(geoR)      # For spatial data simulation
library(fields)    # For spatial GP modeling and visualization

# Set up a grid of locations
set.seed(42)
n <- 100
locs <- expand.grid(x = seq(0, 1, length.out = n), y = seq(0, 1,
  length.out = n))

# Generate spatially correlated temperature data
true_cov <- exp(-rdist(locs) / 2) # Exponential covariance function
temp <- t(chol(true_cov)) %*% rnorm(n * n)

# Plot the spatial data
image(matrix(temp, n, n), main = "Simulated Temperature Data",
  xlab = "Longitude", ylab = "Latitude")
```

In this code:

- We set up a grid of $100^2 = 10,000$ locations over $[0, 1]^2$.

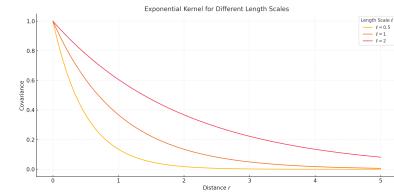


Figure 2.3: The exponential kernel as a function of $r = \|\mathbf{x} - \mathbf{x}'\|$ for different values of ℓ and $\sigma = 1$.

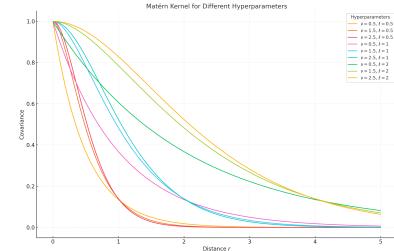


Figure 2.4: The Matérn kernel as a function of $r = \|\mathbf{x} - \mathbf{x}'\|$ for different choices of the hyperparameters.



Figure 2.5: Simulated Gaussian Process on a 2D grid with the exponential kernel.

- We generate spatially correlated temperature data with an exponential covariance function.

In simpler terms, the grid defined points $\mathbf{x}_{ij} \in [0, 1]^2$ for $i, j = 1, \dots, 100$. Now a big covariance matrix can be constructed whose entries correspond to the covariance between the sites \mathbf{x}_{ij} and \mathbf{x}_{kl} . Now we generate a single distribution from this big 10000-dimensional Gaussian distribution. The resulting simulated temperature distribution is given in Figure 2.5.

Remark 2.6.2. *Handling a large Gaussian distribution as described above can be computationally costly, so in practice, the sample is generated entry by entry by leveraging the properties of Gaussian distributions in a smart way.*

Next, we use Gaussian Processes to make predictions about the temperature at unobserved locations by fitting a GP model to our simulated data.

On a high level this is quite straightforward.

1. Add to the earlier grid any additional points you would like to include in the predictions.
2. Using the kernel function, compute the joint distribution of the training and test data.
3. Then, using the standard formulas, compute the conditional distribution of the test data given the training data (this is all Gaussian!).

The conditional mean gives you predictions and conditional covariance helps you quantify uncertainty.

In practice, you would also like to learn hyperparameters of the kernel function. This could be done using the maximum likelihood approach. For computational reasons often working with Gaussian Processes, we rely on Bayesian statistics. There are some standard packages to make this more convenient. Consider for example the following code.

```
library(spBayes)

# Fit a Gaussian Process model with the correct covariance model
fit <- spLM(temp ~ 1,
            coords = as.matrix(locs),
            starting = list("phi" = 1, "sigma.sq" = 1, "tau.sq" = 0.1),
            tuning = list("phi" = 0.1, "sigma.sq" = 0.1, "tau.sq" = 0.1),
            priors = list("phi.Unif" = c(0.01, 1),
                         "sigma.sq.IG" = c(2, 1),
                         "tau.sq.IG" = c(2, 1)),
            cov.model = "exponential", # Specify the covariance model
            n.samples = 1000)
```

```

# Predict at new locations
new_locs <- expand.grid(x = seq(0, 1, length.out = 50), y = seq(0, 1,
  ↵ length.out = 50))
pred <- spPredict(fit, pred.coords = as.matrix(new_locs))

# Visualize the predictions
image(matrix(pred$p.y.predictive, 50, 50),
  main = "Predicted Temperature Data",
  xlab = "Longitude",

```

In this step:

- We fit a Gaussian Process model to the spatial temperature data using the `spLM` function from the `spBayes` package.
- We make predictions at a set of new locations and visualize the GP-based temperature predictions.

2.6.3 Example: Nonparametric regression using GPs*

In machine learning, Gaussian Processes are often used for non-parametric regression. This is outside of scope of this class so we provide this example only to complete the picture but it can be safely omitted.

Consider the following code. Here, unlike in the previous example, we code all relevant formulas from scratch.

```

# Load required libraries
library(MASS)      # For multivariate normal sampling
library(ggplot2) # For fancy plotting

# Define a squared exponential (RBF) kernel
rbf_kernel <- function(x, y, length_scale = 1, sigma_f = 1) {
  sigma_f^2 * exp(-0.5 * (outer(x, y, "-")^2) / length_scale^2)}

# Generate training data from noisy sine function
set.seed(42)
n_train <- 10
x_train <- seq(0, 10, length.out = n_train)
y_train <- sin(x_train) + rnorm(n_train, sd = sqrt(0.1))

# Define test points where predictions are needed
x_test <- seq(0, 10, length.out = 100)

# Compute covariance matrices
length_scale <- 1 # Length scale of the RBF kernel
sigma_f <- 1       # Signal variance
sigma_n <- sqrt(0.1) # Noise variance

K <- rbf_kernel(x_train, x_train, length_scale, sigma_f) + diag(sigma_n^2,
  ↵ n_train)
K_s <- rbf_kernel(x_train, x_test, length_scale, sigma_f)
K_ss <- rbf_kernel(x_test, x_test, length_scale, sigma_f)

# Predictive mean and covariance

```

```

K_inv <- solve(K)
mu_s <- t(K_s) %*% K_inv %*% y_train # Transpose K_s for predictive mean
cov_s <- K_ss - t(K_s) %*% K_inv %*% K_s # Correct covariance calculation

# Sample from the posterior predictive distribution
n_samples <- 5
y_samples <- mvrnorm(n_samples, mu_s, cov_s)

# Create a data frame for the test points and GP predictions
gp_data <- data.frame(x = x_test, GP_Mean = as.vector(mu_s), True_Function =
  sin(x_test))

# Add the GP samples to the data frame
gp_samples_df <- data.frame(x = rep(x_test, times = n_samples), Sample =
  factor(rep(1:n_samples, each = length(x_test))), Value =
  as.vector(t(y_samples)))

# Create a data frame for the training points
train_data <- data.frame(x = x_train, y = y_train)

# Plot using ggplot2
ggplot() +
  # True function
  geom_line(data = gp_data, aes(x = x, y = True_Function),
            color = "black", linetype = "dashed", size = 1, alpha = 0.7) +
  # GP mean
  geom_line(data = gp_data, aes(x = x, y = GP_Mean),
            color = "lightpink2", size = 1.2) +
  # GP samples
  geom_line(data = gp_samples_df, aes(x = x, y = Value, group = Sample),
            color = "springgreen3", alpha = 0.5) +
  # Training points
  geom_point(data = train_data, aes(x = x, y = y),
             color = "blue", size = 2) +
  # Labels and theme
  labs(title = "Gaussian Process Regression",
       x = "x", y = "f(x)") +
  theme_minimal(base_size = 14) +
  theme(legend.position = "none")

```

This code first defines a grid of $n = 10$ equally spaced points x_1, \dots, x_n in the interval $[0, 10]$. For each of these points, we generate an observation $y_i \sim N(f(x_i), \sigma^2)$ with $\sigma^2 = 0.1$. These are the blue dots in Figure 2.6. These are our noisy observations from the function $f(x) = \sin(x)$ (the dashed line).

Having the training data, the test data and a fixed kernel function (here the Gaussian kernel) we can compute the covariance matrix in the underlying Gaussian process (over all the point). Then we can sample from the conditional distribution¹² of the test data given the training data. Doing it multiple times gives us the greenish curves in Figure 2.6. We can also compute the conditional mean and also plot it. This is the pink curve.

¹² All the relevant formulas already appeared earlier!




Figure 2.6: Illustration for GPs applied

2.7 Exercises

Exercise 2.7.1. Two random vectors $X \in \mathbb{R}^p$ and $Y \in \mathbb{R}^q$ are independent if $\mathbb{P}(X \in A, Y \in B) = \mathbb{P}(X \in A)\mathbb{P}(Y \in B)$ for all (measurable) sets $A \subseteq \mathbb{R}^p$, $B \subseteq \mathbb{R}^q$. Let $f : \mathbb{R}^p \rightarrow \mathbb{R}$ and $g : \mathbb{R}^q \rightarrow \mathbb{R}$ be two functions. Show that if $X \perp\!\!\!\perp Y$ then $f(X) \perp\!\!\!\perp g(Y)$.

Exercise 2.7.2. Given $X \sim N_m(\mu, \Sigma)$, show that $X - \mu$ has the same distribution as $\Sigma^{1/2}Z$, where $Z \sim N_m(0, I_m)$.

Exercise 2.7.3. Suppose we want to generate a sample of n observations from $N_m(0, \Sigma)$. Using R or Python, write a code that does it by sampling independently a bunch of univariate $N(0, 1)$ variables and transforming them appropriately.

Exercise 2.7.4. Consider two random vectors $X \sim N_m(\mu_X, \Sigma_X)$ and $Y \sim N_q(\mu_Y, \Sigma_Y)$. Show that if X and Y are independent, the joint distribution of (X, Y) is multivariate normal with mean (μ_X, μ_Y) and block diagonal covariance matrix.

Exercise 2.7.5. Let $X_1, X_2 \sim N_m(\mu, \Sigma)$ with $X_1 \perp\!\!\!\perp X_2$, and define $Y = X_1 + X_2$. Find the distribution of Y .

Exercise 2.7.6. Let \mathbf{Y} be an $n \times p$ matrix of observed responses, where each row \mathbf{y}_i is a p -dimensional random vector from a multivariate normal distribution: $\mathbf{y}_i \sim N_p(\mu, \Sigma)$, and the rows are independent. Assume that we model the mean of \mathbf{Y} as a linear function of a set of m predictors given in the $n \times m$ design matrix \mathbf{X} , i.e.,

$$\mathbb{E}[\mathbf{Y}] = \mathbf{XB},$$

where B is a $m \times p$ matrix of regression coefficients. The residuals are defined as:

$$\mathbf{E} = \mathbf{Y} - \mathbf{XB}\hat{B},$$

where $\hat{B} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}$ is the ordinary least squares estimator of B . The i -th row of \mathbf{E} is denoted as \mathbf{E}_i , which represents the residual vector for the i -th observation. Derive the covariance structure of the residuals by explicitly computing:

$$\text{cov}(\mathbf{E}_i, \mathbf{E}_j),$$

for all $i, j \in \{1, \dots, n\}$. Express your result in terms of Σ , \mathbf{X} , and the projection matrix $\mathbf{P} = \mathbf{X}(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$.

Exercise 2.7.7. Let $X \sim N_m(\mu, \Sigma)$. Use the spectral decomposition of Σ to transform X into independent standard normal variables.

Exercise 2.7.8. Let $X \sim N_m(\mathbf{0}, \Sigma)$ and suppose that Σ is singular (i.e., Σ is not invertible). Prove that in this case, there exists a linear subspace

$L \subseteq \mathbb{R}^m$ such that $\mathbb{P}(X \in L) = 1$. If L is the minimal linear subspace with this property then what is its dimension? Hint: $\text{var}(u^\top X) = 0$ if and only if $\mathbb{P}(u^\top X = 0) = 1$, and algebraically, u is an eigenvector of Σ with eigenvalue 0.

Exercise 2.7.9. Let $X \sim N_m(0, \sigma^2 I_m)$ and let U be any orthogonal matrix. Show that the distribution X is the same as the distribution of UX . Would the answer be the same if the mean of X was not zero?

Exercise 2.7.10. Let $X_1, X_2 \sim N_m(\mu, \Sigma)$ be independent. Consider the random vector $Z = X_1 + X_2$. Prove that the covariance matrix of Z is 2Σ .

Exercise 2.7.11. Suppose $X \sim N_m(\mu, \Sigma)$. Let $X_{\setminus 1}$ denote X with the first entry removed. Consider the linear regression of X_1 on the remaining variables

$$X_1 = \alpha + \mathbf{w}^\top X_{\setminus 1} + \varepsilon \quad \text{with} \quad \varepsilon \perp\!\!\!\perp X_{\setminus 1}.$$

Show that the linear regression coefficients \mathbf{w} can be expressed in terms of the blocks of the covariance matrix Σ . Hint: What is the relationship between these coefficients and conditional variances?

Exercise 2.7.12. Suppose $X \sim N_m(\mu, \Sigma)$. Show that the correlation matrix $R = D^{-1/2}\Sigma D^{-1/2}$, where D is the diagonal matrix of variances, is also positive semidefinite. Show that the space of correlation matrices is a convex set.

Exercise 2.7.13. Prove Proposition 2.1.2. Hint: Let $Z = \Sigma^{-1/2}(X - \mu)$. Show that $Z \sim N_m(\mathbf{0}, I_m)$ and use the fact that the sum of the squares of independent standard normal variables has a chi-squared distribution.

Exercise 2.7.14. Complete all the details of the proof of Proposition 2.2.1.

Exercise 2.7.15. Prove Lemma 2.2.2.

Exercise 2.7.16. Suppose $X_1, X_2 \sim N_m(0, \Sigma)$ are independent and $Y = X_1 - X_2$. Derive the distribution of Y . How does the covariance structure of Y compare to that of X_1 and X_2 individually?

Exercise 2.7.17. Let $X \sim N_m(\mu, \Sigma)$. Show that for any $a \in \mathbb{R}^m$, the probability $\mathbb{P}(a^\top X > c)$ depends on both $a^\top \mu$ and $a^\top \Sigma a$. Derive a formula for this probability in terms of the c.d.f. of the standard normal distribution.

Exercise 2.7.18. Consider the functions $f, g : \mathbb{R}^n \rightarrow \mathbb{R}$ given by $f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x}$ and $g(\mathbf{x}) = \mathbf{x}^\top A \mathbf{x}$ for $\mathbf{a} \in \mathbb{R}^n$ and $A \in \mathbb{S}^n$. Show that $\nabla f(\mathbf{x}) = \mathbf{a}$ and $\nabla g(\mathbf{x}) = 2A\mathbf{x}$.

3

Non-normal distributions

Multivariate Gaussian is obviously not the only distribution widely considered to model multivariate continuous data. In this chapter, we provide a very brief treatment of some of the key ideas involving distributions beyond the multivariate normal.

3.1 Elliptical distributions

3.1.1 Spherical distributions

Recall that $O(m)$ denotes the group^{1,2} of $m \times m$ orthogonal matrices, that is, $U \in \mathbb{R}^{m \times m}$ such that $U^\top U = I_m$. A random vector $X \in \mathbb{R}^m$ is said to have a *spherical distribution* if for any $U \in O(m)$, the random vector X is equal in distribution to UX , i.e.,

$$X \stackrel{d}{=} UX \quad \forall U \in O(m).$$

One example of a spherical distribution is the standard normal distribution $Z \sim N_m(0, I_m)$. More generally, if $Z \sim N_m(0, I_m)$ and $\tau > 0$ is a positive random variable independent of Z , then³

$$X = \frac{1}{\sqrt{\tau}} Z$$

has a spherical distribution. To see this, let $U \in O(m)$. Then

$$UX = \frac{1}{\sqrt{\tau}} UZ \stackrel{d}{=} \frac{1}{\sqrt{\tau}} Z = X,$$

since $UZ \stackrel{d}{=} Z$ and $UZ \perp\!\!\!\perp \tau$. Therefore, X has the same distribution as UX , satisfying the definition of spherical symmetry.

Moment structure: The spherical symmetry implies that the mean of X must be zero and the covariance matrix must be proportional to the identity matrix, i.e.,

$$\mathbb{E}(X) = 0, \quad \text{var}(X) = cI_m$$

for some constant $c \geq 0$. In the case of $X = \frac{1}{\sqrt{\tau}} Z$ with $Z \sim N(0, I_m)$, we have $\text{var}(X) = \mathbb{E}[\tau^{-1}]I_m$.

¹ The set $O(m)$ is a group because $U, V \in O(m)$ then $UV^{-1} \in O(m)$. Check it!

² Note that if $U \in O(m)$ then $|\det(U)| = 1$.

³ This is sometimes called the common variance model. Note that the entries of Z are independent. Although the entries of X will be uncorrelated (can you show it?), they will not be independent.

← Exercise 3.5.2

← Exercise 3.5.4

3.1.2 Independence of $\|X\|$ and $X/\|X\|$

An important property of spherical distributions is the *independence* of the norm of X , denoted $\|X\| := \sqrt{X^\top X}$, and the direction of X , represented by the unit vector $X/\|X\|$. To see this, note that for any $U \in O(m)$,

$$\frac{X}{\|X\|} \stackrel{d}{=} \frac{UX}{\|UX\|} = \frac{UX}{\|X\|} = U \frac{X}{\|X\|}.$$

Thus, $\frac{X}{\|X\|}$ is invariant under orthogonal transformations, meaning it must have a rotationally symmetric distribution. The only rotationally symmetric distribution on the unit sphere is the uniform distribution. Therefore, $\frac{X}{\|X\|}$ has a uniform distribution on the unit sphere in \mathbb{R}^m , independent of the length $\|X\|$.

3.1.3 Formal argument using polar coordinates*

To formally show that $\|X\|$ and $X/\|X\|$ are independent, we use a change of coordinates to *polar coordinates*.

Polar coordinates: In \mathbb{R}^2 , polar coordinates are given by:

$$x_1 = r \cos(\theta), \quad x_2 = r \sin(\theta), \quad r > 0, \quad 0 \leq \theta < 2\pi.$$

This generalizes to \mathbb{R}^m with the transformation:

$$\mathbf{x} = r\mathbf{u}(\boldsymbol{\theta}),$$

where $\boldsymbol{\theta} = (\theta_1, \dots, \theta_{m-1})$ are the angular coordinates, and $r = \|\mathbf{x}\|$ is the radial coordinate. Specifically, for $i = 1, \dots, m$ we have⁴:

⁴ Check that $\|\mathbf{x}\| = r$.

$$u_i(\boldsymbol{\theta}) = \cos(\theta_i) \prod_{j=0}^{i-1} \sin(\theta_j), \quad \sin(\theta_0) = 1, \quad \cos(\theta_m) = 1.$$

The angles θ_j are constrained by:

$$0 \leq \theta_j \leq \pi, \quad j = 1, \dots, m-2, \quad 0 \leq \theta_{m-1} < 2\pi.$$

Recall that for a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, its Jacobian at $\mathbf{x} \in \mathbb{R}^n$, if exists, is the matrix $Jf(\mathbf{x})$ of the partial derivatives $Jf(\mathbf{x}) = [\frac{\partial f_i}{\partial x_j}] \in \mathbb{R}^{m \times n}$. The determinant of the Jacobian matrix of the polar coordinates transformation $\mathbf{x} = r\mathbf{u}(\boldsymbol{\theta})$ is

$$J(r, \boldsymbol{\theta}) = r^{m-1} \prod_{i=2}^{m-1} \sin^{m-i}(\theta_{i-1}).$$

The reason why we care about this determinant is that it appears in the change of variables formula.

Independence of radial and angular components: Suppose X has a density $f(\mathbf{x})$ that only depends on $\|\mathbf{x}\|$, i.e., $f(\mathbf{x}) = g(\|\mathbf{x}\|^2)$ for some function g . When changing variables to polar coordinates, we find:

$$f(\mathbf{x})d\mathbf{x} = g(r^2)r^{m-1}dr \prod_{i=2}^{m-1} \sin^{m-i}(\theta_{i-1})d\theta.$$

Thus, we see that the density separates into a product of a radial component $g(r^2)r^{m-1}dr$ and an angular component, showing that $\|\mathbf{x}\|$ and $\mathbf{x}/\|\mathbf{x}\|$ are independent.

3.1.4 Elliptical distributions as affine transformations of spherical distributions

An *elliptical distribution* is a generalization of the multivariate normal distribution. A random vector $X \in \mathbb{R}^m$ is said to have an elliptical distribution, denoted $X \sim E(\mu, \Sigma)$, if X can be written as:

$$X = \mu + \Sigma^{1/2}Z,$$

where Z is a spherical random vector (e.g., standard normal) and $\Sigma^{1/2}$ is the square root of a positive semi-definite matrix Σ . This form includes the multivariate normal as a special case; recall Proposition 2.1.1.

Why do we care? The main reason is that elliptical distributions allow us to model distributions that look similar to multivariate normal but have much heavier tails. This makes it suitable to model extreme events or data with outliers, which are common in many real-world applications such as finance, insurance, and environmental studies. For example, in finance, the heavy tails of elliptical distributions can capture the higher likelihood of extreme losses or gains compared to the Gaussian model. By maintaining properties like symmetry and linear correlation structures, elliptical distributions strike a balance between flexibility and mathematical tractability, making them a valuable tool for robust modelling in high-dimensional settings.

3.1.5 Covariance and correlation matrix

Note that in elliptical distributions, Σ is called the *scale matrix* and does not directly represent the covariance of X . However, it holds that:

$$\text{Var}(X) = c\Sigma$$

for some constant $c > 0$. Therefore, the correlation structure of X is governed by Σ , even though the exact covariance requires calculating c .

3.1.6 Scale mixture of normals

A key example of elliptical distributions is the **scale mixture of normals**, which admits the stochastic representation:

$$X = \mu + \frac{1}{\sqrt{\tau}} \Sigma^{1/2} Z, \quad (3.1)$$

where $Z \sim N_m(0, I_m)$ and $\tau > 0$ is a random variable independent of Z . This framework captures several important distributions:

- If $\tau \equiv 1$, we recover the multivariate normal distribution.
- If $\tau \sim \frac{1}{k} \chi_k^2$, then X follows a multivariate t -distribution with k degrees of freedom.
- If $k = 1$, we obtain the multivariate Cauchy distribution.
- If $\tau \sim \text{Exp}(1)$, we get the multivariate Laplace distribution.

Note that X may be a complicated distribution but conditionally on τ , X is simply Gaussian. Thus, many efficient approaches to working with scale mixture of normals, use this structure.

Note that, if X satisfies (3.1), the Mahalanobis distance satisfies

$$Y = \|X - \mu\|_\Sigma = \frac{1}{\sqrt{\tau}} \|Z\|^2 \stackrel{d}{=} \frac{1}{\sqrt{\tau}} \chi_m^2$$

with $\tau \perp\!\!\!\perp \chi_m^2$. This generalizes Proposition 2.1.2. We can plot the distribution of $Y = \|X - \mu\|_\Sigma^2$ for different choices of τ and $m = 10$; see Figure 3.2. Note that this picture does entirely show how dramatically different is the tail behavior of all the density plots. The x-axis finishes at 75 and the mass of the χ_1^2 density (corresponding to the Gaussian case) is negligible. This is not however true for some of the other densities. For the multivariate Laplace distribution this probability is 0.124, which is definitely not negligible. As Table 3.1 show, much more extreme events can happen with non-negligible probability.

Case	$P(Y > 75)$	$P(Y > 500)$	$P(Y > 1000)$	$P(Y > 10000)$
Gaussian	0.000	0.000	0.000	0.000
t_{100}	0.000	0.000	0.000	0.000
t_{20}	0.000	0.000	0.000	0.000
t_5	0.019	0.000	0.000	0.000
Laplace	0.124	0.020	0.010	0.001
t_1	0.277	0.109	0.077	0.024

To see why this may be relevant in practice, consider a simple example of returns for four stocks. We can compute the sample version of the Mahalanobis distance and compare its distribution to



Figure 3.1: Examples of 1-d densities that are generalized by the scale mixture of normals. Note that the tails are generally heavier than for the Gaussian distribution.

Table 3.1: Proportion of Samples Exceeding Thresholds

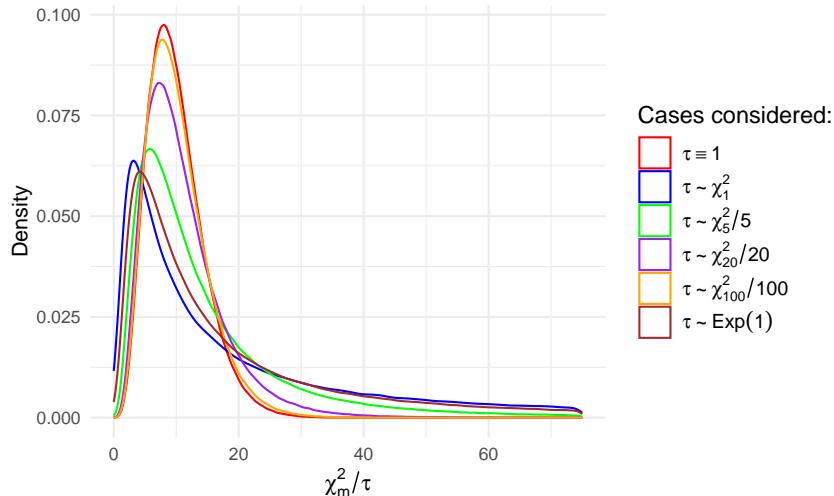


Figure 3.2: The distribution of $\|X - \mu\|_\Sigma^2$ for some choices of τ and with $m = 10$.

χ_4^2 , which corresponds to the Gaussian case. Consider the following code.

```
# Load necessary libraries
library(quantmod) # Fetch stock data
library(MASS) # Mahalanobis distance
library(ggplot2) # Plotting
library(latex2exp) # LaTeX labels
library(dplyr) # Data manipulation

# Define stock tickers and date range
tickers <- c("AAPL", "MSFT", "GOOG", "AMZN") # Add more stocks if needed
start_date <- as.Date("2020-01-01")
end_date <- as.Date("2024-01-01")

# Download stock prices
getSymbols(tickers, src = "yahoo", from = start_date, to = end_date,
           auto.assign = TRUE)

# Convert to adjusted closing prices
prices <- do.call(merge, lapply(tickers, function(ticker)
  ↪ c(get(ticker))))
colnames(prices) <- tickers

# Compute daily log-returns
returns <- na.omit(diff(log(prices)))

# Compute Mahalanobis distances
mu <- colMeans(returns) # Mean vector
sigma <- cov(returns) # Covariance matrix
mahalanobis_distances <- mahalanobis(returns, mu, sigma)

# Degrees of freedom (equal to number of assets)
df <- ncol(returns)

# Generate theoretical chi-squared density
x_seq <- seq(0, max(mahalanobis_distances), length.out = 500)
chi2_density <- dchisq(x_seq, df = df)
```

```

# Compute kernel density estimation (KDE) for Mahalanobis distances
density_mahalanobis <- density(mahalanobis_distances)

# Define tail thresholds
thresholds <- c(10, 20, 50, 75)

# Compute tail probabilities
tail_probs <- data.frame(
  Threshold = thresholds,
  Empirical_Prob = sapply(thresholds, function(t)
    mean(mahalanobis_distances > t)),
  ChiSq_Prob = sapply(thresholds, function(t) 1 - pchisq(t, df = df))
)

# Print the table
print(round(tail_probs, 3))

# Plot: Density of Mahalanobis Distance vs. Chi-squared Distribution
ggplot() +
  geom_line(aes(x = density_mahalanobis$x, y = density_mahalanobis$y),
            color = "blue", linewidth = 1.2, label = "Empirical Density") +
  geom_line(aes(x = x_seq, y = chi2_density), color = "red", linetype =
            "dashed", linewidth = 1, label = "Chi-squared Density") +
  labs(
    title =
      TeX("Density of Mahalanobis Distance vs. $\chi^2_d$ Density"),
    x = TeX("Mahalanobis Distance"),
    y = "Density"
  ) +
  theme_minimal() +
  annotate("text", x = max(density_mahalanobis$x) * 0.6, y =
           max(density_mahalanobis$y) * 0.9,
           label = "Empirical Density", color = "blue", size = 5) +
  annotate("text", x = max(density_mahalanobis$x) * 0.6, y =
           max(chi2_density) * 0.9,
           label = "Chi-squared Density", color = "red", size = 5)

```

In Figure 3.3 we compare the sample density of the Mahalanobis distance with the density of χ_4^2 , which corresponds to the Gaussian case. From the picture it looks that the sample density is much more concentrated around zero. However, we can easily confirm that this distribution has much heavier tails. For example, the mass of the χ_m^2 over the interval $(25, \infty)$ is totally negligible. On the other hand the empirical distribution has there mass 0.022.

3.2 Copula models

A copula is a statistical tool that allows us to model the dependence structure between random variables separately from their marginal distributions. This is especially useful when modeling multivariate data where the marginal distributions are not Gaussian.⁵

Definition 3.2.1. A *copula function* (or simply a *copula*) is a multivariate distribution function with uniform marginals. A copula function in m dimensions is typically written $C(\mathbf{u})$, $\mathbf{u} = (u_1, \dots, u_m)$.

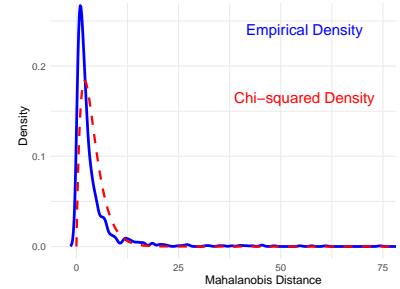


Figure 3.3: Comparison of the sample distribution of the Mahalanobis distance and the density of χ_4^2 .

⁵ Recall that the CDF of the uniform $U(0, 1)$ distribution is $F(u) = u$ and so $C : [0, 1]^m \rightarrow [0, 1]$ must satisfy $C(u_1, 1, \dots, 1) = u_1$.

3.2.1 Sklar's theorem

The foundation of copula theory is *Sklar's Theorem*.

Theorem 3.2.2 (Sklar, 1959). *Given a continuous random vector $X = (X_1, \dots, X_m)$ with joint c.d.f. F and marginal c.d.f.s F_1, \dots, F_m , there exists a unique copula C such that for all $\mathbf{x} = (x_1, \dots, x_m)$*

$$F(x_1, \dots, x_m) = C(F_1(x_1), \dots, F_m(x_m)). \quad (3.2)$$

Conversely, for any given one-dimensional c.d.f.s F_1, \dots, F_m and a copula C , (3.2) defines an m -variate c.d.f. F with marginal c.d.f.s F_1, \dots, F_m .

The copula $C : [0, 1]^m \rightarrow [0, 1]$ captures the dependence between the components of \mathbf{X} , while the marginal distributions F_1, \dots, F_m describe the behavior of each component independently.

To understand Sklar's theorem, consider the case when $m = 1$. Here $C(u) = u$ for $u \in [0, 1]$ and the CDF of X is indeed $C(F(x)) = F(x)$. Note also that, if X is a univariate continuous random variable with c.d.f. F^6 then $F(X) \sim U(0, 1)$. Indeed,

$$\mathbb{P}(F(X) \leq u) = \mathbb{P}(X \leq F^{-1}(u)) = F(F^{-1}(u)) = u.$$

Let now $X = (X_1, \dots, X_m)$ be a random vector with c.d.f.

$$F(x_1, \dots, x_m) = \mathbb{P}(X_1 \leq x_1, \dots, X_m \leq x_m).$$

Set $U_i = F_i(X_i)$. Then $U = (U_1, \dots, U_m)$ has marginal uniform distributions with c.d.f.

$$\begin{aligned} \mathbb{P}(U_1 \leq u_1, \dots, U_m \leq u_m) &= \mathbb{P}(X_1 \leq F_1^{-1}(u_1), \dots, X_m \leq F_m^{-1}(u_m)) \\ &= F(F_1^{-1}(u_1), \dots, F_m^{-1}(u_m)) \\ &=: C(\mathbf{u}). \end{aligned} \quad (3.3)$$

Note that C defined here is the same as C in (3.2) and Sklar's theorem states that such C is defined uniquely.

3.2.2 Simple example

Consider the following joint cumulative distribution function (CDF):

$$F_{X,Y}(x, y) = \begin{cases} 0 & \text{if } x < 0 \text{ or } y < 0, \\ x^2y^2 & \text{if } 0 \leq x \leq 1 \text{ and } 0 \leq y \leq 1, \\ 1 & \text{if } x > 1 \text{ and } y > 1, \\ \min(x^2, y^2) & \text{otherwise.} \end{cases}$$

The marginal CDFs are:

$$F_X(x) = \begin{cases} 0 & \text{if } x < 0, \\ x^2 & \text{if } 0 \leq x \leq 1, \\ 1 & \text{if } x > 1, \end{cases}$$

⁶If X is continuous then F is strictly increasing almost surely.

and

$$F_Y(y) = \begin{cases} 0 & \text{if } y < 0, \\ y^2 & \text{if } 0 \leq y \leq 1, \\ 1 & \text{if } y > 1. \end{cases}$$

The inverse CDFs are for $u, v \in (0, 1)$:

$$F_X^{-1}(u) = \sqrt{u}, \quad F_Y^{-1}(v) = \sqrt{v}.$$

Thus the copula is given by:

$$C(u, v) = F_{X,Y}(F_X^{-1}(u), F_Y^{-1}(v)).$$

Substituting $F_X^{-1}(u) = \sqrt{u}$ and $F_Y^{-1}(v) = \sqrt{v}$, we get:

$$C(u, v) = \begin{cases} uv & \text{if } u \leq 1, v \leq 1, \\ \min(u, v) & \text{otherwise.} \end{cases}$$

3.2.3 Invariance under monotone transformations

Suppose that $Y_i = f_i(X_i)$, where each f_i is a strictly increasing (and so invertible) transformation. It is important to note that X and Y define the same copula. Indeed, let F, G be the CDFs of X and Y respectively with individual CDFs denoted as F_i, G_i . In the light of (3.3), it is enough to show that

$$F(F_1^{-1}(u_1), \dots, F_m^{-1}(u_m)) = G(G_1^{-1}(u_1), \dots, G_m^{-1}(u_m))$$

Now note that

$$G_i(y_i) = \mathbb{P}(Y_i \leq y_i) = \mathbb{P}(X_i \leq f_i^{-1}(y_i)) = F_i(f_i^{-1}(y_i)),$$

which shows that $G_i = F_i \circ f_i^{-1}$, or equivalently, $G_i^{-1} = f_i \circ F_i^{-1}$.

Similarly,

$$G(y_1, \dots, y_m) = F(f_1^{-1}(y_1), \dots, f_m^{-1}(y_m))$$

Using these two equations we get that

$$G(G_1^{-1}(u_1), \dots, G_m^{-1}(u_m)) = F(f_1^{-1}(G_1^{-1}(u_1)), \dots, f_m^{-1}(G_m^{-1}(u_m))) = F(F_1^{-1}(u_1), \dots, F_m^{-1}(u_m))$$

as claimed.

3.2.4 Density of a copula

The p.d.f. $c(\mathbf{u})$ of a copula $C(\mathbf{u})$ is obtained by differentiation. We have

$$c(\mathbf{u}) = \frac{\partial^m C(\mathbf{u})}{\partial u_1 \cdots \partial u_m}.$$

Using the fact that $C(\mathbf{u}) = F(F_1^{-1}(u_1), \dots, F_m^{-1}(u_m))$ and the chain rule, we have

$$c(\mathbf{u}) = \frac{f(\mathbf{x})}{\prod_{i=1}^m f_i(x_i)}. \tag{3.4}$$

Recall that for a continuous distribution with CDF $F(x_1, \dots, x_m)$, its density $f(x_1, \dots, x_m)$ is given by

$$f(x_1, \dots, x_m) = \frac{\partial^m}{\partial x_1 \cdots \partial x_m} F(x_1, \dots, x_m).$$

3.2.5 Gaussian copula

The *Gaussian copula* is derived from the multivariate normal distribution. Suppose $X \sim N_m(\mu, \Sigma)$ is a normal random vector. Since a copula is unchanged if components undergo a monotone transformation, w.l.o.g. we may assume $\mu_i = 0$ and $\Sigma_{ii} = 1$ for all i so that the marginal distributions are standard normal. The CDF of $N(0, 1)$ is denoted by $\Phi(x)$.

Using formula in (3.4) the corresponding density is

← Exercise 3.5.21

$$c(u; \Sigma) = \det(\Sigma)^{-1/2} \exp\left\{-\frac{1}{2}\mathbf{x}^\top (\Sigma^{-1} - I_m)\mathbf{x}\right\}, \quad \text{where } x_i = \Phi^{-1}(u_i). \quad (3.5)$$

The most popular choice for Σ is the equicorrelation matrix $\Sigma = (1 - \rho)I_m + \rho\mathbf{1}\mathbf{1}^\top$, where $-\frac{1}{m-1} \leq \rho \leq 1$. For this choice $\det(\Sigma) = (1 - \rho)^{m-1}(1 + (m-1)\rho)$. So for example the density of a two-dimensional Gaussian copula is

$$c(u; \rho) = \frac{1}{\sqrt{1 - \rho^2}} \exp\left\{-\frac{\rho^2(x_1^2 + x_2^2) - 2\rho x_1 x_2}{2(1 - \rho^2)}\right\}. \quad (3.6)$$

Example 3.2.3. We can use Gaussian copula to model the joint distribution of $X_1 \sim \text{Exp}(1)$ and $X_2 \sim \text{Exp}(1)$. The c.d.f. of $\text{Exp}(1)$ is $F(x) = 1 - e^{-x}$. The Gaussian copula is defined as: $C_\rho(u_1, u_2) = \Phi_2(\Phi^{-1}(u_1), \Phi^{-1}(u_2); \rho)$, where Φ_2 is the c.d.f. of the bivariate normal distribution with variances 1 and correlation ρ . By Sklar's theorem, the joint c.d.f. of X_1 and X_2 is then:

$$F_{X_1, X_2}(x_1, x_2) = C_\rho(F_{X_1}(x_1), F_{X_2}(x_2)).$$

To simulate data from this model:

1. Generate a sample from a 2D normal distribution with correlation ρ :

$$(Z_1, Z_2) \sim N_2\left(\mathbf{0}, \begin{bmatrix} 1 & \rho \\ \rho & 1 \end{bmatrix}\right).$$

2. Transform the marginals to the uniform distribution via the standard normal c.d.f.:

$$U_1 = \Phi(Z_1), \quad U_2 = \Phi(Z_2).$$

3. Transform the uniform marginals to the exponential distribution:

$$X_1 = -\ln(1 - U_1), \quad X_2 = -\ln(1 - U_2).$$

3.2.6 Archimedean copula*

In addition to Gaussian copulas, there are other families of copulas, such as *Archimedean copulas*, which are popular in applications due

to their simplicity. An Archimedean copula is defined using a generator function $\phi : (0, 1) \rightarrow [0, \infty)$ that is strictly decreasing and convex with $\lim_{x \rightarrow 0^+} \phi(x) = \infty$ and $\lim_{x \rightarrow 1^-} \phi(x) = 0$. Set $\psi = \phi^{-1}$, which is monotone decreasing on $(0, \infty)$ with $\lim_{x \rightarrow 0^+} \psi(x) = 1$ and $\lim_{x \rightarrow \infty} \psi(x) = 0$. Assume that ψ is completely monotone function, that is, suppose that ψ is infinitely differentiable, and that the derivatives alternate in sign: $(-1)^k \psi^{(k)}(x) > 0$ for all $x \in (0, \infty)$, $k \geq 1$.

The m -variate Archimedean copula generated by ϕ is given by:

$$C(u_1, \dots, u_m) = \psi(\phi(u_1) + \dots + \phi(u_m)).$$

It can be shown that C defines a valid c.d.f. with density

$$c(\mathbf{u}) = \psi^{(m)}(v) \prod_{i=1}^m \phi'(u_i), \quad \text{where } v_i = \phi(u_i), v = v_1 + \dots + v_m.$$

Some well-known Archimedean copulas include:

- Clayton copula: $\phi(t) = \frac{1}{\theta}(t^{-\theta} - 1)$ for $\theta > 0$.
- Gumbel copula: $\phi(t) = (-\log t)^\theta$ for $\theta \geq 1$.
- Frank copula: $\phi(t) = -\log \left(\frac{e^{-\theta t} - 1}{e^{-\theta} - 1} \right)$ for $\theta \neq 0$.

3.2.7 Example

Suppose that we observe some data $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$ from F , which we model with a copula. Just to illustrate how this can be used in real examples, we use copulas to model the dependence structure of returns of five random stocks. First plot their correlation structure.

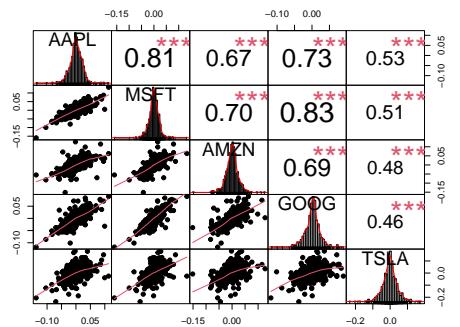
```
library(quantmod)
library(copula)
library(PerformanceAnalytics)

# Define stock tickers and date range
tickers <- c("AAPL", "MSFT", "AMZN", "GOOG", "TSLA")
start_date <- "2020-01-01"
end_date <- "2023-01-01"

# Download adjusted closing prices
getSymbols(tickers, from = start_date, to = end_date, src = "yahoo")
prices <- do.call(merge, lapply(tickers, function(x) Ad(get(x))))

# Calculate daily log returns
returns <- na.omit(ROC(prices, type = "continuous"))
colnames(returns) <- tickers

# Plot the returns
chart.Correlation(returns, histogram = TRUE, pch = 19)
```



We fit a Gaussian copula to the multivariate log returns. We start by fixing a copula model – in this case Gaussian. By (3.2) we can express the CDF F of the data in terms of the copula and the marginal

CDFs F_i . However, the marginal CDFs of the data are not known. Given the sample, it is natural to consider the empirical CDF

$$\hat{F}_i(x_i) := \frac{1}{n} \sum_{t=1}^n \mathbb{1}\{\mathbf{x}_i^{(t)} \leq x_i\}.$$

In particular,

$$\hat{F}_i(\mathbf{x}_i^{(t)}) := \frac{\text{rank of } \mathbf{x}_i^{(t)} \text{ in the } j\text{-th column of } \mathbf{X}}{n}.$$

Thus, we first convert returns to pseudo-observations by mapping through \hat{F}_i 's, $\mathbf{u}_i^{(t)} := \hat{F}_i(\mathbf{x}_i^{(t)})$. This results in a new transformed data matrix \mathbf{U} . Consider the following code.

```
# Convert to pseudo-observations
u <- pobs(as.matrix(returns))

# Fit a Gaussian copula
cop_model <- normalCopula(dim = ncol(returns)) # Gaussian copula
fit <- fitCopula(cop_model, u, method = "ml") # Maximum likelihood

# Extract fitted parameters (correlation matrix)
rho_matrix <- getSigma(fit@copula)
print("Fitted correlation matrix:")
print(rho_matrix)
```

The correlation matrix looks very simple but remember that we still have the marginal distributions, which adds additional flexibility. We can now simulate new returns based on the fitted copula model and the original marginal distributions.

```
# Simulate from the fitted Gaussian copula
simulated_copula <- rCopula(nrow(returns), fit@copula)

# Transform copula samples to the original marginals
simulated_returns <- data.frame(
  AAPL = qnorm(simulated_copula[,1], mean = mean(returns$AAPL), sd =
    + sd(returns$AAPL)),
  MSFT = qnorm(simulated_copula[,2], mean = mean(returns$MSFT), sd =
    + sd(returns$MSFT)),
  AMZN = qnorm(simulated_copula[,3], mean = mean(returns$AMZN), sd =
    + sd(returns$AMZN)),
  GOOG = qnorm(simulated_copula[,4], mean = mean(returns$GOOG), sd =
    + sd(returns$GOOG)),
  TSLA = qnorm(simulated_copula[,5], mean = mean(returns$TSLA), sd =
    + sd(returns$TSLA)))

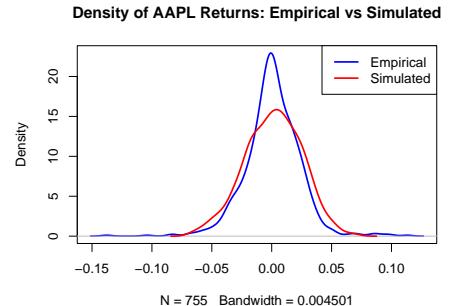
# Compare empirical and simulated correlations
empirical_corr <- cor(returns)
simulated_corr <- cor(simulated_returns)

cat("Empirical Correlation Matrix:\n")
print(empirical_corr)
cat("\nSimulated Correlation Matrix:\n")
print(simulated_corr)
```

```
> print(rho_matrix)
 [,1] [,2] [,3] [,4] [,5]
 [1,] 1.000 0.655 0.655 0.655 0.655
 [2,] 0.655 1.000 0.655 0.655 0.655
 [3,] 0.655 0.655 1.000 0.655 0.655
 [4,] 0.655 0.655 0.655 1.000 0.655
 [5,] 0.655 0.655 0.655 0.655 1.000
```

```
# Scatterplot matrix of simulated returns
pairs(simulated_returns, main = "Simulated Returns (Gaussian Copula)", pch =
  ↪  = 19, col = rgb(0, 0, 1, 0.5))

# Overlaid density plots for one stock (e.g., AAPL)
plot(density(returns$AAPL), col = "blue", main =
  ↪  "Density of AAPL Returns: Empirical vs Simulated", lwd = 2)
lines(density(simulated_returns$AAPL), col = "red", lwd = 2)
legend("topright", legend = c("Empirical", "Simulated"), col = c("blue",
  ↪  "red"), lwd = 2)
```



3.2.8 Applications of copula models

Copulas have a wide range of applications, especially in fields where understanding the dependence between random variables is crucial. Some important applications include:

- **Finance and risk management:** In finance, copulas are used to model the joint behavior of asset returns or risk factors. They are particularly useful for modeling tail dependencies, which refer to the joint likelihood of extreme losses or gains. The Gaussian copula, for instance, was famously used in pricing collateralized debt obligations (CDOs), although the limitations of copulas in capturing extreme co-movements were highlighted during the financial crisis of 2007-2008.
- **Insurance and actuarial science:** In insurance, copulas allow actuaries to model dependencies between different types of claims, such as health claims and life insurance claims. By understanding the joint risk, insurers can set more appropriate premiums and reserves for correlated risks.
- **Environmental sciences:** Copulas are used to model the dependence structure between extreme environmental events, such as floods or droughts. This is important for designing infrastructure that can withstand correlated extreme events. For instance, the joint probability of high rainfall and river overflow can be better understood with copulas, leading to more informed decisions in flood risk management.
- **Multivariate survival analysis:** In medical statistics and survival analysis, copulas are used to model the time until the occurrence of multiple types of events (such as death, disease relapse, or recovery) that may be dependent on each other. Archimedean copulas are particularly popular in this context due to their tractability.

3.3 Gaussian mixture models

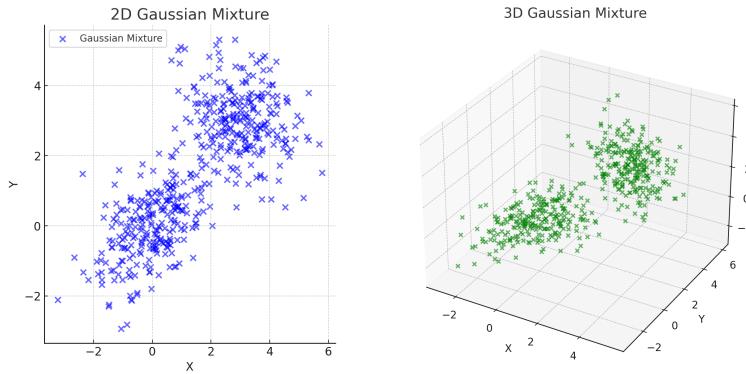
A Gaussian Mixture Model (GMM) is a probabilistic model that assumes the data is generated from a mixture of several Gaussian distributions, each with its own parameters. Formally, a random vector $X \in \mathbb{R}^m$ is said to follow a GMM if its density function is a weighted sum of multivariate Gaussian densities:

← Exercise 3.5.25

$$f(\mathbf{x}) = \sum_{k=1}^K \pi_k N_m(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k), \quad (3.7)$$

where $\pi_k \geq 0$ are the mixture weights such that $\sum_{k=1}^K \pi_k = 1$, and $N_m(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$ denotes the multivariate normal density with mean $\boldsymbol{\mu}_k \in \mathbb{R}^m$ and covariance matrix $\boldsymbol{\Sigma}_k \in \mathbb{S}_+^m$.

Figure 3.3 provides two simple examples of 2D and 3D datasets that are generated from mixtures of two Gaussian components.



Remark 3.3.1. At this point it is not clear how to handle such a complicated distribution efficiently. For example, given fixed parameters values, how can we generate samples? We will see in Section 3.3.2 that this model admits a simple latent variable representation, which can be exploited in various ways.

← Exercise 3.5.26

3.3.1 Why use Gaussian mixtures?

Gaussian mixtures offer several advantages:

- **Flexibility:** They can model complex distributions by combining simple Gaussian components. A GMM can approximate almost any continuous probability distribution by increasing the number of components K .
- **Multimodality:** GMMs are particularly useful for modeling data with multiple modes (clusters), where each Gaussian component models one mode.

- **Easily interpretable:** Each Gaussian component has an intuitive interpretation (mean, covariance, weight). It is a natural model if the population comes from different sub-populations.

Gaussian mixtures form a popular model based clustering method. In this context, for simplicity, we often assume $\Sigma_k = \Sigma$ for $k = 1, \dots, K$.

3.3.2 Gaussian Mixture as a latent variable model

Let $Z \in \{1, \dots, K\}$ be a discrete random variable such that $\mathbb{P}(Z = k) = \pi_k$ for $k = 1, \dots, K$. Moreover, suppose that the conditional distribution of X given $Z = k$ is $N_m(\mu_k, \Sigma_k)$. Then the marginal distribution of X is the same as given in (3.7) (check!).

Remark 3.3.2. Many popular probabilistic machine learning methods model data distribution $p(x)$ from a joint $p(x, z)$, which is specified by setting the marginal distribution $p(z)$ of a latent variable and the conditional $p(x|z)$. Models of that form lead to many highly tractable exact and approximate estimation procedures.

3.3.3 Likelihood inference for GMMs

Suppose we have data $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^m$ from the Gaussian mixture model with a known number of K components. Fitting a GMM to a dataset involves estimating the parameters $\theta = \{\pi_k, \mu_k, \Sigma_k\}_{k=1}^K$, which by default is done by maximizing the log-likelihood function

$$\ell(\theta) = \sum_{i=1}^n \log f(\mathbf{x}_i) = \sum_{i=1}^n \log \left(\sum_{k=1}^K \pi_k N_m(\mathbf{x}_i; \mu_k, \Sigma_k) \right).$$

The likelihood function for GMMs is a complicated multimodal function of the parameters θ and maximizing it directly may be hard. In fact it is not even bounded! However, we can use the latent representation of the GMM to perform the Expectation-Maximization (EM) algorithm.

← Exercise 3.5.29

Remark 3.3.3. Note that we assume that K is fixed throughout the whole analysis. In practice, K is not always known. In this case, learning K is a separate task, which we do not discuss here in detail.

3.3.4 Appendix: Expectation-Maximization (EM) Algorithm

Suppose we have data $\mathbf{x}_1, \dots, \mathbf{x}_n$ with distribution $f_X(\mathbf{x}; \theta)$. The goal is to estimate the parameter θ that maximizes the likelihood. In this section we deal with the situation when the model for X is specified as the margin of a larger model incorporating a latent vector Z . Many models in statistics and machine learning are of this form.

The EM algorithm is particularly useful when the marginal likelihood of X is difficult to compute, but the complete data likelihood (i.e., the joint likelihood of X and Z) is easier to work with.

Consider the complete data $(\mathbf{x}_1, \mathbf{z}_1), \dots, (\mathbf{x}_n, \mathbf{z}_n)$, where we observe only $\mathbf{x}_1, \dots, \mathbf{x}_n$ as before. Assume that the joint model depends on the parameter vector θ .

The algorithm iteratively estimates θ by alternating between two steps:

(o) Initialize θ^0 , a reasonable initial guess for the parameter.

(E-step) Compute the expected value of the complete data log-likelihood ℓ_F

$$\ell_F(\theta) = \sum_{i=1}^n \log f(\mathbf{x}_i, \mathbf{z}_i; \theta)$$

with respect to the conditional distribution of the latent variables Z , given the observed data X and the current parameter estimate θ^0 . This gives the function:

$$Q(\theta; \theta^0) = \mathbb{E}_{Z|X, \theta^0} [\ell_F(\theta) | \mathbf{X}, \theta^0] = \sum_{i=1}^n \mathbb{E}_{Z|X, \theta^0} [\log f(\mathbf{x}_i, Z; \theta) | X = \mathbf{x}_i, \theta^0].$$

(M-step) Maximize the expected complete log-likelihood with respect to θ :

$$\theta^1 = \arg \max_{\theta} Q(\theta; \theta^0).$$

Update $\theta^0 \leftarrow \theta^1$ and repeat until convergence.

Under mild conditions, it can be shown that each iteration of the EM algorithm increases the observed data log-likelihood. Specifically, after each E-step and M-step, the likelihood satisfies the inequality:

$$\ell(\theta^{t+1}) \geq \ell(\theta^t).$$

This property ensures that the algorithm makes progress toward maximizing the likelihood at every step.

While the EM algorithm is guaranteed to increase the likelihood at each step, it may not necessarily converge to the global maximum. Instead, the algorithm converges to a **local maximum** of the likelihood function. Convergence can be understood through the following results:

3.3.5 EM Algorithm for Gaussian Mixture Models (GMM)

The complete data log-likelihood for GMMs is given by:

$$\ell_F(\theta) = \sum_{i=1}^n \sum_{k=1}^K \mathbb{1}(\mathbf{z}_i = k) [\log N_m(\mathbf{x}_i; \mu_k, \Sigma_k) + \log \pi_k].$$

Here, $N_m(\mathbf{x}; \mu_k, \Sigma_k)$ denotes the Gaussian probability density function for component k with mean μ_k and covariance matrix Σ_k .

E-step

In the E-step, we compute the expected value of the complete data log-likelihood, given the current parameter estimates $\theta^0 = \{\pi_k^0, \mu_k^0, \Sigma_k^0\}$. Specifically,

$$\begin{aligned} Q(\theta; \theta_0) &= \sum_{i=1}^n \sum_{k=1}^K \mathbb{E}_{Z|X, \theta_0} [\mathbb{1}(Z=k)|X=x_i, \theta_0] [\log N_m(\mathbf{x}_i; \mu_k, \Sigma_k) + \log \pi_k] \\ &= \sum_{i=1}^n \sum_{k=1}^K \mathbb{P}(z_i = k | X = \mathbf{x}_i, \theta_0) [\log N_m(\mathbf{x}_i; \mu_k, \Sigma_k) + \log \pi_k] \end{aligned}$$

we compute the posterior probabilities, or the responsibilities, that each point \mathbf{x}_i belongs to component k :

$$w_{ik} = \mathbb{P}(z_i = k | \mathbf{x}_i, \theta^0) = \frac{\pi_k^0 N_m(\mathbf{x}_i; \mu_k^0, \Sigma_k^0)}{\sum_{l=1}^K \pi_l^0 N_m(\mathbf{x}_i; \mu_l^0, \Sigma_l^0)}.$$

Using these responsibilities, the expected complete log-likelihood is:

$$Q(\theta; \theta_0) = \sum_{i=1}^n \sum_{k=1}^K w_{ik} [\log N_m(\mathbf{x}_i; \mu_k, \Sigma_k) + \log \pi_k].$$

M-step

In the M-step, we maximize $Q(\theta; \theta_0)$ with respect to the parameters $\theta = \{\pi_k, \mu_k, \Sigma_k\}$. This leads to the following updates:

$$\begin{aligned} \pi_k &= \frac{1}{n} \sum_{i=1}^n w_{ik}, \\ \mu_k &= \frac{\sum_{i=1}^n w_{ik} \mathbf{x}_i}{\sum_{i=1}^n w_{ik}}, \\ \Sigma_k &= \frac{\sum_{i=1}^n w_{ik} (\mathbf{x}_i - \mu_k)(\mathbf{x}_i - \mu_k)^\top}{\sum_{i=1}^n w_{ik}}. \end{aligned}$$

These updates have intuitive interpretations: $\hat{\mu}_k$ is the weighted average of the data points assigned to component k , where the weights are the posterior probabilities w_{ik} ; $\hat{\Sigma}_k$ is the weighted covariance matrix of the points assigned to component k ; $\hat{\pi}_k$ is the proportion of points assigned to component k .

The EM algorithm iterates between the E-step and M-step until convergence, at which point the log-likelihood no longer increases significantly.

Deriving the updates for π_k and μ_k is relatively straightforward. The update on Σ_k requires either some persistence or some experience with calculus of matrix-functions. We leave this as an exercise.

3.3.6 Example: EM Algorithm for Gaussian Mixture Models (GMM)

Suppose we have a dataset where each observation is drawn from a mixture of two Gaussian distributions. The goal is to estimate the

parameters (means, covariances, and mixing proportions) of these Gaussian components using the Expectation-Maximization (EM) algorithm.

Assume the data consists of n points in 2D, where each observation x_i comes from one of two Gaussian distributions. Let π_k represent the mixing proportion of the k -th Gaussian component, μ_k the mean vector, and Σ_k the covariance matrix for component k , with $k = 1, 2$. The dataset is a mixture of these distributions with some overlap.

Data Simulation We can generate a 2D dataset from two Gaussian distributions, one with

$$\mu_1 = \begin{bmatrix} 2 \\ 2 \end{bmatrix}, \quad \Sigma_1 = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$$

and the other with

$$\mu_2 = \begin{bmatrix} 7 \\ 7 \end{bmatrix}, \quad \Sigma_2 = \begin{bmatrix} 1 & -0.3 \\ -0.3 & 1 \end{bmatrix}$$

The mixing parameter is $\pi_1 = 0.4 = 1 - \pi_2$. We use the following R code:

```
set.seed(123)
n <- 300 # number of points
pi_true <- c(0.4, 0.6) # true mixing proportions
mu_true <- list(c(2, 2), c(7, 7)) # true means
sigma_true <- list(matrix(c(1, 0.5, 0.5, 1), 2), matrix(c(1, -0.3, -0.3, 1), 2)) # covariances

# Sample labels from a multinomial distribution
z <- rbinom(n, size = 1, prob = pi_true[2])
x <- matrix(0, nrow = n, ncol = 2)

# Generate points from the two Gaussians
for (i in 1:n) {
  if (z[i] == 0) {
    x[i, ] <- MASS::mvrnorm(1, mu_true[[1]], sigma_true[[1]])
  } else {x[i, ] <- MASS::mvrnorm(1, mu_true[[2]], sigma_true[[2]])}
}

# Plotting the data points with color corresponding to the group
plot(x, col = z + 1, pch = 19, main =
  "Generated Data from a Gaussian Mixture", xlab = "X1", ylab = "X2")
legend("topright", legend = c("Group 1", "Group 2"), col = c(1, 2), pch =
  19)
```

This code generates a dataset from two Gaussian distributions with different means and covariances.

EM Algorithm for GMMs Now we apply the EM algorithm to estimate the parameters of the Gaussian mixture model. The following R code implements the EM algorithm:




Figure 3.4: This is how the data looks like when we color points by group assignment. In practice, this group assignment is not observed.

```

# Function to calculate Gaussian density
dGaussian <- function(x, mu, sigma) {exp(-0.5 * t(x - mu) %*% solve(sigma)
  ↪ %*% (x - mu)) / sqrt(det(2 * pi * sigma))}

# Initial estimates
K <- 2 # number of components
pi_hat <- c(0.5, 0.5) # initial guess for mixing proportions
mu_hat <- list(c(0, 0), c(1, -1)) # initial guess for means
sigma_hat <- list(diag(2), diag(2)) # initial guess for covariances

# Modified EM algorithm to extract responsibilities at specific iterations
em_gmm_with_responsibilities <- function(x, K, pi_hat, mu_hat, sigma_hat,
  ↪ steps_to_plot, max_iter = 100, tol = 1e-4) {
  n <- nrow(x)
  d <- ncol(x)

  log_likelihoods <- numeric(max_iter)
  responsibility_snapshots <- list()

  for (iter in 1:max_iter) {
    # E-step: Compute the responsibilities
    gamma <- matrix(0, n, K)
    for (i in 1:n) {
      for (k in 1:K) {
        gamma[i, k] <- pi_hat[k] * dGaussian(x[i, ], mu_hat[[k]],
          ↪ sigma_hat[[k]])
      }
      gamma[i, ] <- gamma[i, ] / sum(gamma[i, ]) # Normalize
    }

    # Save responsibilities if requested
    if (iter %% steps_to_plot) {
      responsibility_snapshots[[as.character(iter)]] <- gamma
    }

    # M-step: Update parameters
    N_k <- colSums(gamma)
    pi_hat <- N_k / n
    for (k in 1:K) {
      mu_hat[[k]] <- colSums(gamma[, k] * x) / N_k[k]
      sigma_hat[[k]] <- matrix(0, d, d)
      for (i in 1:n) {
        sigma_hat[[k]] <- sigma_hat[[k]] + gamma[i, k] * (x[i, ] -
          ↪ mu_hat[[k]]) %*% t(x[i, ] - mu_hat[[k]])
      }
      sigma_hat[[k]] <- sigma_hat[[k]] / N_k[k]
    }

    # Log-likelihood computation
    log_likelihood <- 0
    for (i in 1:n) {
      temp <- 0
      for (k in 1:K) {
        temp <- temp + pi_hat[k] * dGaussian(x[i, ], mu_hat[[k]],
          ↪ sigma_hat[[k]])
      }
      log_likelihood <- log_likelihood + log(temp)
    }

    log_likelihoods[iter] <- log_likelihood
  }
}

```

```

# Check for convergence
if (iter > 1 && abs(log_likelihoods[iter] - log_likelihoods[iter - 1]) 
  < tol) {
  log_likelihoods <- log_likelihoods[1:iter]
  break
}

list(pi = pi_hat, mu = mu_hat, sigma = sigma_hat, log_likelihood =
  log_likelihoods, responsibilities = responsibility_snapshots)
}

# Run the EM algorithm and save responsibilities after specific steps
steps_to_plot <- c(1, 2, 30)
result <- em_gmm_with_responsibilities(x, K, pi_hat, mu_hat, sigma_hat,
  steps_to_plot)

# Plotting responsibilities
par(mfrow = c(3, 1))
for (step in steps_to_plot) {
  gamma <- result$responsibilities[[as.character(step)]]
  colors <- rgb(gamma[, 1], 0, gamma[, 2], alpha = 0.6) # Red for
  # component 1, Blue for component 2
  plot(x, col = colors, pch = 19, main = paste("Step", step), xlab = "X1",
    ylab = "X2")
}
par(mfrow = c(1, 1))

# Plot log-likelihood progression
plot(result$log_likelihood, type = "o", main =
  "Log-Likelihood Progression", xlab = "Iteration", ylab =
  "Log-Likelihood")

```

When playing around with this simulation you can pay closer attention to the following problems:

- Are the estimated values of the parameters close to the true value?
- Is this model identifiable?
- In this simulation we chose $K = 2$ both in simulations and in the estimation step. In practice we do not know K so we may want to check a bunch of different values. Think how this could be done but in an ad hoc and formal way.

3.4 Higher order moments and cumulants*

For a random vector $X = (X_1, \dots, X_m)$ its moments are the expressions

$$m_{k_1 \dots k_m} := \mathbb{E}[X_1^{k_1} \cdots X_m^{k_m}]$$

and $k := k_1 + \dots + k_m$ is the order of this moment. The corresponding central moments are

$$\mathbb{E}[(X_1 - \mathbb{E}(X_1))^{k_1} \cdots (X_m - \mathbb{E}(X_m))^{k_m}].$$




Figure 3.5: Colorings of the data points according to their posterior probabilities after 1, 2, and 30 iterations of the EM algorithm.

Note that the second order central moments are covariances.

Moments can be obtained from the moment generating function

$$M_X(\mathbf{t}) = \mathbb{E}e^{\mathbf{t}^\top X}$$

by computing the corresponding higher order partial derivatives and evaluating at $\mathbf{t} = \mathbf{0}$

$$m_{k_1 \dots k_m} = \left. \frac{\partial^k}{\partial^{k_1} t_1 \dots \partial^{k_m} t_m} M_X(\mathbf{t}) \right|_{\mathbf{t}=0}.$$

Cumulants are different quantities, closely related to moments. They are defined in the same way as moments through the **cumulant generating function** $K_X(\mathbf{t}) = \log M_X(\mathbf{t})$

$$c_{k_1 \dots k_m} := \left. \frac{\partial^k}{\partial^{k_1} t_1 \dots \partial^{k_m} t_m} K_X(\mathbf{t}) \right|_{\mathbf{t}=0}.$$

It can be shown that the first order cumulants are the means, $\mathbb{E}X_i$, the second order cumulants are the covariances $\text{cov}(X_i, X_j)$, and the third order cumulants are equal to the third order central moments.

It is not hard to check that the moment generating function of $X \sim N_m(\mu, \Sigma)$ is

$$M_X(\mathbf{t}) = \exp\{\mu^\top \mathbf{t} + \frac{1}{2}\mathbf{t}^\top \Sigma \mathbf{t}\}$$

and consequently, the cumulant generating function is a quadratic polynomial

$$K_X(\mathbf{t}) = \mu^\top \mathbf{t} + \frac{1}{2}\mathbf{t}^\top \Sigma \mathbf{t}.$$

One important consequence of this fact is that all cumulants of order three or higher are zero.

One fundamental characterization of the Gaussian distribution is the following.

Theorem 3.4.1 (Marcinkiewicz). *The cumulant generating function $K_X(\mathbf{t})$ is a polynomial if and only if X is Gaussian.*

This result allows us to conceptually model departure from Gaussianity by allowing nonzero higher order cumulants. Naturally, most of the literature focuses on the third order cumulants also known as skewness.

3.5 Exercises

Exercise 3.5.1. Suppose that $\mathbf{x}, \mathbf{y} \in \mathbb{R}^m$ satisfy $\|\mathbf{x}\| = \|\mathbf{y}\|$. Show that there exists an orthogonal matrix U such that $\mathbf{y} = U\mathbf{x}$. Conclude that if $f(\mathbf{x}) = f(U\mathbf{x})$ for all $U \in O(m)$ then f depends on \mathbf{x} only through its

norm $\|\mathbf{x}\|$. Hint: Consider an orthonormal basis with $\mathbf{x}/\|\mathbf{x}\|$ as one of the basis vectors and another orthonormal basis with $\mathbf{y}/\|\mathbf{y}\|$ as one of the basis vectors.

Exercise 3.5.2. Show that if X has spherical distribution, then $\mathbb{E}X = 0$ and $\text{var}(X) = cI_m$ for some $c \geq 0$.

Exercise 3.5.3. Let $S^{m-1} = \{\mathbf{x} \in \mathbb{R}^m : \|\mathbf{x}\| = 1\}$ be the unit sphere in \mathbb{R}^m . Let Z have uniform distribution on S^{m-1} . Show that $\mathbb{E}Z = 0$ and $\text{var}(Z) = \frac{1}{m}I_m$. Hint: Z is spherical. Moreover, $\text{tr}(\text{var}(Z)) = \text{tr}(\mathbb{E}(ZZ^\top)) = \mathbb{E}(Z^\top Z)$.

Exercise 3.5.4. Suppose $X = \frac{1}{\sqrt{\tau}}Z$ with $Z \sim N(0, I_m)$, where $\tau > 0$ is a random variable independent of Z . Show that $\mathbb{E}X = 0$ and $\text{var}(X) = \mathbb{E}[\tau^{-1}]I_m$.

Exercise 3.5.5. If Z has a spherical distribution show that

1. The characteristic function $\psi_Z(\mathbf{t})$ is of the form $\phi(\|\mathbf{t}\|^2)$ for some ϕ .
2. Show that each component has characteristic function $\psi_i(s) = \phi(s^2)$.
3. Show that if one component of Z is Gaussian then Z is Gaussian.

Exercise 3.5.6 (*). Suppose X has elliptical distribution. Use the previous exercise to show that if one component of X is Gaussian then X is Gaussian.

Exercise 3.5.7 (*). Show that if Z is spherical with independent components then Z is standard normal. Hint: You will need to use the Cauchy's functional equation and the fact that the characteristic function is always continuous.

Exercise 3.5.8. Let $X \sim N_m(0, \Sigma)$. Show that $\Sigma^{-1/2}X$ has a spherical distribution. Explain why this property supports the use of spherical distributions as a generalization of the multivariate normal.

Exercise 3.5.9. Let $Z \in \mathbb{R}^m$ have spherical distribution. Show that BZ is elliptical for any fixed matrix $B \in \mathbb{R}^{p \times m}$.

Exercise 3.5.10. Suppose $X \sim E(\mu, \Sigma)$, an elliptical distribution. Show that any linear transformation $AX + b$, where $A \in \mathbb{R}^{n \times m}$ and $b \in \mathbb{R}^n$, is also elliptically distributed.

Exercise 3.5.11. Prove that if $X \sim E(0, \Sigma)$, then the correlation structure of X is determined by Σ . Explain why Σ is called the scale matrix rather than the covariance matrix.

Exercise 3.5.12. Let $X = \frac{1}{\sqrt{\tau}}Z$, where $Z \sim N_m(0, I_m)$ and $\tau \sim \chi_k^2/k$. Show that X follows a multivariate t-distribution with k degrees of freedom.

Exercise 3.5.13. The multivariate Laplace distribution can be derived as a scale mixture of normals. Specifically, for a random vector $\mathbf{X} \in \mathbb{R}^d$ following a multivariate Laplace distribution, the density can be expressed as:

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{2}{(2\pi)^{m/2} |\Sigma|^{1/2}} \int_0^\infty \exp\left(-\frac{\mathbf{x}^\top \Sigma^{-1} \mathbf{x}}{2s}\right) \frac{e^{-s/2}}{s^{(m+2)/2}} ds,$$

where Σ is the covariance matrix and s is the scaling variable.

- (a) Based on this representation, conclude that the multivariate Laplace distribution is constructed by scaling a normal distribution with an exponentially distributed random variable:

$$X | S = s \sim N_m(\mathbf{0}, s\Sigma), \quad S \sim \text{Exp}(1/2).$$

- (b) For the univariate case ($m = 1$), simplify the above expression to derive the probability density function of the univariate Laplace distribution with scale parameter $b > 0$.

- (c) Use the PDF from part (a) to compute the cumulative distribution function (CDF) of the univariate Laplace distribution. Verify that the resulting CDF is:

$$F_X(x) = \begin{cases} \frac{1}{2}e^{x/b} & \text{if } x \leq 0, \\ 1 - \frac{1}{2}e^{-x/b} & \text{if } x > 0. \end{cases}$$

Exercise 3.5.14. Consider the polar coordinates transformation for $\mathbf{X} \in \mathbb{R}^3$. Write the explicit transformation formulas and compute the Jacobian determinant for this case.

Exercise 3.5.15. Simulate data from a three-component Gaussian mixture model in \mathbb{R}^2 . Fit a GMM to the data using the EM algorithm and compare the estimated parameters to the true values. Interpret the results.

Exercise 3.5.16. Let $Z = (Z_1, Z_2) \sim N_2(0, \Sigma)$ with correlation ρ . Derive the Gaussian copula density in the bivariate case and explain how ρ governs the dependence structure.

Exercise 3.5.17. Using Sklar's theorem, show how a Gaussian copula can be used to model the joint distribution of X_1 and X_2 that both have the c.d.f. of the form $F(x) = 1 - (\frac{1}{x})^\alpha$ for some fixed shape parameter $\alpha > 0$ (Pareto distribution with scale parameter 1). Simulate data from this model and compute empirical and theoretical correlations.

Exercise 3.5.18. Derive the relationship between the density of a multivariate elliptical distribution $E(\mu, \Sigma)$ and its associated spherical distribution.

Exercise 3.5.19. Simulate data from a Clayton copula and compare its tail dependence properties to a Gaussian copula. Discuss the implications for modeling extreme events.

Exercise 3.5.20. Consider a multivariate t-distribution with k degrees of freedom. Derive its marginal distribution for a single component X_i , and compute its mean and variance.

Exercise 3.5.21. Verify (3.5) that gives the density of the Gaussian copula model. Then check also (3.6).

Exercise 3.5.22 (Properties of Copulas). Let X_1, X_2 be two continuous random variables with cumulative distribution functions (CDFs) $F_1(x)$ and $F_2(x)$, respectively.

- (a) Show that the random variables $U_1 = F_1(X_1)$ and $U_2 = F_2(X_2)$ are uniformly distributed on $[0, 1]$.
- (b) Explain why the function $C(u_1, u_2) = \mathbb{P}(U_1 \leq u_1, U_2 \leq u_2)$ defines a valid copula.

Exercise 3.5.23 (Gaussian Copula). Consider a bivariate normal random vector $(X_1, X_2) \sim N_2(0, \Sigma)$ with correlation ρ .

- (a) Derive the Gaussian copula density formula for this model.
- (b) Generate 5000 samples from a Gaussian copula with $\rho = 0.7$ and transform them to have marginal exponential distributions with mean 1. Plot the resulting scatterplot.

Exercise 3.5.24 (Fitting a Copula to Data). Download stock return data for 3 stocks (e.g., AAPL, MSFT, AMZN) and estimate their empirical copula.

- (a) Compute the empirical copula function.
- (b) Fit a Gaussian copula and compare its dependence structure to the empirical copula.
- (c) Use a Kendall's tau test to check the fit of the copula model.

Exercise 3.5.25. Equation (3.7) provides the density of the Gaussian mixture model. Verify this function indeed integrates to 1.

Exercise 3.5.26. Propose an algorithm to generate samples from this model for fixed values of all parameters. Hint: Check Section 3.3.2.

Exercise 3.5.27 (Mixture of Two Gaussians). Consider a Gaussian mixture model (GMM) in \mathbb{R}^2 with two components:

$$f(x) = \pi N_2(\mu_1, \Sigma_1) + (1 - \pi)N_2(\mu_2, \Sigma_2).$$

(a) Show that the density $f(x)$ is not necessarily unimodal.

(b) Simulate 5000 samples from a GMM with

$$\mu_1 = (2, 2), \quad \Sigma_1 = I_2, \quad \mu_2 = (-2, -2), \quad \Sigma_2 = I_2, \quad \pi = 0.5.$$

Plot the scatterplot of the samples.

(c) Apply the Expectation-Maximization (EM) algorithm to estimate the parameters.

Exercise 3.5.28 (Likelihood of a Gaussian Mixture Model). Consider a Gaussian mixture model with K components:

$$f(x) = \sum_{k=1}^K \pi_k N_m(x; \mu_k, \Sigma_k).$$

(a) Write down the log-likelihood function for a sample $\mathbf{x}_1, \dots, \mathbf{x}_n$.

(b) Show that the log-likelihood function is not convex in general.

(c) Explain why the EM algorithm can still find reasonable parameter estimates.

Exercise 3.5.29. Consider the Gaussian mixture model with a fixed number of components K . Suppose that a sample $\mathbf{x}_1, \dots, \mathbf{x}_n$ is observed and take the corresponding log-likelihood function. We will show that this function is not bounded.

1. Consider the density $f(x; \mu, \sigma^2)$ of the uniform normal $N(\mu, \sigma^2)$. What happens to $f(x; \mu, \sigma^2)$ when $\sigma^2 \rightarrow 0$ for a fixed μ ? Specifically, evaluate $f(x; \mu, \sigma^2)$ when $x = \mu$.
2. Now consider a GMM with $K = 2$ and $\pi_1 = 0.5, \pi_2 = 0.5$. Let \mathbf{x}_1 be one of the data points. Show that if one component, say $N_m(x; \mu_1, \Sigma_1)$, collapses onto \mathbf{x}_1 (i.e., $\mu_1 = \mathbf{x}_1, \Sigma_1 \rightarrow \mathbf{0}$), then the likelihood becomes arbitrarily large.

Exercise 3.5.30 (Exercise 3.2.2, Mardia et al.).⁷ Let u_0, u_1, \dots, u_m be univariate Poisson random variables with parameters $\lambda_0, \lambda_1 - \lambda_0, \dots, \lambda_m - \lambda_0$, respectively. Write down the joint distribution of $x_i = u_0 + u_i, i = 1, \dots, m$, and show that the marginal distributions of x_1, \dots, x_m are all Poisson. For $m = 2$ and writing $(x_1, x_2) = (x, y)$, show that the p.m.f. is given by

$$f(x, y) = \exp(-\lambda_1 - \lambda_2 + \lambda_0) \frac{a^x b^y}{x! y!} \sum_{r=0}^s \frac{x^{(r)}}{a^r} \frac{y^{(r)}}{b^r} \frac{\lambda_0^{(r)}}{r!},$$

where $s = \min\{x, y\}$, $a = \lambda_1 - \lambda_0$, $b = \lambda_2 - \lambda_0$, $\lambda_1 > \lambda_0 > 0$, $\lambda_2 > \lambda_0$, and $x^{(r)} = x(x-1)\cdots(x-r+1)$. Furthermore

$$\mathbb{E}(y|x) = b + \frac{\lambda_0}{\lambda_1} x, \quad \text{var}(y|x) = b + \frac{a\lambda_0}{\lambda_1^2} x.$$

⁷ Kanti V Mardia, John T Kent, and Charles C Taylor. *Multivariate analysis*, volume 88. John Wiley & Sons, 2024

Exercise 3.5.31 (Exercise 3.2.5, Mardia et al.). A random vector has a multivariate t distribution if its density is of the form

$$g_v(\mathbf{t}; \boldsymbol{\mu}, \Sigma) = \frac{c_p \det \Sigma^{-1/2}}{[1 + v^{-1}(\mathbf{t} - \boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{t} - \boldsymbol{\mu})]^{(v+p)/2}},$$

where

$$c_p = \frac{\Gamma((v+p)/2)}{(\pi v)^{p/2} \Gamma(v/2)}.$$

and v is known as the number of degrees of freedom.

- (i) Let $x_j = y_j / (S/\sqrt{v})$, $j = 1, \dots, p$, where $\mathbf{y} \sim N_p(\mathbf{0}, I_p)$, $S^2 \sim \chi_v^2$, and \mathbf{y} and S are independent. Show that \mathbf{x} has the density $g_v(\mathbf{x}; \mathbf{0}, I_p)$.
- (ii) For $v = 1$, the distribution is known as a multivariate Cauchy distribution. Show that its characteristic function is

$$\exp\{i\boldsymbol{\mu}^\top \mathbf{t} - (\mathbf{t}^\top \Sigma \mathbf{t})^{1/2}\}.$$

Exercise 3.5.32 (Exercise 3.2.7, Mardia et al.). Let \mathbf{e}_i be the i -th canonical unit vector in \mathbb{R}^m . Let \mathbf{y}_j for $j = 1, 2, \dots$ be a sequence of i.i.d. random vectors such that $\mathbb{P}(\mathbf{y} = \mathbf{e}_i) = a_i$ for $i = 1, \dots, m$; we have $a_i \geq 0$ and $\sum_i a_i = 1$.

- (i) Show that $\mathbb{E}(\mathbf{y}) = \mathbf{a}$, $\text{var}(\mathbf{y}) = \text{diag}(\mathbf{a}) - \mathbf{a}\mathbf{a}^\top$,
- (ii) Let $\mathbf{x} = \sum_{j=1}^n \mathbf{y}_j$. Show that \mathbf{x} has the multinomial distribution.
- (iii) Show that $\mathbb{E}(\mathbf{x}) = n\mathbf{a}$, $\text{var}(\mathbf{x}) = n[\text{diag}(\mathbf{a}) - \mathbf{a}\mathbf{a}^\top]$.
- (i) Verify that $[\text{diag}(\mathbf{a}) - \mathbf{a}\mathbf{a}^\top]\mathbf{1} = -$, and hence, conclude that $\text{var}(\mathbf{x})$ is singular.

Exercise 3.5.33 (Kurtosis and Moments). The kurtosis of a random variable X is defined as

$$\kappa = \frac{\mathbb{E}[X^4]}{(\mathbb{E}[X^2])^2}.$$

- (a) Compute the kurtosis of the standard normal distribution $N(0, 1)$.
- (b) Compute the kurtosis of the Laplace distribution with mean 0 and variance 1.
- (c) Compare the results and explain why the Laplace distribution has heavier tails.

Exercise 3.5.34 (Cumulants of a Normal Distribution). The cumulant generating function of a random variable X is defined as

$$K_X(t) = \log \mathbb{E}[e^{tX}].$$

- (a) Compute $K_X(t)$ for $X \sim N(0, 1)$.
- (b) Show that all cumulants of order $k \geq 3$ are zero.
- (c) Explain why this property characterizes the normal distribution.

Exercise 3.5.35 (Cumulants and Heavy Tails). Consider a Student's t -distribution with k degrees of freedom.

- (a) Compute the third cumulant of a t -distributed random variable.
- (b) Explain why higher-order cumulants capture heavy-tailed behavior.
- (c) Compare the behavior of the t_3 and $N(0, 1)$ distributions using cumulants.

Exercise 3.5.36 (Dependence in Gaussian Mixture Models). Suppose $X \sim N_2(0, I_2)$ with correlation ρ . Now, let $Z \sim \text{Bernoulli}(0.5)$, and define

$$Y = (1 - Z)X + ZW,$$

where $W \sim N_2(0, I_2)$ is independent of X .


- (a) Show that Y follows a Gaussian mixture model.
- (b) Compute the correlation between Y_1 and Y_2 .
- (c) Explain why the dependence structure is different from that of a Gaussian model.

4

Principal Component Analysis

The primary challenge in multivariate statistics is managing a large number of variables, potentially in the millions. One approach is to model these variables directly using high-dimensional statistical techniques. Alternatively, dimensionality reduction methods can be applied to derive a smaller set of variables that capture the most significant relationships in the data. These derived variables often serve as effective substitutes for the original data.

This chapter provides a comprehensive introduction to Principal Component Analysis (PCA). While the lecture slides illustrate simple examples of how PCA can project a multivariate dataset onto two dimensions for visualization and insight, the following sections carefully develop the underlying theory. For a nice motivating example see Figure 4.1 taken from the paper¹.



¹ Novembre et al, Genes mirror geography within Europe, Nature, 2009

Figure 4.1: Principal Component Analysis (PCA) applied to genetic data from individuals across Europe, illustrating how the first two principal components mirror the geographical distribution of populations. Each point represents an individual, color-coded by their country of origin, with clusters forming a map-like structure of Europe.

4.1 Principal components

We start with the following problem. Given a random m -vector $X \sim (\mu, \Sigma)$ find a scalar random variable $Z = u^\top X$ with the largest variance. This is not well defined if we allow u to be any vector in \mathbb{R}^m . Thus, we constrain to $u \in \mathbb{R}^m$ such that $\|u\| = 1$. We say that u is a unit vector. Unit vectors define directions.

Recall $\text{var}(u^\top X) = u^\top \Sigma u$. Thus we solve a constrained problem

$$\text{maximize } u^\top \Sigma u \text{ subject to } u^\top u = 1.$$

Since we optimize a continuous function over a compact set, this optimum must be attained. We encode the constraint $\|u\| = 1$ as $\|u\|^2 = u^\top u = 1$ to avoid dealing with square roots.

The Lagrangian is

$$L = u^\top \Sigma u - \lambda(u^\top u - 1)$$

and every optimum of the original problem must be a stationary point of the Lagrangian. We get

$$\nabla_u L = 2\Sigma u - 2\lambda u.$$

This shows that every stationary point of L must correspond to an eigenvector of Σ with λ being the corresponding eigenvalue. If u is such an eigenvector then the value of the original function is

$$u^\top \Sigma u = \lambda u^\top u = \lambda.$$

Thus, the stationary point of L that gives the maximal value of the original function is the eigenvector corresponding to the maximal eigenvalue of Σ (which will be always strictly positive).

Suppose that we have now identified u_1 as the first principal direction and λ_1 as the corresponding variance.

We can now look for other linear combinations uncorrelated with $u_1^\top X$ with maximum variance. The uncorrelation condition is

$$0 = \text{cov}(u_1^\top X, u^\top X) = u_1^\top \Sigma u = \lambda_1 u_1^\top u.$$

Since $\lambda_1 > 0$, equivalently $u_1^\top u = 0$. This amounts to solving

$$\text{maximize } u^\top \Sigma u \text{ subject to } u^\top u = 1 \text{ and } u_1^\top u = 0.$$

Here the Lagrangian is

$$L = u^\top \Sigma u - \lambda(u^\top u - 1) - \nu u_1^\top u.$$

The stationary points satisfy $\Sigma u - \lambda u = \frac{\nu}{2} u_1$. If $\nu \neq 0$ we get a contradiction with $u^\top u_1 = 0$ ² and so $\nu = 0$. But in this case, we

We write $X \sim (\mu, \Sigma)$ to denote that $\mathbb{E}X = \mu$ and $\text{var}(X) = \Sigma$.

We derived these gradient formulas earlier.

² Multiplying by u_1^\top from the left, we get $u_1^\top \Sigma u - \lambda u_1^\top u = \nu/2$, or equivalently, $(\lambda_1 - \lambda)u_1^\top u = \nu/2$.

easily show that u_2 has to be the eigenvector of Σ corresponding to the second largest eigenvalue λ_2 . And we can continue in this way obtaining m random variables $Z_1 = u_1^\top X, \dots, Z_m = u_m^\top X$. By construction, the variables Z_i are uncorrelated and they are ordered by the size of their variance.

Recall the spectral theorem in Section 1.1.5, which states that we can decompose $\Sigma = U\Lambda U^\top$, where U is an orthogonal matrix and Λ is diagonal with positive diagonal entries. The *columns* of U are eigenvectors and the diagonal entries of Λ are the corresponding eigenvalues. If the entries of Λ are in decreasing order, the columns of U are the subsequent principal directions and the PCA transformation $Z = (Z_1, \dots, Z_m)$ satisfies

$$Z = U^\top X.$$

4.2 Principal Components for Observed Data

Suppose we have samples $x_1, \dots, x_n \in \mathbb{R}^m$. Write them as rows in a data matrix $X \in \mathbb{R}^{n \times m}$. We then standardize the data so that each column has mean zero and variance one³. Denote the normalized data matrix by \tilde{X} .

We can then compute the sample **correlation** matrix $S_n = \frac{1}{n}\tilde{X}^\top \tilde{X}$ and proceed exactly as above, replacing Σ with S_n . We again write the spectral decomposition $S_n = U\Lambda U^\top$, where U is the matrix of loadings.

We can now use it to project the data to a d -dimensional subspace with $d < m$. Let u_1, \dots, u_d be the first d principal directions⁴. The **scores** are the images of the centered data points $\tilde{x}_1, \dots, \tilde{x}_n$ in the coordinate system defined by the d loadings. These are the points, y_1, \dots, y_n , given by

$$y_j = (u_1^\top \tilde{x}_j, \dots, u_d^\top \tilde{x}_j).$$

In other words, $Y \in \mathbb{R}^{n \times d}$ with rows y_1, \dots, y_n is obtained by taking the first d columns of $\tilde{X}U$.

4.3 Simple Graphics for PCA

4.3.1 Scree Plot and Fraction of Variance Explained

Suppose that U is the matrix of principal directions, that is, $S_n = U\Lambda U^\top$. Then $\text{tr}(S_n) = \lambda_1 + \dots + \lambda_m$. It follows that

$$S_{11} + \dots + S_{mm} = \lambda_1 + \dots + \lambda_m$$

and so the total variance is preserved by PCA.

³ In practice, we normalize data so that the diagonal entries of S_n are all equal to 1. To see why this is important, perform a simple experiment: Take any preferred dataset and multiply one of the observed variables by a large number (e.g., by 1000 when changing units from kilograms to grams). See how this affects the results.

⁴ These are the first d columns of U , assuming the diagonal entries of Λ are in the decreasing order.

Recall from linear algebra that if u_1, \dots, u_d is an orthogonal basis of a linear subspace $V \subseteq \mathbb{R}^m$ and $v \in V$, then $v = (v^\top u_1)u_1 + \dots + (v^\top u_m)u_m$

Note that $S_{ii} = \text{Var}(e_i^\top X)$ and so $\lambda_1 \geq \max_i(S_{ii})$.

Often, most of the variance is explained by a small number of principal components. The amount of the total variance explained by the first k principal components is

$$\frac{\lambda_1 + \dots + \lambda_k}{\lambda_1 + \dots + \lambda_m} \in [0, 1]$$

and we can choose k , for example, so that this ratio is at least 0.9.

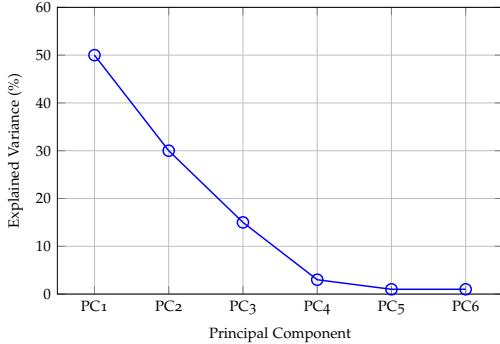


Figure 4.2: An example of the Scree plot. Here the first three components capture most of the variance. If $\Sigma_{ii} = 1$ for all i , then $\frac{\lambda_i}{m} \times 100\%$ is the percentage of the variance explained by the i -th component.

In R, simply run `princomp(data, cor=TRUE)`.

4.3.2 Biplot

In the slides, we presented a number of biplots. What is this? A PCA biplot is a graphical representation that shows both the observations (data points) and the principal components' directions (loadings) in a single plot. It provides insight into how the original variables contribute to the principal components and how the observations relate to each other in the space of the principal components.

The idea is simple. We compute the first two principal directions $\mathbf{u}_1, \mathbf{u}_2 \in \mathbb{R}^m$ and each observation in the dataset is projected onto the principal components. These projections were called scores in Section 4.2. They are computed as $(\mathbf{u}_1^\top \mathbf{x}_i, \mathbf{u}_2^\top \mathbf{x}_i)$ for $i = 1, \dots, n$. This is our 2-dimensional dataset.

The loadings represent the contribution of each original variable to the principal components. Loadings are the coefficients of the linear combinations that define each principal component. Concretely, these are obtained from the first two rows of the loading matrix U . In a biplot, these are typically shown as arrows or vectors, indicating the direction and strength of each variable's influence on the principal components.

4.4 PCA and Affine Approximating Subspaces

As before, let $\mathbf{X} \in \mathbb{R}^{n \times m}$ be the data matrix representing n points \mathbf{x}_i in \mathbb{R}^m . We again assume the data points in \mathbf{X} are centered and

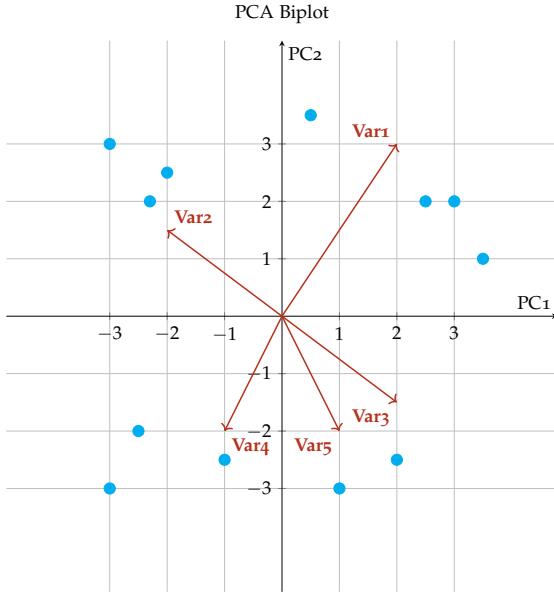


Figure 4.3: An example of a biplot with scores as blue dots and loadings as brick-red arrows.

normalized. In this case $S_n = \frac{1}{n}\mathbf{X}^\top\mathbf{X}$ is a correlation matrix. Recall that $S_n = U\Lambda U^\top$ is the spectral decomposition of S_n and so the SVD of \mathbf{X} must be of the form $\mathbf{X} = VDU^\top$ for some orthogonal V and the diagonal entries of D are the singular values of \mathbf{X} . In this way $\frac{1}{n}\mathbf{X}^\top\mathbf{X} = \frac{1}{n}UD^\top DU$ and we get the relation $\Lambda = \frac{1}{n}DD^\top$.

Every d -dimensional affine subspace of \mathbb{R}^m can be written as $\mu + \mathcal{L}$, where \mathcal{L} is a d -dimensional linear subspace of \mathbb{R}^m and $\mu \in \mathbb{R}^m$. Without loss we can assume $\mu \in \mathcal{L}^\perp$ ⁵.

Let $\{w_1, \dots, w_d\}$ be an orthonormal basis of \mathcal{L} . Then, each $\mathbf{x} \in \mu + \mathcal{L}$ is of the form

$$\mathbf{x} = \mu + Wz,$$

where W is the matrix with columns w_1, \dots, w_d and $z \in \mathbb{R}^d$. By orthogonality, we have $W^\top W = I_d$. Since $\mu \in \mathcal{L}^\perp$, $W^\top \mu = \mathbf{0}$.

To compute the distance of \mathbf{x}_i to $\mu + \mathcal{L}$, we find the orthogonal projection of \mathbf{x}_i onto this subspace. Analytically, we solve

$$\text{minimize } \|\mathbf{x}_i - (\mu + Wz_i)\|^2 \text{ subject to } z_i \in \mathbb{R}^d.$$

We easily check that the optimal z_i satisfies

$$\hat{z}_i = W^\top(\mathbf{x}_i - \mu) = W^\top \mathbf{x}_i$$

and so the projection of \mathbf{x}_i onto the affine subspace $\mu + \mathcal{L}$ is

$$\hat{\mathbf{x}}_i := \mu + WW^\top \mathbf{x}_i.$$

The goal is to find an affine subspace of dimension d that minimizes the sum of squared distances between the points \mathbf{x}_i and their

⁵Indeed, we can uniquely decompose μ as $u + v$ with $u \in \mathcal{L}$ and $v \in \mathcal{L}^\perp$ and so $\mu + \mathcal{L} = v + \mathcal{L}$.

projections onto this subspace. That is, we aim to solve the following optimization problem:

$$\text{minimize } \sum_{i=1}^n \|\hat{\mathbf{x}}_i - \mathbf{x}_i\|^2 = \sum_{i=1}^n \|\mu - (I_m - WW^\top)\mathbf{x}_i\|^2$$

subject to $\mu \in \mathbb{R}^m$, $W \in \mathbb{R}^{m \times d}$, $W^\top W = I_d$. Optimizing first with respect to μ we get

$$\hat{\mu} = (I_m - WW^\top)\bar{\mathbf{x}}_n = 0,$$

where we used the fact that the data are centred. Thus, equivalently we have

$$\text{minimize } \sum_{i=1}^n \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2 \text{ subject to } W^\top W = I_d, \quad (4.1)$$

where $\hat{\mathbf{x}}_i = WW^\top \mathbf{x}_i$.

Let $\|A\|_F = \sqrt{\text{tr}(AA^\top)}$ denote the Frobenius norm of A ⁶. We have

$$\sum_{i=1}^n \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2 = \|\mathbf{X} - \hat{\mathbf{X}}\|_F^2 = \|\mathbf{X} - \mathbf{X}WW^\top\|_F^2.$$

Note that $\mathbf{X}WW^\top$ has rank $\leq d$. Thus, the following result is relevant.

Theorem 4.4.1 (Eckart-Young Theorem). *Given $\mathbf{X} \in \mathbb{R}^{n \times m}$ with SVD $\mathbf{X} = VDU^\top$ and the problem: minimize $\|\mathbf{X} - M\|_F$ subject to $\text{rank}(M) \leq d$, the solution is given by $\hat{M} = VD_dU^\top$, where D_d is obtained from D by zeroing out all but the first d diagonal entries.*

For a proof and a more detailed discussion, see Section I.9 in ⁷.

Using this result, it is enough to show that there is a W with $W^\top W = I_d$ such that $\hat{M} = \mathbf{X}WW^\top$ or, in other words,

$$VD_dU^\top = VDU^\top WW^\top. \quad (4.2)$$

Taking $W = U_d$ as the first d columns of U we easily verify (4.2). Indeed,

$$VDU^\top U_d U_d^\top = VD \begin{bmatrix} I_d \\ \mathbf{0} \end{bmatrix} U_d^\top = VD \begin{bmatrix} U_d^\top \\ \mathbf{0} \end{bmatrix} = VD_d \begin{bmatrix} U_d^\top \\ \mathbf{0} \end{bmatrix} = VD_d U^\top.$$

4.5 Principal Component Regression (PCR) and Probabilistic PCA (PPCA)

4.5.1 Principal Component Regression

In some cases, PCA can be used as a preprocessing step in regression to reduce the dimensionality of the data and remove multicollinearity. By using the first few principal components as predictors in the

⁶ The space $\mathbb{R}^{n \times m}$ admits a natural inner product $\langle A, B \rangle = \text{tr}(AB^\top)$. The Frobenius norm is the induced norm: $\|A\|_F = \sqrt{\langle A, A \rangle} = \sqrt{\sum_{i,j} A_{ij}^2}$

← Exercise 4.8.19

⁷ Gilbert Strang. *Linear Algebra and Learning from Data*. Wellesley-Cambridge Press, Wellesley, MA, 2019. ISBN 9780692196380

regression model, we can obtain more stable estimates of the regression coefficients and improve the interpretability of the model.

For example, suppose we have a dataset with m variables and we perform PCA to extract the first d principal components. We can then fit a regression model using these d components as predictors:

$$\hat{Y} = \alpha + \beta_1 Z_1 + \beta_2 Z_2 + \cdots + \beta_d Z_d,$$

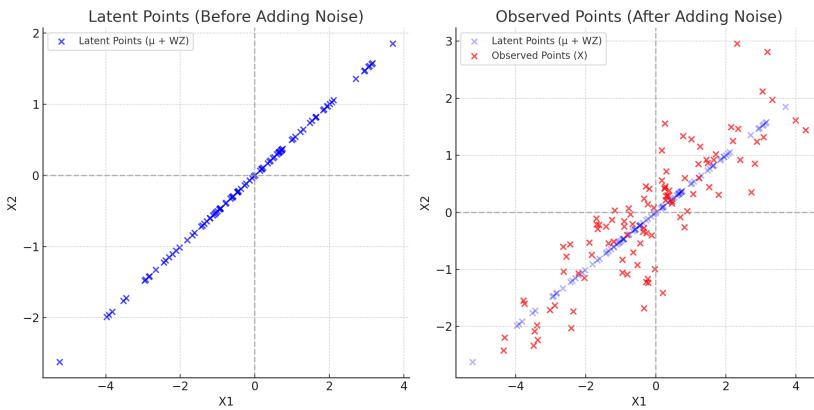
where Z_1, Z_2, \dots, Z_d are the first d principal components. This is known as principal component regression (PCR), and it is often used when the number of predictors is large relative to the number of observations.

4.5.2 Probabilistic PCA

There is a related generative model in which we model a random vector X through the linear relation

$$X = \mu + WZ + \varepsilon, \quad Z \sim N_r(0, I_r), \quad \varepsilon \sim N_m(0, \sigma^2 I_m), \quad Z \perp\!\!\!\perp \varepsilon.$$

To have a concrete example, suppose $r = 1, m = 2, \mu = (0, 0)$, and $W = (1, \frac{1}{2})$. Sampling Z from $N(0, 1)$ produces blue points in Figure 4.5.2. Additionally, sampling $\varepsilon \sim N_2(0, I_2)$ and adding it to the corresponding blue points gives the red points, which are the samples from the underlying PPCA model.



It is not hard to see that $X|Z$ is Gaussian and so X is also marginally Gaussian with mean μ and covariance $\sigma^2 I_m + WW^\top$. It turns out that this model admits a closed-form formula for the MLE, which mimics the PCA solution.

Theorem 4.5.1. *With the notation as above, let $S_n = U\Lambda U^\top$ with the diagonal entries of Λ satisfying $\lambda_1 \geq \cdots \geq \lambda_m$. The MLE in the PPCA*

model satisfies

$$\begin{aligned}\hat{\mu} &= \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i, \\ \hat{\sigma}^2 &= \frac{1}{m-r} \sum_{i=r+1}^m \lambda_i, \\ \hat{W} &= U_r (\Lambda_r - \hat{\sigma}^2 I_r)^{\frac{1}{2}} Q,\end{aligned}$$

for any orthogonal matrix Q .

Note that the column span of \hat{W} is the same as the column span of U_d and so the MLE corresponds to PCA.

4.6 PCA and Matrix Completion*

Suppose that now $\mathbf{X} \in \mathbb{R}^{n \times m}$ contains some missing data denoted as `NA`. We would like to be able to:

1. Estimate the missing entries.
2. Compute the PCA in the presence of missing data.

We saw in Theorem 4.4.1 that SVD(r) solves

$$\text{minimize} \|\mathbf{X} - M\|_F^2 \quad \text{subject to } \text{rank}(M) \leq r.$$

Let $\Omega \in \{0,1\}^{n \times m}$ be such that $\Omega_{ij} = 0$ if the corresponding entry of \mathbf{X} is `NA` and $\Omega_{ij} = 1$ otherwise. Define $\Omega \circ \mathbf{X}$ as the Hadamard (entrywise) product:

$$(\Omega \circ \mathbf{X})_{ij} = \begin{cases} 0 & \text{if } \mathbf{X}_{ij} = \text{NA}, \\ \mathbf{X}_{ij} & \text{otherwise.} \end{cases}$$

The optimization problem for estimating missing entries and performing PCA is:

$$\text{minimize} \|\Omega \circ (\mathbf{X} - M)\|_F^2 \quad \text{subject to } \text{rank}(M) \leq r.$$

4.6.1 Hardimpute Algorithm

The `hardimpute` algorithm iteratively solves the matrix completion problem:

1. Initialize $M^0 = \mathbf{0}_{n \times m}$.
2. Iterate until convergence:
 - $\mathbf{X}^* = \Omega \circ \mathbf{X} + (\mathbf{1}_{n \times m} - \Omega) \circ M^{j-1}$.
 - $M^j = \arg \min_{M: \text{rank}(M) \leq r} \|\mathbf{X}^* - M\|_F^2$.

4.7 Appendix: Covariance matrix estimation in high dimensions

In PCA we use the sample covariance matrix S_n as a natural estimator of the population covariance matrix Σ . We note however that this is only a good estimator when n is much larger than m . Otherwise, we may need to use an alternative estimator. In this section we very briefly overview some of the existing techniques.

4.7.1 Motivating alternative estimators

As noted in Section 1.3.3, the sample covariance matrix is a natural estimator due to its small bias, which diminishes rapidly as n increases. Furthermore, it can be shown to be asymptotically Gaussian under the assumption that the underlying dimension m remains fixed. However, this asymptotic analysis becomes inadequate in scenarios where n is not significantly larger than m , as the standard assumptions no longer hold.

For concreteness, suppose that we measure quality of this estimator by analyzing the distance $\|S_n - \Sigma\|$ in the operator norm

$$\|S_n - \Sigma\| := \max_{\|\mathbf{u}\|=1} \|(S_n - \Sigma)\mathbf{u}\|.$$

Under mild moment conditions, an argument based on the classical law of large numbers (and continuous mapping theorem) can be used to show that the difference $\|S_n - \Sigma\|$ converges to zero in probability as $n \rightarrow \infty$. Consequently, S_n is a consistent estimator of the population covariance Σ in the classical setting.

What happens when we also allow m to tend to infinity? To make this question a bit more formal, consider sequences of problems (S_n, Σ) indexed by the pair (n, m) , and suppose that we allow both n and m grow with their ratio remaining fixed, say $m/n = \gamma \in (0, 1)$. In Figure 4.4 we plot the results of simulations for the case when $\Sigma = I_m$ with each sample $\mathbf{x}^{(i)} \sim N(\mathbf{0}, I_m)$.

Using n samples we compute S_n and compute its eigenvalues. They should be all nonnegative and actually positive if $n \geq p$. We then construct a histogram for these eigenvalues. Since in our simulations n is large (we take $n = 1000$), classical theory suggests, all eigenvalues should be close to 1 as S_n should be close to I_m . But this is not what we observe. It actually seems that the eigenvalues follow a distribution that is well spread out; cf. Figure 4.4.

This situation has been well studied in random matrix theory. The underlying eigenvalue distribution is called the Marchenko-Pastur law. Define $\lambda_{\min} = (1 - \sqrt{\gamma})^2$ and $\lambda_{\max} = (1 + \sqrt{\gamma})^2$ then the law

admits density

$$f_{MP}(\lambda) = \frac{1}{2\pi\gamma\lambda} \sqrt{(\lambda_{\max} - \lambda)(\lambda - \lambda_{\min})} \quad \text{for } \lambda_{\min} \leq \lambda \leq \lambda_{\max}.$$

If you want to play around with this, this is R code that produces one of the plots in Figure 4.4.

```
# Load necessary libraries
library(MASS)
library(ggplot2)

# Simulation parameters
n <- 1000          # Number of samples
p <- 200           # Dimensionality
n_sim <- 100       # Number of simulations

# Function to compute eigenvalues of the sample covariance matrix
compute_eigenvalues <- function(n, p) {
  data <- mvrnorm(n = n, mu = rep(0, p), Sigma = diag(p))      sample_cov <-
  → cov(data) # Compute sample covariance matrix
  eigen(sample_cov)$values # Return eigenvalues
}

# Perform simulations
set.seed(42) # For reproducibility
eigenvalues_list <- replicate(n_sim, compute_eigenvalues(n, p), simplify =
→ FALSE)
all_eigenvalues <- unlist(eigenvalues_list)

# Parameters for the Marchenko-Pastur density
gamma <- p / n # Aspect ratio
lambda_min <- (1 - sqrt(gamma))^2 # Lower bound of the MP support
lambda_max <- (1 + sqrt(gamma))^2 # Upper bound of the MP support

# Marchenko-Pastur density function
mp_density <- function(lambda, gamma) {
  if (lambda >= lambda_min && lambda <= lambda_max) {
    return((1 / (2 * pi * gamma)) * sqrt((lambda_max - lambda) * (lambda -
    → lambda_min)) / lambda)
  } else {return(0)}
}

# Generate MP density values
lambda_seq <- seq(lambda_min, lambda_max, length.out = 500)
mp_values <- sapply(lambda_seq, mp_density, gamma = gamma)

# Plot normalized histogram and MP density
ggplot(data = data.frame(eigenvalue = all_eigenvalues), aes(x =
→ eigenvalue)) +
  geom_histogram(aes(y = ..density..), bins = 50, color = "black", fill =
  → "skyblue", alpha = 0.5) +
  geom_line(data = data.frame(lambda = lambda_seq, density = mp_values),
  → aes(x = lambda, y = density), color = "deepskyblue4", size = 1.2) +
  labs(title =
  → "Comparison of Eigenvalue Histogram and Marchenko-Pastur Density",
  x = "Eigenvalue", y = "Density") + theme_minimal()
```




Figure 4.4: The histogram of eigenvalues and the corresponding Marchenko-Pastur law for $n = 1000$ and $p = 200$ (equiv. $\gamma = 0.2$) on the top, $p = 500$ (equiv. $\gamma = 0.5$) in the bottom.

The Marchenko-Pastur law is an asymptotic result, albeit with a non-classical flavor, as it allows both the sample size n and the

dimension m to diverge simultaneously. Similar to classical asymptotics, it provides an approximation for the distribution of eigenvalues even when n and m are finite. When m is not negligible relative to n , this approximation significantly outperforms the classical one.

4.7.2 Review of alternative estimators

If the covariance matrix has some special structure, incorporating this structure in the estimation procedures can greatly help in high-dimensional scenarios. Regularization introduces constraints or penalties to improve the stability and accuracy of covariance estimates. Key approaches include:

Linear shrinkage. The linear shrinkage estimator is given by:

$$\hat{\Sigma}_{\text{ridge}} = (1 - \lambda)S_n + \lambda I_m,$$

where $\lambda \in [0, 1]$ is the shrinkage parameter. Choosing λ appropriately (e.g., via cross-validation) balances bias and variance, ensuring better performance in high-dimensional settings.

The intuition behind shrinking toward the identity matrix arises from several considerations. First, as we already noted, in high-dimensional settings where $p \approx n$ or $p > n$, the sample covariance matrix S_n is often unstable, highly variable, or even singular (non-invertible). Adding λI_m stabilizes the estimate by ensuring that $\hat{\Sigma}_{\text{ridge}}$ is well-conditioned and invertible. Second, the shrinkage estimator introduces bias by pulling S_n toward I_m , but it reduces variance, which can dominate in high dimensions. A well-chosen λ balances this trade-off, leading to improved estimation accuracy for the covariance matrix.

Graphical Lasso (Sparse Precision Matrix Estimation). For many high-dimensional datasets, it is reasonable to assume that variables are conditionally independent given others. In Section 2.3.2 we showed that, at least in the Gaussian case, certain type of conditional independences correspond to zeros in the inverse covariance matrix. In situations, where $K = \Sigma^{-1}$ is indeed sparse we can leverage this by using the graphical lasso:

$$\hat{K} = \arg \min_{\Theta \succ 0} \{ \text{tr}(S_n K) - \log \det(K) + \lambda \|K\|_1 \},$$

where penalizing $\|K\|_1 = \sum_{i \neq j} |K_{ij}|$ with penalty parameter $\lambda > 0$ promotes sparsity. The resulting estimate \hat{K} can be inverted to approximate Σ . We discuss this estimator in much more detail in Chapter 8.

Factor Models. In many applications, the data can be modeled as having a low-rank structure plus noise. A covariance matrix in such settings can be decomposed as:

$$\Sigma = LL^\top + \Psi,$$

where $L \in \mathbb{R}^{p \times r}$ ($r \ll p$) is a low-rank factor matrix, and Ψ is a diagonal matrix capturing idiosyncratic variances. Estimation could proceed by minimizing reconstruction error:

$$\min_{L, \Psi} \|S_n - LL^\top - \Psi\|_F^2.$$

We discuss the closely related Factor Analysis model in Chapter 7.

Thresholding-Based Methods. A simple and effective approach for high-dimensional covariance estimation is to threshold small entries in S_n , which helps reduce noise and enforce sparsity. For a given threshold $\tau > 0$, the thresholded covariance estimate is:

$$\hat{\Sigma}_{\text{thresh}} = [(S_n)_{ij} \cdot \mathbb{I}(|(S_n)_{ij}| > \tau)]_{i,j}.$$

Thresholding methods rely on the assumption that most entries in the covariance matrix are small or zero, which is common in many applications.

Banding and Tapering for Structured Covariance Matrices. For structured data, such as time series or spatial data, covariance matrices often exhibit banded or tapering structures, where correlations decay with distance. It is then natural to impose this structure in the estimation process. For an illustration what we mean here, consider the problem in Exercise .

Tyler's scatter estimator. This estimator came from the robust statistics and it assumes that the data come from an elliptical distribution centred at the origin. The objective is to estimate the scale matrix Σ .⁸

Suppose $X \sim E(0, \Sigma)$ and set $Y = \Sigma^{-1/2}X$. Then Y has a spherical distribution, so that $Y/\|Y\|$ has a uniform distribution on the unit sphere in \mathbb{R}^m . It is easy to show that the covariance matrix of such a vector is $\frac{1}{m}I_m$. Hence

$$\mathbb{E}\left[\frac{1}{\|Y\|^2}YY^\top\right] = \mathbb{E}\left[\frac{1}{X^\top\Sigma^{-1}X}\Sigma^{-1/2}XX^\top\Sigma^{-1/2}\right] = \frac{1}{m}I_m,$$

or equivalently,

$$\mathbb{E}\left[\frac{1}{X^\top\Sigma^{-1}X}XX^\top\right] = \frac{1}{m}\Sigma. \quad (4.3)$$

Now suppose we observe a random sample $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$ with the same distribution as X . The sample version of (4.3) is given by⁹

⁸ Recall from Section 3.1.5 that the covariance matrix in the elliptical distribution is a multiple of the scale matrix Σ , $\text{var}(X) = c\Sigma$ for some $c > 0$.

← Exercise 3.5.3

We could use our Mahalanobis distance notation: $\|X\|_\Sigma^2 = X^\top\Sigma^{-1}X$.

⁹ To estimate $\mathbb{E}f(X)$ is it natural to use $\frac{1}{n} \sum_i f(\mathbf{x}^{(i)})$.

$$\sum_i \frac{1}{\mathbf{x}^{(i)^\top} \Sigma^{-1} \mathbf{x}^{(i)}} \mathbf{x}^{(i)} \mathbf{x}^{(i)^\top} = \frac{n}{m} \Sigma. \quad (4.4)$$

Under mild conditions on the data (e.g. $n \geq m$ and no linear dependence in the data), it turns out that (4.4) has a unique solution $\hat{\Sigma}$, say, defined up to a scalar multiple. The estimator is typically computed using an iterative fixed-point algorithm:

1. Start with an initial positive definite matrix $\hat{\Sigma}^{(0)}$, such as the sample covariance matrix.
2. Iterate the update rule:

$$\hat{\Sigma}^{(k+1)} = \frac{m}{n} \sum_{i=1}^n \frac{\mathbf{x}^{(i)} \mathbf{x}^{(i)^\top}}{\mathbf{x}^{(i)^\top} (\hat{\Sigma}^{(k)})^{-1} \mathbf{x}^{(i)}}.$$

4.8 Exercises

Exercise 4.8.1. Consider a dataset with n observations and m variables. Explain why PCA finds an orthonormal basis for the data space. How does this basis relate to the directions of maximum variance?

Exercise 4.8.2. Suppose you apply PCA to a dataset and find that the first two principal components explain 95% of the variance. What does this tell you about the structure of the data?

Exercise 4.8.3. Let $\mathbf{X} \sim N_m(\mathbf{0}, \Sigma)$. Show that PCA transformations preserve normality, i.e., the principal components also follow a normal distribution. Show that in this case the principal components are independent.

Exercise 4.8.4. Explain how the PCA can be obtained from the SVD of the underlying data matrix \mathbf{X} . For simplicity assume that the data points are centered.

Exercise 4.8.5. Why is it necessary to standardize the data (zero mean, unit variance) before applying PCA when variables are measured on different scales?

Exercise 4.8.6. Show that the principal components are uncorrelated by computing the covariance matrix of the transformed variables.

Exercise 4.8.7. Suppose we perform PCA on a dataset where one variable is a linear combination of the others. What will happen to the eigenvalues, and how will this affect the principal components?

Exercise 4.8.8. Let $\mathbf{X} \sim N(\mathbf{0}, I_m)$ and consider the principal component transformation $\mathbf{Z} = \mathbf{U}^\top \mathbf{X}$, where \mathbf{U} is the eigenvector matrix of I_m . Show that \mathbf{Z} has the same distribution as \mathbf{X} .

Exercise 4.8.9. Given the eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_m$, define the cumulative proportion of variance explained and determine how many principal components are needed to retain at least 90% of the variance.

Exercise 4.8.10. Show that PCA minimizes the sum of squared distances between the original data points and their projections onto the first k principal components.

Exercise 4.8.11. Explain the relationship between Singular Value Decomposition (SVD) and PCA. Why does the SVD of the data matrix yield the same principal components as the eigendecomposition of the covariance matrix?

Exercise 4.8.12. Suppose we obtain the first k principal components of a dataset. How can these be used for dimensionality reduction? What are the consequences of choosing a very small or very large k ?

Exercise 4.8.13. Perform PCA on the following dataset:

$$X = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix}$$

Compute the covariance matrix, find the eigenvalues and eigenvectors, and project the data onto the first principal component.

Exercise 4.8.14. If the first principal component explains 80% of the total variance, what does this imply about the structure of the data? Provide an example where most of the variance is captured by one component.

Exercise 4.8.15. Why do the eigenvalues of the covariance matrix correspond to the variances of the principal components? Derive this result formally.

Exercise 4.8.16. In high-dimensional settings, why might PCA be used for regularization? Discuss how PCA can help in reducing overfitting when $m \gg n$.

Exercise 4.8.17. Let $X \sim N(\mu, \Sigma)$ and consider the transformation $Z = U^\top(X - \mu)$, where U is the eigenvector matrix of Σ . Show that the components of Z are uncorrelated.

Exercise 4.8.18. Discuss the interpretation of PCA in terms of data reconstruction: If only the first k principal components are used, how does this affect the approximation of the original data? Derive an expression for the reconstruction error.

Exercise 4.8.19. Show that the matrix AB with $A \in \mathbb{R}^{k \times l}$ and $B \in \mathbb{R}^{l \times m}$ has rank $\leq l$.

Exercise 4.8.20. Consider a sample $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$ from the three-dimensional Gaussian distribution $N_3(0, \Sigma)$, where Σ is assumed to take the form

$$\Sigma = \begin{bmatrix} a & b & b \\ b & a & b \\ b & b & a \end{bmatrix} = (a - b)I_3 + b\mathbf{1}_3\mathbf{1}_3^\top.$$

Propose an estimator for the covariance matrix and check its performance using simulations. Bonus question: The Sylvester's determinant theorem states that for any two matrices $A, B \in \mathbb{R}^{m \times n}$ it holds that $\det(I_m + AB^\top) = \det(I_n + B^\top A)$. Use this fact to find conditions on a, b that assure that Σ is positive definite.

5

Some Other Dimension Reduction Methods

Dimensionality reduction techniques beyond PCA are often necessary when the underlying data structure is nonlinear or when variance-based methods like PCA do not capture enough of the essential relationships in the data. The goal of this chapter is to show what are the possible ways to handle non-linearity with preserving some of the computational advantages of linear methods. We do it by introducing two state-of-the-art techniques: Uniform Manifold Approximation and Projection (UMAP). This part of the lecture will be presented on slides and will be relatively high-level.

5.1 Multidimensional Scaling (MDS)

Multidimensional Scaling (MDS) is a statistical technique used to represent data points in a lower-dimensional space while preserving their pairwise distances as faithfully as possible.

5.1.1 Problem Setup

Given n objects, let $\Delta = [\delta_{ij}]$ be an $n \times n$ **dissimilarity matrix**. A dissimilarity matrix is simply a symmetric matrix with zero diagonal entries and nonnegative off-diagonal entries. Its entry δ_{ij} represents the dissimilarity between objects i and j . In our setting, the objects will be the data points $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^m$ and δ_{ij} will be their mutual distances, but this setting is much more general.

The goal of MDS is to find a configuration of points $\{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\}$ in a d -dimensional space ($d \ll m$) such that:

$$d_{ij} = \|\mathbf{y}_i - \mathbf{y}_j\| \approx \delta_{ij}.$$

5.1.2 Classical MDS Solution

Classical MDS assumes that for some $m \in \mathbb{N}$ and $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^m$ such that the dissimilarities δ_{ij} are Euclidean distances, $\delta_{ij} =$

$\|\mathbf{x}_i - \mathbf{x}_j\|$. Let $\Delta \odot \Delta = [\delta_{ij}^2]$ be the matrix of squared dissimilarities (Hadamard product of Δ with itself). Note that the entries of $\mathbf{X}\mathbf{X}^\top$ satisfy $(\mathbf{X}\mathbf{X}^\top)_{i,j} = \mathbf{x}_i^\top \mathbf{x}_j$ and

$$\delta_{ij}^2 = \|\mathbf{x}_i - \mathbf{x}_j\|^2 = (\mathbf{X}\mathbf{X}^\top)_{i,i} + (\mathbf{X}\mathbf{X}^\top)_{j,j} - 2(\mathbf{X}\mathbf{X}^\top)_{i,j}.$$

Equivalently, using matrix algebra,

$$\Delta \odot \Delta = \text{diag}(\mathbf{X}\mathbf{X}^\top)\mathbf{1}\mathbf{1}^\top + \mathbf{1}\mathbf{1}^\top \text{diag}(\mathbf{X}\mathbf{X}^\top) - 2\mathbf{X}\mathbf{X}^\top.$$

Recall that the centering matrix $H = I_n - \frac{1}{n}\mathbf{1}\mathbf{1}^\top$ satisfies $H\mathbf{1} = \mathbf{0}$ and so

$$B := -\frac{1}{2}H(\Delta \odot \Delta)H = H\mathbf{X}(H\mathbf{X})^\top = \tilde{\mathbf{X}}\tilde{\mathbf{X}}^\top.$$

A natural way to get a lower-dimensional embedding of the centered data $\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_n$ in a d -dimensional space that approximately preserves the mutual distances is by approximating the matrix B with a rank d matrix $\mathbf{Y}\mathbf{Y}^\top$ with $\mathbf{Y} \in \mathbb{R}^{n \times d}$. We already saw in Theorem 4.4.1 how to solve this problem. In this way we obtain $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^d$ such that $\sum_{i,j} (\langle \mathbf{x}_i, \mathbf{x}_j \rangle - \langle \mathbf{y}_i, \mathbf{y}_j \rangle)^2$ is minimized

As suggested by Theorem 4.4.1, we first perform eigenvalue decomposition: $B = V\Lambda V^\top$, or equivalently, singular value decomposition of $H\mathbf{X}$. Then, we select the top d eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_d$ (assumed positive) and their corresponding eigenvectors $\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_d$. The $n \times d$ configuration matrix is:

$$\mathbf{Y} = U_d \Lambda_d^{1/2},$$

where U_d contains the top d eigenvectors, and $\Lambda_d^{1/2}$ is the diagonal matrix of the square roots of the top d eigenvalues.

Below is a simple example demonstrating classical MDS using simulated data:

```
# Generate sample data and apply classical MDS
set.seed(42)
library(MASS)
library(ggplot2)

# Generate sample data
X <- mvrnorm(100, mu = rep(0, 5), Sigma = diag(5))

# Compute distance matrix and apply MDS
dist_matrix <- dist(X)
mds_result <- cmdscale(dist_matrix, k = 2)

# Plot the results
data <- data.frame(X1 = mds_result[, 1], X2 = mds_result[, 2])
ggplot(data, aes(x = X1, y = X2)) +
  geom_point(color = "blue", alpha = 0.7) +
  ggtitle("Classical MDS Projection") +
  xlab("MDS Dimension 1") +
  ylab("MDS Dimension 2") +
  theme_minimal()
```

Do not confuse $\mathbf{X}\mathbf{X}^\top$ with $\mathbf{X}^\top \mathbf{X}$.

The above code generates 5-dimensional data and applies classical MDS to project it onto a two-dimensional space, preserving the pairwise distances as much as possible. Extend this code to compare the original distances to the one obtained by MDS.

5.1.3 Duality between Classic MDS and PCA

The solution for classical MDS was given by the top d eigenvectors of $B = H\mathbf{X}(H\mathbf{X})^\top \in \mathbb{S}^n$. On the other hand, the d principal directions in PCA with data \mathbf{X} were defined as the top d eigenvectors of $(H\mathbf{X})^\top H\mathbf{X} \in \mathbb{S}^m$.

Suppose $n \leq m$. Let $\mathbf{u}_i \in \mathbb{R}^m$ denote the i -th principal component loading vector, where $\|\mathbf{u}_i\| = 1$ and $i = 1, \dots, m$. Thus, we have $(H\mathbf{X})^\top H\mathbf{X} = U\Lambda U^\top$, where $U = [\mathbf{u}_1, \dots, \mathbf{u}_m]$ and $\Lambda \in \mathbb{S}^m$ is the diagonal matrix with the corresponding eigenvalues. By the singular value decomposition we have, $H\mathbf{X} = V\tilde{\Lambda}^{1/2}U^\top$ for some $V \in O(n)$. Here $\tilde{\Lambda}^{1/2} \in \mathbb{R}^{n \times m}$ is a matrix such that $(\tilde{\Lambda}^{1/2})_{ii} = \sqrt{\lambda_i}$ for $i = 1, \dots, m$ and all the other entries are zero. Thus,

$$H\mathbf{X}(H\mathbf{X})^\top = V \begin{bmatrix} \Lambda & \mathbf{0} \\ \mathbf{0} & \mathbf{0} \end{bmatrix} V^\top.$$

Recall from Section 4.2 that the scores in PCA after taking the first d principal directions are given by the first d columns of $H\mathbf{X}U = V\tilde{\Lambda}^{1/2}$. In other words, we get the following result.

Theorem 5.1.1. *Consider the data $\mathbf{X} \in \mathbb{R}^{n \times m}$. The classical MDS problem on $\delta_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|$ for $i \neq j$ is equivalent to the PCA problem on the centred data $H\mathbf{X}$.*

5.2 Laplacian Eigenmaps (Spectral Embedding)

Laplacian Eigenmaps, also known as spectral embedding, is a non-linear dimensionality reduction technique that constructs low-dimensional representations of data while preserving local geometric structure. It is widely used in manifold learning and graph-based machine learning.

5.2.1 Problem Setup

Let $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ be data points in a high-dimensional space \mathbb{R}^m . The goal is to map these points into a low-dimensional space \mathbb{R}^d ($d \ll m$) such that local proximity relationships in the original space are preserved.

We start by constructing a weighted graph $G = (V, E, W)$, where:

- $V = \{1, 2, \dots, n\}$ represents the data points.

- E represents edges between points based on some notion of proximity (e.g., k -nearest neighbors or ϵ -neighborhoods).
- W is the weight matrix, where W_{ij} represents the similarity between points \mathbf{x}_i and \mathbf{x}_j .

For simplicity of the exposition we assume that G is connected and all edge weights are strictly positive.

Two common methods to define W are:

- Binary Weights:

$$W_{ij} = \begin{cases} 1, & \text{if } \mathbf{x}_j \text{ is a neighbor of } \mathbf{x}_i, \\ 0, & \text{otherwise.} \end{cases}$$

- Heat Kernel:

$$W_{ij} = \begin{cases} \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right), & \text{if } \mathbf{x}_j \text{ is a neighbor of } \mathbf{x}_i, \\ 0, & \text{otherwise.} \end{cases}$$

Here, $\sigma > 0$ controls the scale of the similarity.

5.2.2 Graph Laplacian

To compute the Laplacian Eigenmaps, we first define the graph Laplacian, which captures the structure of the graph.

The degree matrix D is a diagonal matrix where $D_{ii} = \sum_{j=1}^n W_{ij}$ represents the degree of vertex i . Since G is connected and all weights are positive, also all diagonal entries of D are strictly positive. The **graph Laplacian** is defined by

← Exercise 5.5.2

$$L := D - W. \quad (5.1)$$

In Exercise 5.5.2 you will show that L is a positive semi-definite matrix with non-trivial kernel, $L\mathbf{1} = \mathbf{0}$. All its off-diagonal entries are non-positive. Under the assumption that G is connected and all edge weights are strictly positive we can also show that $\text{rank}(L) = n - 1$.

The **normalized graph Laplacian** is

$$L_N = D^{-1/2} L D^{-1/2}.$$

Note that L_N has 1s on the diagonal.

5.2.3 Objective Function

Laplacian Eigenmaps minimize an objective that preserves local distances by penalizing large differences between embeddings of

connected points. Let $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^d$ be the embedding coordinates of the points and let $\mathbf{Y} \in \mathbb{R}^{n \times d}$ be the corresponding data matrix. The objective is:

$$\text{minimize } \frac{1}{2} \sum_{i,j} W_{ij} \|\mathbf{y}_i - \mathbf{y}_j\|^2 \quad (5.2)$$

with respect to $\mathbf{Y} \in \mathbb{R}^{n \times d}$. The interpretation here is clear. If the weight W_{ij} is high, that is when \mathbf{x}_i and \mathbf{x}_j are close, it enforces \mathbf{y}_i and \mathbf{y}_j to be close too.

To efficiently deal with this objective function, we formulate the following result.

Lemma 5.2.1. *We have $\frac{1}{2} \sum_{i,j} W_{ij} \|\mathbf{y}_i - \mathbf{y}_j\|^2 = \text{tr}(\mathbf{Y}^\top L \mathbf{Y})$.*

Proof. We use similar arguments as in Section 5.1.2. Note that $(\mathbf{Y}\mathbf{Y}^\top)_{i,j} = \mathbf{y}_i^\top \mathbf{y}_j$ and that the diagonal entries of $\mathbf{1}\mathbf{1}^\top W$ and $W\mathbf{1}\mathbf{1}^\top$ are in both cases equal to the diagonal entries of D^1 . Also, directly by definition, $W\mathbf{1} = D\mathbf{1}$. We also easily verify that the matrix

$$E = \text{diag}(\mathbf{Y}\mathbf{Y}^\top)\mathbf{1}\mathbf{1}^\top + \mathbf{1}\mathbf{1}^\top \text{diag}(\mathbf{Y}\mathbf{Y}^\top) - 2\mathbf{Y}\mathbf{Y}^\top$$

¹Indeed, $(\mathbf{1}\mathbf{1}^\top W)_{i,i} = \mathbf{e}_i^\top \mathbf{1}\mathbf{1}^\top W \mathbf{e}_i = \mathbf{1}^\top W \mathbf{e}_i$ and $(W\mathbf{1}\mathbf{1}^\top)_{i,i} = \mathbf{e}_i^\top W\mathbf{1}\mathbf{1}^\top \mathbf{e}_i = \mathbf{e}_i^\top W\mathbf{1}$. Both quantities are equal because W is symmetric.

has entries $E_{ij} = \|\mathbf{y}_i - \mathbf{y}_j\|^2$ and so $\frac{1}{2} \sum_{i,j} W_{ij} \|\mathbf{y}_i - \mathbf{y}_j\|^2 = \frac{1}{2} \text{tr}(WE)$. Since D is a diagonal matrix and E has zeros on the diagonal, we get that $\text{tr}(WE) = -\text{tr}(LE)$. Since $L\mathbf{1} = \mathbf{0}$ we get also that $\text{tr}(LE) = -2\text{tr}(L\mathbf{Y}\mathbf{Y}^\top)$. It follows that

$$\frac{1}{2} \sum_{i,j} W_{ij} \|\mathbf{y}_i - \mathbf{y}_j\|^2 = \frac{1}{2} \text{tr}(WE) = -\frac{1}{2} \text{tr}(LE) = \text{tr}(L\mathbf{Y}\mathbf{Y}^\top),$$

which concludes the proof. \square

By Lemma 5.2.1, the optimization problem in (5.2) is equivalent to minimizing $\text{tr}(\mathbf{Y}^\top L \mathbf{Y})$ with respect to $\mathbf{Y} \in \mathbb{R}^{n \times d}$. To avoid trivial solutions (e.g., $\mathbf{Y} = \mathbf{0}$), we impose:

- Orthogonality constraint: $\mathbf{Y}^\top D \mathbf{Y} = I_d$.
- Centering constraint (optional): $\mathbf{Y}^\top D \mathbf{1} = \mathbf{0}$.

The optimization problem becomes:

$$\text{minimize } \text{tr}(\mathbf{Y}^\top L \mathbf{Y}) \text{ subject to } \mathbf{Y}^\top D \mathbf{Y} = I_d \text{ and } \mathbf{Y}^\top D \mathbf{1} = \mathbf{0}.$$

Denoting $\tilde{\mathbf{Y}} = D^{1/2} \mathbf{Y}$ we equivalently get

$$\text{minimize } \text{tr}(\tilde{\mathbf{Y}}^\top L_N \tilde{\mathbf{Y}}) \text{ subject to } \tilde{\mathbf{Y}}^\top \tilde{\mathbf{Y}} = I_d \text{ and } \tilde{\mathbf{Y}}^\top D^{1/2} \mathbf{1} = \mathbf{0}. \quad (5.3)$$

Denote by $\mathbf{y}_1, \dots, \mathbf{y}_d$ the **columns** of $\tilde{\mathbf{Y}}$. The constraints in (5.3) say that $\mathbf{y}_1, \dots, \mathbf{y}_d$ are orthonormal and they are all orthogonal to the vector $\mathbf{y}_0 = D^{1/2} \mathbf{1}$. This vector is special as it corresponds to the eigenvector of L_N with eigenvalue 0.

To link (5.3) to a known optimization problem, note that

$$\text{tr}(\tilde{\mathbf{Y}}^\top L_N \tilde{\mathbf{Y}}) = \mathbf{y}_1^\top L_N \mathbf{y}_1 + \cdots + \mathbf{y}_d^\top L_N \mathbf{y}_d.$$

We know from our discussion PCA that the optimal solution can be found by first optimizing $\mathbf{y}_1^\top L_N \mathbf{y}_1$ subject to $\|\mathbf{y}_1\| = 1$ and $\mathbf{y}_1 \perp \mathbf{y}_0$, then proceed recursively, at each step $t = 2, \dots, d$ optimizing $\mathbf{y}_t^\top L_N \mathbf{y}_t$ subject to the restriction that $\|\mathbf{y}_t\| = 1$ and $\mathbf{y}_t \perp \mathbf{y}_0, \dots, \mathbf{y}_{t-1}$. The solution is obtained by finding the first $d + 1$ eigenvectors of the eigenproblem:

$$L_N \mathbf{y} = \lambda \mathbf{y},$$

corresponding to $d + 1$ smallest eigenvalues, disregarding the first eigenvector, and studding the reminaing d vectors as columns of $\tilde{\mathbf{Y}}$.

5.2.4 Example: Twisted curve

To have a simple illustration of this method, consider data that align along a twisted curve as depicted in Figure 5.1. The following code takes the corresponding data (blue points on the curve) and feeds them both to PCA and Laplacian eigenmap method.

```
# Load libraries
library(Rdimtools)
library(ggplot2)
library(scatterplot3d)
library(gridExtra)

# Generate twisted curve data
set.seed(123)
n <- 1000
t <- 1.5 * pi * (1 + 2 * runif(n)) # Angle
x <- t * cos(t)
y <- t * sin(t)
z <- 100*t # Add height to the roll
twisted <- data.frame(x = x, y = y, z = z)

# Plot 3D twisted curve
scatterplot3d(twisted$x, twisted$y, twisted$z, color = 'blue',
               main = "Twisted curve", xlab = "x", ylab = "y", zlab = "z")

# Convert twisted curve data to matrix
twisted_matrix <- as.matrix(twisted)

# Apply PCA
pca_result <- prcomp(twisted_matrix, center = TRUE, scale. = TRUE)
pca_2d <- as.data.frame(pca_result$x[, 1:2]) # Select first 2 principal
# components
colnames(pca_2d) <- c("PC1", "PC2")

# Apply Laplacian Eigenmaps
lapeig_result <- do.lapeig(twisted_matrix, ndim = 2, type =
  c("proportion", 0.1))
lapeig_2d <- as.data.frame(lapeig_result$Y)
colnames(lapeig_2d) <- c("LE1", "LE2")
```

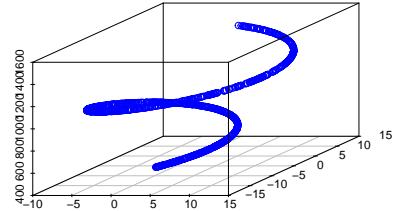


Figure 5.1: Twisted curve in 3D.

```

# Plot PCA result
p1 <- ggplot(pca_2d, aes(x = PC1, y = PC2, color = z)) +
  geom_point() +
  labs(title = "PCA Projection",
       x = "Principal Component 1", y = "Principal Component 2") +
  theme_minimal()

# Plot Laplacian Eigenmaps result
p2 <- ggplot(lapeig_2d, aes(x = LE1, y = LE2, color = z)) +
  geom_point() +
  labs(title = "Eigenmaps Projection",
       x = "Laplacian Dimension 1", y = "Laplacian Dimension 2") +
  theme_minimal()

# Display plots side by side using gridExtra
grid.arrange(p1, p2, ncol = 2)

```

SOME OTHER DIMENSION REDUCTION METHODS

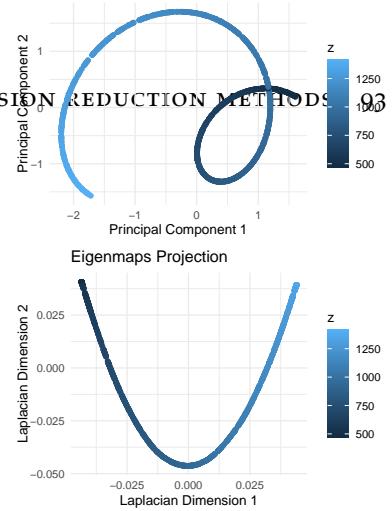


Figure 5.2: PCA and Laplacian eigenmap projections of the twisted curve data. The coloring corresponds to the value of the z-axis.

As observed in Figure 5.2 the principal directions chosen by PCA project the data so that points that are very far away from each other get glued together. On the other hand the nonlinear embedding defined by the eigenmap separated all the points. In other words, the method is able to learn the underlying manifold and propose an embedding that preserves the geometry.

5.2.5 Theoretical background story*

I did not provide any explanation behind the Laplacian eigenmap method. This story is however very beautiful and I recommend you to read about this if you have the necessary background. On a very high level: The Laplace operator Δ is a second-order differential operator in \mathbb{R}^m . For each twice differentiable function it is defined by

$$\Delta f = \sum_{i=1}^m \frac{\partial^2 f}{\partial x_i^2}.$$

The spectrum of the Laplace operator consists of all eigenvalues $\lambda \in \mathbb{R}$ for which there is a corresponding eigenfunction f with $\Delta f = \lambda f$. This is known as the Helmholtz equation. If Ω is a bounded domain in \mathbb{R}^m , then the eigenfunctions of the Laplacian are an orthonormal basis for the Hilbert space $L^2(\Omega)$.

Now, more generally, on a smooth Riemannian manifold \mathcal{M} , the Laplace-Beltrami operator $\Delta_{\mathcal{M}}$ is a differential operator that generalizes the Laplacian from Euclidean space. Eigenfunctions of $\Delta_{\mathcal{M}}$ (solutions to $\Delta_{\mathcal{M}}f = \lambda f$) reveal the intrinsic geometry of the manifold. When we do not have explicit access to the manifold \mathcal{M} , we work with a discrete approximation using a graph constructed from the data points $\{\mathbf{x}_i\} \subset \mathcal{M}$. The graph Laplacian acts as a discrete version of the Laplace-Beltrami operator. As the number of data points

increases and the graph becomes finer, L converges to $\Delta_{\mathcal{M}}$ under certain conditions. In particular, the eigenvectors of L approximate the eigenfunctions of $\Delta_{\mathcal{M}}$.

5.3 Uniform Manifold Approximation and Projection (UMAP)

One of the limitations of PCA is that it works well only when the first 2–3 principal components capture most of the variability in the data. Multidimensional scaling (MDS) and Laplacian eigenmaps attempt to address this by preserving local distances and capturing some nonlinear patterns. However, both methods have notable drawbacks. MDS becomes slow and impractical for large datasets, while Laplacian eigenmaps tend to focus too much on local structure and often distort the global arrangement, causing distant points to collapse together (a problem known as ‘crowding’).

UMAP (Uniform Manifold Approximation and Projection) was developed to address these issues by offering a better balance between local and global structure. Unlike eigenmaps, UMAP retains both local neighborhoods and the broader shape of the data, producing more meaningful embeddings. Additionally, UMAP scales efficiently to large datasets and is more robust to parameter choices. This makes UMAP a powerful tool for visualizing complex data and uncovering patterns that methods like PCA or eigenmaps might miss.

This exposition is based on the [excellent tutorial](#) by Ghojogh et al.², while the [original UMAP paper](#) offers more technical details.

5.3.1 Key idea

Consider data $\mathbf{x}_1, \dots, \mathbf{x}_n$ in \mathbb{R}^m , where \mathbb{R}^m can be very big. Like for eigenmaps, the key idea of UMAP and many similar methods is to embed the data points into a much lower dimensional subspace \mathbb{R}^d , say $\mathbf{y}_i = \text{UMAP}(\mathbf{x}_i)$ so that two points \mathbf{y}_i and \mathbf{y}_j are close to each other if and only if the corresponding $\mathbf{x}_i, \mathbf{x}_j$ are close to each other.

5.3.2 Example: Applying UMAP to the Iris Dataset

Before we present the method, we will see it in action.

The Iris dataset is a classic example in multivariate statistics. It consists of 150 observations of iris flowers, with four continuous features (sepal length, sepal width, petal length, petal width) and a categorical label (species).

A linear method like PCA can give some separation between species in the first two principal components, but this separation is not always very distinct. UMAP, on the other hand, attempts

² Benyamin Ghojogh, Ali Ghodsi, Fakhri Karray, and Mark Crowley. Uniform manifold approximation and projection (umap) and its variants: tutorial and survey. *arXiv:2109.02508*, 2021

to preserve both local and global structures of the data in a low-dimensional embedding. As a result, UMAP often produces a more visually separable and intuitively meaningful representation.

The following code compares the performance of both PCA and UMAP on the Iris dataset.

```
# Load necessary libraries
library(umap)

# Load the Iris data
data(iris)
iris_data <- iris[, 1:4]          # numeric features
iris_labels <- iris$Species      # species labels


set.seed(123)

#### UMAP ####
umap_result <- umap(iris_data)

# UMAP Plot
plot(umap_result$layout, col = as.integer(iris_labels), pch = 19, xlab =
  "UMAP1", ylab = "UMAP2", main = "UMAP Projection of the Iris Dataset")
legend("topright", legend = levels(iris_labels), col = 1:3, pch = 19)

#### PCA ####
pca_result <- prcomp(iris_data, scale. = TRUE)
pca_scores <- pca_result$x

# PCA Plot using the first two principal components
plot(pca_scores[, 1], pca_scores[, 2], col = as.integer(iris_labels), pch =
  19, xlab = "PC1", ylab = "PC2", main =
  "PCA Projection of the Iris Dataset")
legend("topright", legend = levels(iris_labels), col = 1:3, pch = 19)
```



In the resulting plots, we see that observations from the same species form relatively distinct clusters. Compared to a simple PCA plot, these groups may be more compact and better separated. This suggests that UMAP can capture nonlinear relationships and subtle differences between classes that PCA might not.

5.3.3 High-level description of the algorithm

The UMAP algorithm performs dimensionality reduction in three key steps:

1. construction of a k -nearest neighbor (kNN) graph,
2. initialization using Laplacian eigenmap embedding,
3. optimization via stochastic gradient descent (SGD) to minimize a carefully designed cost function.

Below, we provide more details for each step.

5.3.4 Step 1: Data Graph in the Input Space

Consider a training dataset $\mathbf{X} \in \mathbb{R}^{n \times m}$, where n is the sample size. We first construct a k -Nearest Neighbors (kNN) graph for this dataset, with $k = 15$ by default. More precisely, we construct a graph such that each node, represented by a data point \mathbf{x}_i , is connected to another point \mathbf{x}_j if \mathbf{x}_j is among the k closest neighbors of \mathbf{x}_i . Here the distance is arbitrary but we often use the Euclidean distance. For each data point \mathbf{x}_i , let \mathcal{N}_i denote its set of k nearest neighbors.

We further refine this procedure as follows. For each pair of points \mathbf{x}_i and its neighbour $\mathbf{x}_j \in \mathcal{N}_i$, we compute the probability $p_{j|i} \in (0, 1)$ that \mathbf{x}_i selects \mathbf{x}_j as a neighbor:

$$p_{j|i} = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\| - \rho_i}{\sigma_i}\right), \quad (5.4)$$

where ρ_i is the distance from \mathbf{x}_i to its nearest neighbour³, and σ_i is a scaling factor computed such that the total similarity to neighbours is normalized:

$$\sum_{j=1}^n p_{j|i} = \log_2(k).$$

This is a directional similarity measure, but we symmetrize it to ensure symmetry between points:

$$p_{ij} = p_{j|i} + p_{i|j} - p_{j|i}p_{i|j}. \quad (5.5)$$

5.3.5 Data Graph in the Embedding Space

UMAP aims to create a lower-dimensional embedding $\mathbf{Y} = [\mathbf{y}_1, \dots, \mathbf{y}_n] \in \mathbb{R}^{n \times d}$ of the original data $\mathbf{X} \in \mathbb{R}^{n \times m}$, where $d \ll m$. The similarity between two points in this embedding space is measured as:

$$q_{ij} = \frac{1}{1 + a\|\mathbf{y}_i - \mathbf{y}_j\|^{2b}}, \quad (5.6)$$

where $a > 0$ and $b > 0$ are hyperparameters typically set to $a \approx 1.929$ and $b \approx 0.7915$.

We still need to explain better how the embedding $\mathbf{y}_1, \dots, \mathbf{y}_n$ is defined. The goal of UMAP is to make the data graph in the low-dimensional embedding space as similar as possible to the graph in the input space. This is achieved by minimizing the fuzzy cross-entropy cost:

$$c(\mathbf{y}_1, \dots, \mathbf{y}_n) = \sum_{i \neq j} \left(p_{ij} \log \frac{p_{ij}}{q_{ij}} + (1 - p_{ij}) \log \frac{1 - p_{ij}}{1 - q_{ij}} \right). \quad (5.7)$$

This cost function has two terms: the attractive force (first term), which brings neighbors closer in the embedding, and the repulsive

← Exercise 5.5.3

³ So that the closest point in \mathcal{N}_i becomes a neighbour with probability 1.

← Exercise 5.5.4

← Exercise 5.5.5

← Exercise 5.5.6

force (second term), which pushes non-neighbor points apart. Minimizing this function ensures that points that are neighbors in the input space are also neighbors in the embedding space and vice versa.

← Exercise 5.5.7

5.3.6 Step 2: Initialization using the spectral embedding

UMAP initializes the embedding $\mathbf{y}_1, \dots, \mathbf{y}_n$ in a low-dimensional space using a Laplacian eigenmap discussed in Section 5.2. The spectral embedding step ensures that the initial positions of points in the low-dimensional space approximate the global structure of the graph, reducing the computational burden on the subsequent optimization phase.

5.3.7 Step 3: Training Algorithm

To refine the initial embedding, UMAP minimizes the cost function in (5.7), which depends on $\mathbf{y}_1, \dots, \mathbf{y}_n$ through (5.6). UMAP employs stochastic gradient descent (SGD) to optimize the cost function c .

Unlike traditional gradient descent, which updates all points simultaneously, SGD: 1. Samples pairs of points $(\mathbf{y}_i, \mathbf{y}_j)$ randomly in mini-batches. 2. Updates \mathbf{y}_i and \mathbf{y}_j incrementally based on the gradient of c_1 with respect to their positions.

5.4 Autoencoders

Autoencoders provide another approach to dimensionality reduction, offering a neural network-based alternative to methods like PCA, MDS, eigenmaps, and UMAP. While UMAP is effective at preserving local and global structures, it is a non-parametric method that does not directly produce a model for embedding new data. Autoencoders address this limitation by learning a parametric mapping from the input space to a low-dimensional representation, making it possible to quickly embed new data without retraining.

UMAP and autoencoders differ in their approaches to dimensionality reduction. UMAP focuses on preserving local neighborhoods using a graph-based approach, while autoencoders minimize reconstruction error through neural networks. UMAP often produces more interpretable global structures when the manifold is well-defined, but autoencoders can model more complex nonlinear patterns, especially when using deep networks. Additionally, UMAP is non-parametric, meaning it requires re-computation for new data points, whereas autoencoders, as parametric models, can quickly embed new data once trained. Unlike UMAP, autoencoders can also incorporate labels during training to produce class-aware embeddings. Moreover,

it is possible to combine the two approaches: an autoencoder can be trained on UMAP embeddings, blending UMAP's neighborhood preservation with the generalization capabilities of a neural network.

An autoencoder consists of two neural networks: an encoder, which maps input data $\mathbf{X} \in \mathbb{R}^{n \times d}$ to a lower-dimensional representation $\mathbf{Z} \in \mathbb{R}^{n \times p}$ with $p < d$, and a decoder, which reconstructs the input from \mathbf{Z} . The training objective is to minimize the reconstruction error, typically measured using mean squared error (MSE):

$$L(\mathbf{X}, \hat{\mathbf{X}}) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2,$$

where $\hat{\mathbf{X}}$ is the output of the decoder.

In summary, autoencoders offer a parametric, flexible approach that can learn mappings from high-dimensional data to low-dimensional representations. They are valuable for supervised learning, efficient generalization to new data, and can complement methods like UMAP to combine local neighborhood preservation with powerful, learned embeddings.

We can use the code below to do some basic dimensionality reduction on the MNIST dataset. We provide UMAP results for comparison.

```
library(keras)
library(umap)
library(ggplot2)
library(gridExtra)

# Load MNIST dataset
mnist <- dataset_mnist()
x_train <- mnist$train$x / 255
x_test <- mnist$test$x / 255

# Flatten the images (from 28x28 to 784)
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test <- array_reshape(x_test, c(nrow(x_test), 784))

##### ----- Autoencoder ----- #####
# Autoencoder architecture
input_layer <- layer_input(shape = c(784))

# Encoder
encoded <- input_layer %>%
  layer_dense(units = 128, activation = 'relu') %>%
  layer_dense(units = 64, activation = 'relu') %>%
  layer_dense(units = 2, activation = 'linear') # 2D embedding

# Decoder
decoded <- encoded %>%
  layer_dense(units = 64, activation = 'relu') %>%
  layer_dense(units = 128, activation = 'relu') %>%
  layer_dense(units = 784, activation = 'sigmoid')
```

```

# Compile Autoencoder
autoencoder <- keras::model(input_layer, decoded)
autoencoder %>% keras::compile(
  optimizer = keras::optimizer_adam(amsgrad = TRUE),
  loss = keras::loss_mean_squared_error()
)

# Train Autoencoder
history <- autoencoder %>% fit(
  x_train, x_train,
  epochs = 20,
  batch_size = 256,
  validation_data = list(x_test, x_test)
)

# Encoder Model for Dimensionality Reduction
encoder <- keras::model(input_layer, encoded)
low_dim_embeddings <- encoder %>% predict(x_test)

# Plot Autoencoder results
embedding_df <- data.frame(
  X1 = low_dim_embeddings[,1],
  X2 = low_dim_embeddings[,2],
  label = as.factor(mnist$test$y)
)

p1 <- ggplot(embedding_df, aes(x = X1, y = X2, color = label)) +
  geom_point(alpha = 0.6) +
  labs(title = "2D Representation (Autoencoder)") +
  theme_minimal()

##### ----- UMAP ----- #####
# Run UMAP for comparison
umap_result <- umap(x_test, n_neighbors = 15, min_dist = 0.1, n_components
                     = 2)

# Plot UMAP results
umap_df <- data.frame(
  X1 = umap_result$layout[,1],
  X2 = umap_result$layout[,2],
  label = as.factor(mnist$test$y)
)

p2 <- ggplot(umap_df, aes(x = X1, y = X2, color = label)) +
  geom_point(alpha = 0.6) +
  labs(title = "2D Representation (UMAP)") +
  theme_minimal()

##### ----- Comparison Plot ----- #####
grid.arrange(p1, p2, ncol = 2)

```

The resulting plot in given in Figure 5.4. From this picture we see that Autoencoder does not do a great job separating the groups. UMAP seems to be much better at this. On the other hand, Autoencoder learns the function that gives this lower-dimensional embedding. For any new dataset, we can apply



As an exercise, you can play around with different network architectures (e.g. the number of layers to see if the performance of the Autoencoder can be improved.

5.5 Exercises

Exercise 5.5.1. Suppose $C \in \bar{\mathbb{S}}^n$. Show that $C_{ii} + C_{jj} - 2C_{ij} \geq 0$. Show that the distances $d_{ij}^2 = C_{ii} + C_{jj} - 2C_{ij}$ satisfy the triangle inequality.

Exercise 5.5.2. Show that the graph Laplacian matrix L in (5.1) is positive definite and it has at least one zero eigenvalue.

Exercise 5.5.3. Show that $p_{j|i}$ defined in (5.4) lies in the interval $(0, 1)$.

Exercise 5.5.4. Show that symmetrized probability p_{ij} in UMAP defined in (5.5) is nonnegative.

Exercise 5.5.5. Let p, q be two probability distributions over some space \mathcal{X} . Define the Kullback-Leibler divergence $\text{KL}(p, q) = \mathbb{E}_p \log \frac{p(X)}{q(X)}$. Show that $\text{KL}(p, q) \geq 0$ and it is equal to zero if and only if p, q define the same distribution. Hint: Use the Jensen's inequality and the fact that $-\log$ is a strictly convex function.

Exercise 5.5.6. Suppose $X_i \sim \text{Bern}(p_i)$ and $Y_j \sim \text{Bern}(q_j)$ are all independent. Show that the Kullback-Leibler divergence between the distribution p of $X = (X_1, \dots, X_m)$ and the distribution q of $Y = (Y_1, \dots, Y_m)$ is given

by the formula

$$\text{KL}(p, q) = \sum_{i=1}^m \left(p_i \log \frac{p_i}{q_i} + (1 - p_i) \log \frac{1-p_i}{1-q_i} \right).$$

Exercise 5.5.7. Show that the expression in (5.7) is always nonnegative and it is equal to zero if and only if $p_{ij} = q_{ij}$ for all $i \neq j$.

Exercise 5.5.8. Implement classical MDS on a random dataset. Compare the distances between the original points and the MDS embedding.

Exercise 5.5.9. Compare PCA, UMAP, and Laplacian Eigenmaps on the Iris dataset. Which method better separates the species?

Exercise 5.5.10. Modify the Laplacian Eigenmaps code to explore the effect of changing the number of neighbors in the graph.

Exercise 5.5.11. Use an autoencoder on the MNIST dataset to reduce dimensionality. Compare the results with UMAP.

Exercise 5.5.12. Explain why classical MDS and PCA give the same embedding for Euclidean distances.

Exercise 5.5.13. What is the significance of the first nonzero eigenvalue of the graph Laplacian in Laplacian Eigenmaps?

Exercise 5.5.14. Show that the graph Laplacian L is positive semi-definite and has at least one zero eigenvalue.

6

Canonical Correlation Analysis

In this chapter, we consider two datasets $\mathbf{X} \in \mathbb{R}^{n \times p}$ and $\mathbf{Y} \in \mathbb{R}^{n \times q}$ that represent observations of vectors $X \in \mathbb{R}^p$ and $Y \in \mathbb{R}^q$, respectively. The goal is to find non-zero linear combinations $\mathbf{a}^\top \mathbf{X}$ and $\mathbf{b}^\top \mathbf{Y}$ with the largest possible correlation. Canonical Correlation Analysis (CCA) addresses this problem by finding pairs of linear combinations from each dataset that are maximally correlated.

Example 6.0.1 (Human Connectome Project). *The Human Connectome Project collects data on brain connectivity patterns (e.g., from functional MRI scans) and behavioral measures (e.g., cognitive performance, personality traits, and lifestyle factors). Let \mathbf{X} represent brain connectivity features and \mathbf{Y} represent behavioral and demographic variables. Canonical Correlation Analysis (CCA) can uncover pairs of linear combinations from these two sets that are maximally correlated, revealing how certain patterns of brain activity relate to cognitive abilities or mental health outcomes. For instance, one canonical pair may indicate that increased connectivity in specific brain regions is strongly associated with better working memory and fluid intelligence.*

6.1 Population CCA

We proceed similarly like for PCA by first discussing the population case.

6.1.1 Principal Correlation Vectors

Let $\Sigma_{11} = \text{var}(\mathbf{X})$, $\Sigma_{22} = \text{var}(\mathbf{Y})$, and $\Sigma_{12} = \text{cov}(\mathbf{X}, \mathbf{Y})$. The correlation between $\mathbf{a}^\top \mathbf{X}$ and $\mathbf{b}^\top \mathbf{Y}$ is given by

$$\rho(\mathbf{a}, \mathbf{b}) := \frac{\mathbf{a}^\top \Sigma_{12} \mathbf{b}}{\sqrt{\mathbf{a}^\top \Sigma_{11} \mathbf{a} \mathbf{b}^\top \Sigma_{22} \mathbf{b}}}.$$

This expression can be challenging to optimize directly. To simplify, we reparametrize the problem using $\boldsymbol{\alpha} = \Sigma_{11}^{1/2} \mathbf{a}$ and $\boldsymbol{\beta} = \Sigma_{22}^{1/2} \mathbf{b}$. This

transforms the correlation expression into

$$\rho(\boldsymbol{\alpha}, \boldsymbol{\beta}) = \frac{\boldsymbol{\alpha}^\top \Sigma_{11}^{-1/2} \Sigma_{12} \Sigma_{22}^{-1/2} \boldsymbol{\beta}}{\|\boldsymbol{\alpha}\| \|\boldsymbol{\beta}\|} = \frac{\boldsymbol{\alpha}^\top M \boldsymbol{\beta}}{\|\boldsymbol{\alpha}\| \|\boldsymbol{\beta}\|} = \left(\frac{\boldsymbol{\alpha}}{\|\boldsymbol{\alpha}\|} \right)^\top M \left(\frac{\boldsymbol{\beta}}{\|\boldsymbol{\beta}\|} \right),$$

where we define the matrix

$$M = \Sigma_{11}^{-1/2} \Sigma_{12} \Sigma_{22}^{-1/2}.$$

To maximize the correlation, we now solve the following optimization problem:

$$\max_{\boldsymbol{\alpha}, \boldsymbol{\beta}} \boldsymbol{\alpha}^\top M \boldsymbol{\beta} \quad \text{subject to} \quad \|\boldsymbol{\alpha}\| = \|\boldsymbol{\beta}\| = 1.$$

The Lagrangian for this constrained optimization is given by:

$$L = 2\boldsymbol{\alpha}^\top M \boldsymbol{\beta} - \sigma(\boldsymbol{\alpha}^\top \boldsymbol{\alpha} - 1) - \sigma'(\boldsymbol{\beta}^\top \boldsymbol{\beta} - 1).$$

At a stationary point, we get the following system of equations:

$$M \boldsymbol{\beta} = \sigma \boldsymbol{\alpha} \quad \text{and} \quad M^\top \boldsymbol{\alpha} = \sigma' \boldsymbol{\beta}.$$

By computing $\boldsymbol{\alpha}^\top M \boldsymbol{\beta}$ from both expressions, we conclude that $\sigma = \sigma'$. From this, it follows that

$$M^\top M \boldsymbol{\beta} = \sigma^2 \boldsymbol{\beta} \quad \text{and} \quad MM^\top \boldsymbol{\alpha} = \sigma^2 \boldsymbol{\alpha}.$$

Thus, the optimal $\boldsymbol{\alpha}$ is the eigenvector of MM^\top corresponding to the largest eigenvalue, and the optimal $\boldsymbol{\beta}$ is the eigenvector of $M^\top M$ corresponding to the same eigenvalue. This eigenvalue, say λ_1 , satisfies $\lambda_1 = \sigma_1^2$, where σ_1 is the largest singular value of M .

Remark 6.1.1. If we decompose M using the Singular Value Decomposition (SVD), $M = UDV^\top$, where U and V are orthogonal matrices and D is zero outside of the diagonal, then $MM^\top = UDD^\top U^\top$ and $M^\top M = VD^\top DV^\top$. The eigenvectors of MM^\top are the left singular vectors of M , and the eigenvectors of $M^\top M$ are the right singular vectors of M . The singular values are the canonical correlations.

6.1.2 First r Principal Correlation Vectors

As in PCA, we can recursively find more canonical correlation directions after obtaining the first one. Let $\boldsymbol{a}_1, \boldsymbol{b}_1$ be the first pair of canonical correlation vectors. To find the next pair $(\boldsymbol{a}_2, \boldsymbol{b}_2)$, we solve the following constrained problem:

$$\max_{\boldsymbol{\alpha}, \boldsymbol{\beta}} \boldsymbol{\alpha}^\top M \boldsymbol{\beta} \quad \text{subject to} \quad \|\boldsymbol{\alpha}\| = \|\boldsymbol{\beta}\| = 1 \quad \text{and} \quad \boldsymbol{\alpha}^\top \boldsymbol{a}_1 = \boldsymbol{\beta}^\top \boldsymbol{b}_1 = 0.$$

The solution corresponds to the second largest singular values and their corresponding left and right singular vectors. Repeating this process provides the subsequent pairs $\boldsymbol{a}_3, \boldsymbol{b}_3, \dots, \boldsymbol{a}_r, \boldsymbol{b}_r$.

6.1.3 Correlations Between Canonical Correlation Variables

Let $\eta_1 = \mathbf{a}_1^\top \mathbf{X}$ and $\phi_1 = \mathbf{b}_1^\top \mathbf{Y}$ be the first pair of canonical correlation variables. The covariance between η_1 and ϕ_1 is given by

$$\text{cov}(\eta_1, \phi_1) = \mathbf{a}_1^\top \Sigma_{12} \mathbf{b}_1 = \mathbf{a}_1^\top M \boldsymbol{\beta}_1 = \sigma_1,$$

where σ_1 is the largest canonical correlation. Moreover,

$$\text{cov}(\eta_1, \phi_2) = \mathbf{a}_1^\top \Sigma_{12} \mathbf{b}_2 = \mathbf{a}_1^\top M \boldsymbol{\beta}_2 = \sigma_2 \mathbf{a}_1^\top \mathbf{a}_2 = 0.$$

This can be easily generalized.

Theorem 6.1.2. Let $\eta_i = \mathbf{a}_i^\top \mathbf{X}$ and $\phi_i = \mathbf{b}_i^\top \mathbf{Y}$ for $i = 1, \dots, r$ be the first r canonical correlation variables. The covariance matrix of $\boldsymbol{\eta} = (\eta_1, \dots, \eta_r)$ and $\boldsymbol{\phi} = (\phi_1, \dots, \phi_r)$ is block-diagonal:

$$\text{var} \left(\begin{bmatrix} \boldsymbol{\eta} \\ \boldsymbol{\phi} \end{bmatrix} \right) = \begin{bmatrix} I_r & D_r \\ D_r & I_r \end{bmatrix},$$

where $D_r = \text{diag}(\sigma_1, \dots, \sigma_r)$ contains the first r canonical correlations.

Hence, the variables η_1, \dots, η_r and ϕ_1, \dots, ϕ_r are uncorrelated across different pairs but correlated within each pair.

6.2 Sample CCA

In practice, we do not know the population covariance matrices $\Sigma_{11}, \Sigma_{22}, \Sigma_{12}$. Instead, we estimate them from data. Let $\mathbf{X} \in \mathbb{R}^{n \times p}$ and $\mathbf{Y} \in \mathbb{R}^{n \times q}$ represent the two datasets, and let S_{11}, S_{22}, S_{12} denote the sample covariance matrices of \mathbf{X} and \mathbf{Y} ¹. As long as S_{11} and S_{22} are invertible, we can apply the same procedure as for the population case, using the sample estimates².

The sample CCA maximizes the sample correlation between linear combinations of \mathbf{X} and \mathbf{Y} :

$$\hat{\rho}(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a}^\top S_{12} \mathbf{b}}{\sqrt{\mathbf{a}^\top S_{11} \mathbf{a} \mathbf{b}^\top S_{22} \mathbf{b}}}.$$

The sample canonical correlation vectors can be found by solving the generalized eigenvalue problem for the matrix

$$\hat{M} = S_{11}^{-1/2} S_{12} S_{22}^{-1/2}.$$

6.3 Practical Considerations and Applications

6.3.1 Data Preprocessing

Before applying CCA, both datasets \mathbf{X} and \mathbf{Y} should be centered to ensure that the covariance matrices reflect the relationships between

¹ Note that in this setting we need joint observations of \mathbf{X} and \mathbf{Y} . If such joint observations are not available (e.g. these variables come from different domains and so cannot be observed jointly) we need to employ other techniques that assume some latent joint structure.

² Recall from Section 4.7 that there may be better ways to estimate the underlying covariance matrix. This is especially important in high-dimensional situations.

the variables, not the differences in their means. If the variables have different scales, it may also be necessary to standardize them to have unit variance.

6.3.2 Applications

Canonical Correlation Analysis (CCA) is widely applied in fields where understanding relationships between two datasets is crucial. In neuroscience, CCA is used to uncover associations between brain activity, such as EEG or fMRI signals, and behavioral responses or external stimuli. This helps researchers understand which neural patterns correlate with specific cognitive or motor activities.

In genomics, CCA identifies relationships between gene expression levels and phenotypic traits. By analyzing how combinations of genes relate to observable characteristics, it aids in discovering genetic influences on traits or diseases.

Marketing research also benefits from CCA. When consumer survey data, such as preferences and opinions, are paired with purchasing behavior records, CCA reveals which combinations of attitudes drive purchasing patterns. This insight helps businesses tailor their marketing strategies effectively.

Additionally, CCA is vital for multimodal data fusion. For example, when both image and text data are collected for the same observations, CCA highlights how visual patterns align with textual descriptions, supporting tasks such as cross-modal retrieval and analysis.

In economics and finance, CCA explores connections between economic indicators, such as inflation and unemployment, and financial metrics, like stock prices and interest rates. Such analyses help reveal how macroeconomic trends jointly influence financial markets.

6.4 Interpretation and Visualization of CCA Results

Interpreting CCA results requires clear visualization techniques that reveal relationships between canonical variables and the original data. Below are three common approaches:

6.4.1 Biplots

Biplots illustrate the relationship between original variables and canonical variables. Typically, pairs of canonical variables (η_1, ϕ_1) , (η_2, ϕ_2) , and so on are plotted in 2D. Original variables are shown as vectors, with their length and direction indicating their correlation with the canonical variables.

6.4.2 Heatmaps of Canonical Correlations

A heatmap of the canonical correlation matrix $D_r = \text{diag}(\sigma_1, \sigma_2, \dots, \sigma_r)$ provides a visual summary of correlations between (η_i, ϕ_j) . Darker regions on the heatmap indicate stronger canonical correlations.

6.5 Example: Real Data Analysis in R

To demonstrate Canonical Correlation Analysis, we will use the `mtcars` dataset, which contains various automobile attributes. We will split the dataset into two groups of variables:

- Group 1 (X); Variables related to performance: `mpg` (Miles per gallon), `hp` (Gross horsepower), `drat` (Rear axle ratio).
- Group 2 (Y); Variables related to design: `wt` (Weight (1000 lbs)), `qsec` (1/4 mile time), `gear` (Number of forward gears).

Our goal is to identify pairs of canonical variables that capture the strongest relationships between these two groups of variables.

Step 1: Load and Preprocess Data

We load the `mtcars` dataset and split it into two groups of variables. All variables are standardized to ensure comparability.

```
#Load necessary library and data
library(CCA)
data(mtcars)

#Group 1: Performance-related variables
X <- mtcars[, c('mpg', 'hp', 'drat')]

#Group 2: Design-related variables
Y <- mtcars[, c('wt', 'qsec', 'gear')]

#Standardize the datasets
X <- scale(X)
Y <- scale(Y)
```

Step 2: Perform Canonical Correlation Analysis

We use the `cancor` function in R to compute the canonical correlations and the corresponding canonical variables.

```
# Perform CCA
cca_result <- cancor(X, Y)

#Print the canonical correlations
print(cca_result$cor)
```

The canonical correlations are the correlations between the canonical variables derived from the two datasets. For this dataset, the values are: 0.94, 0.75, 0.22, indicating a strong relationship between the first pair of canonical variables.

Step 3: Extract Canonical Variables

The canonical variables for X and Y can be computed using the canonical weights.

```
# Extract canonical weights
x_weights <- cca_result$xcoef
y_weights <- cca_result$ycoef

# Create a data frame for weights
weights <- data.frame(
  Variable = c(colnames(X), colnames(Y)),
  Canonical1 = c(x_weights[, 1], y_weights[, 1]),
  Canonical2 = c(x_weights[, 2], y_weights[, 2]),
  Canonical3 = c(x_weights[, 3], y_weights[, 3]))

# Compute canonical scores
x_scores <- as.matrix(X) %*% x_weights
y_scores <- as.matrix(Y) %*% y_weights

# Create a data frame for visualization
scores_df <- data.frame(
  Canonical_X1 = x_scores[, 1],
  Canonical_Y1 = y_scores[, 1],
  Canonical_X2 = x_scores[, 2],
  Canonical_Y2 = y_scores[, 2])
```


Now we can plot scatter plots for the first two canonical variables in both groups.

```
library(ggplot2)

# Scatterplot for the first pair of canonical variables
ggplot(scores_df, aes(x = Canonical_X1, y = Canonical_Y1)) +
  geom_point(color = "blue", size = 3, alpha = 0.7) +
  geom_smooth(method = "lm", color = "red", linetype = "dashed") +
  labs(
    title = "Canonical Variables: First Pair",
    x = "Canonical Variable 1 (X group)",
    y = "Canonical Variable 1 (Y group)") + theme_minimal()

# Scatterplot for the second pair of canonical variables
ggplot(scores_df, aes(x = Canonical_X2, y = Canonical_Y2)) +
  geom_point(color = "green", size = 3, alpha = 0.7) +
  geom_smooth(method = "lm", color = "red", linetype = "dashed") +
  labs(
    title = "Canonical Variables: Second Pair",
    x = "Canonical Variable 2 (X group)",
```

The plots above shows the relationship between the first pair and the second pair of canonical variables. A strong linear relationship is evident, reflecting the high canonical correlation of 0.94 and 0.75.



7

Factor analysis and Independent Component Analysis

Factor Analysis (FA) and Independent Component Analysis (ICA) are two foundational methods for understanding dependencies and uncovering latent structures in multivariate data. FA models the dependence between observed variables by assuming that their variation can be largely attributed to a few common latent factors. ICA, on the other hand, seeks to decompose observed signals into statistically independent components, making it particularly useful in applications like blind-source separation. Both methods share conceptual connections with Probabilistic Principal Component Analysis (PCA), discussed in Section 4.5.2. In this chapter, we motivate and explore these models, highlighting their theoretical foundations, practical applications, and relationships to modern multivariate analysis techniques.

7.1 Motivating examples

In this section we provide some examples to build some intuition behind factor analysis and related models. In all this examples, latent representations of the data have an intuitive interpretation and allow to significantly reduce the underlying dimensionality.

7.1.1 FA: Motivating Examples

Example 7.1.1 (Capital Asset Pricing Model (CAPM)). *In finance, the Capital Asset Pricing Model (CAPM) provides a foundational framework for understanding stock returns. Suppose we observe the returns of m stocks denoted by X_1, X_2, \dots, X_m . According to CAPM, the returns of individual stocks can largely be explained by their sensitivity to a single common factor, often referred to as the market return. This is mathematically represented as:*

$$X_i = \mu_i + w_i Z + \varepsilon_i, \quad i = 1, \dots, m,$$

where Z is the market return (common factor), w_i is the sensitivity (or loading) of stock i to the market, μ_i is the mean return of stock i , and ε_i is

the idiosyncratic error, assumed to be uncorrelated across stocks.

As we will see in this chapter, FA generalizes CAPM by allowing for multiple latent factors that jointly explain the dependence structure of stock returns. For instance, additional factors could account for sector-specific trends or macroeconomic variables like interest rates or inflation.

The origins of FA can be traced back to early studies on human intelligence. We now present this classical example.

Example 7.1.2 (Human Intelligence and Pearson's Work). *In the early 20th century, Charles Spearman and Karl Pearson investigated relationships between scores on different cognitive tasks. Spearman hypothesized the existence of a single general intelligence factor (Z) that underlies performance in various tests, such as mathematics, verbal reasoning, and spatial visualization; we denote these tasks by (X_1, \dots, X_m) . This relationship can be expressed as:*

$$X_i = \mu_i + w_i Z + \varepsilon_i, \quad i = 1, \dots, m,$$

where X_i represents the score on the i -th test, Z is the general intelligence factor, w_i is the loading of test i on Z , and ε_i represents unique abilities or measurement noise.

Later studies expanded the model to include multiple latent factors representing specific abilities like memory, language, or reasoning. FA formalizes these models, enabling researchers to estimate the latent factors and their relationships with observed variables.

The next example links to recommender systems and the famous Netflix challenge. This is not directly related to the Factor Analysis model discussed here but it uses similar principles.

Example 7.1.3 (Recommender systems). *Recommender systems deal with a large user-item interaction matrix, where rows correspond to users, columns correspond to items (e.g., movies, books, or products), and entries represent interactions such as ratings, purchases, or clicks. However, the observed data is often sparse, with many missing entries.*

It is typically assumed that the observed interactions (e.g., user ratings) are influenced by a small number of underlying latent factors:

- *Latent user preferences: Each user has a set of preferences or tastes, such as a liking for action movies, romantic comedies, or documentaries.*
- *Latent item attributes: Each item can be described by its attributes, like genre, production quality, or popularity.*

We then model the interaction X_{ij} (e.g., the rating of user i for item j) as a linear combination of:

$$X_{ij} = \mathbf{z}_i^\top \boldsymbol{\lambda}_j + \varepsilon_{ij},$$

where \mathbf{z}_i is the latent vector for user i , λ_j is the latent vector for item j , ϵ_{ij} captures noise or unique factors not explained by the latent structure. Such a latent model may provide massive dimensionality reduction making the model computationally efficient. By estimating the latent factors \mathbf{z}_i and λ_j , the model predicts missing entries in the matrix (e.g., whether a user will like an unwatched movie).

7.1.2 ICA: Motivating Examples

Independent Component Analysis (ICA) aims to separate observed signals into independent components. Here are two motivating examples that highlight its importance:

Example 7.1.4 (Cocktail Party Problem). Consider the classic cocktail party problem, where multiple people are speaking simultaneously, and you have several microphones placed around the room. Each microphone captures a mixture of the voices due to overlapping sound waves. We model this as:

$$\mathbf{X} = \mathbf{WZ},$$

where \mathbf{X} is the observed mixed signal (from the microphones), \mathbf{Z} is the vector of independent source signals (e.g., individual voices), \mathbf{W} is the mixing matrix describing how the source signals combine. The goal of ICA is to recover the (unobserved) independent signals \mathbf{Z} and the mixing matrix \mathbf{W} using only the observed mixed signals \mathbf{Z} , without prior knowledge of \mathbf{W} or \mathbf{Z} .

Example 7.1.5 (Biomedical Signal Processing (EEG/MEG Data)). In neuroscience, ICA is widely used to analyze data from electroencephalography (EEG) or magnetoencephalography (MEG), which record electrical or magnetic activity in the brain. The observed signals are mixtures of:

- Neural signals originating from different brain regions.
- Artifacts, such as eye movements, muscle activity, or electrical interference.

ICA helps disentangle these sources by assuming that the neural and artifact signals are statistically independent. This enables researchers to isolate brain activity and remove unwanted artifacts, significantly improving the interpretability of the data.

7.2 Definitions

In factor analysis (FA), the idea is to explain the correlations among a set of m observed variables by introducing a set of r unobserved latent variables, or factors, where $r \leq m$. The FA model for a random

vector $X \in \mathbb{R}^m$ assumes that each observed variable can be expressed as a linear combination of these common factors, plus a unique error term. Specifically, the FA model can be written as:

$$X = \mu + WZ + \varepsilon, \quad (7.1)$$

where $Z \in \mathbb{R}^r$ is the vector of common factors (latent variables), assumed to follow a standard normal distribution $Z \sim N_r(0, I_r)$, $W \in \mathbb{R}^{m \times r}$ is the factor loading matrix, which relates the common factors Z to the observed variables X , and $\varepsilon \in \mathbb{R}^m$ is the vector of unique factors (specific errors) with $\varepsilon \sim N_m(0, \Psi)$, where Ψ is a diagonal matrix. We assume $Z \perp\!\!\!\perp \varepsilon$.

The factor loadings in W represent the relationship between each observed variable and the common factors Z . The unique variances in Ψ represent the variance in each observed variable that is not explained by the common factors. If an entry of Ψ is small, it means that the corresponding variable is well explained by the common factors. The interpretation of W is less clear and we will discuss this in more detail in Section 7.3.2 with a notable example given in Section 7.6.

Note that this model is almost exactly the same as the Probabilistic PCA model in Section 4.5.2 with the only difference that in PCA $\Psi = \sigma^2 I_m$ for some σ^2 . One important consequence of simpler assumptions of PPCA is that the model results in closed formulas for the MLE. In the case of FA, maximizing the likelihood function requires the EM algorithm or some other numerical scheme.

The Independent Component Analysis model assumes a very similar stochastic representation¹

$$X = \mu + WZ,$$

which effectively corresponds to the FA model with $\varepsilon \equiv 0$ ($\Psi = \mathbf{0}_{m \times m}$). Here however the main difference is that typically it is assumed that $r = m$ with W invertible and also that $Z \sim (0, I_r)$ with independent components but not necessarily jointly Gaussian. We will comment more on that in a second.

¹ In practice, both before applying ICA and FA, the data are centred. So, with no loss of generality, we could assume $\mu = 0$ throughout.

7.3 Model Implications and identifiability

7.3.1 The covariance matrix

Since X in (7.1) is an affine combination of independent Gaussian vectors, it is Gaussian itself. It is clear that $\mathbb{E}X = \mu$. Moreover, the covariance matrix Σ of the observed data X can be derived from the

model as follows:

$$\begin{aligned}\Sigma &= \text{var}(WZ + \varepsilon) \\ &= W \text{var}(Z)W^\top + \text{var}(\varepsilon) \\ &= WW^\top + \Psi.\end{aligned}\tag{7.2}$$

Thus, the model imposes the constraint that the covariance matrix Σ can be decomposed into two parts: a low-rank² matrix WW^\top , which captures the shared variance among the variables due to the common factors, and a diagonal matrix Ψ that captures the unique variances of each observed variable (i.e., the variance not explained by the common factors).

² Note that $\text{rank}(WW^\top) \leq r$.

7.3.2 Lack of identifiability and Varimax notation

Unlike in PPCA, the loadings in Factor Analysis are assumed to have a meaning. This follows because the latent variable in this contexts often has a clear interpretation as in the example we gave earlier. In this section we show that we should be very careful trying to interpret the estimated loading matrix because there are infinitely many loading matrices that explain the data equally well.

To see this, note that for any orthogonal matrix $U \in \mathcal{O}(r)$ it holds that $WW^\top = WU(WU)^\top$. Thus, in the factor analysis model the parameters (W, Ψ) lead to the same observed distribution as the parameters (WU, Ψ) for any $U \in \mathcal{O}(r)$. We say that the model is not identifiable.

There are several approaches to pick the canonical loading matrix from the set of all equivalent ones. Usually this is resolved by rotating the factor loadings to satisfy an arbitrary constraint such as

$$W^\top \Psi^{-1} W \quad \text{is diagonal,}\tag{7.3}$$

where the diagonal elements are written in decreasing order.

The varimax method of orthogonal rotation is another way to get a unique W , proposed by Kaiser in 1958. Its rationale is to provide axes with a few large loadings and as many near-zero loadings as possible. This is accomplished by an iterative maximization of a quadratic function of the loadings.

Let $W \in \mathbb{R}^{m \times r}$ be a matrix of unrotated loadings. This could be, for example, the output of the maximum likelihood estimation procedure. Let $U \in \mathcal{O}(r)$ leading to the rotated loading matrix WU . We choose U using the varimax criterion that maximizes the sum of variances of the squared loadings within each column of the loading matrix, where each row is properly normalized.

To get some idea why this quantity makes sense note that we can rewrite (7.1) by defining $\tilde{\varepsilon} = \Psi^{-1/2}\varepsilon \sim N(0, I_m)$:

$$\tilde{X} = \tilde{\mu} + \tilde{W}Z + \tilde{\varepsilon}.$$

Similarly as in PCA, it is natural to assume that the columns of W are orthogonal. In other words $\tilde{W}^\top \tilde{W}$ is diagonal.

More concretely, let $M \in \mathbb{R}^{m \times r}$ be the matrix with entries

$$M_{ij} = \frac{(WU)_{ij}^2}{\sum_{k=1}^r (WU)_{ik}^2}.$$

Note that the entries of M are nonnegative and, by definition, $M\mathbf{1}_r = \mathbf{1}_m$. In other words, M is a stochastic matrix³. In the varimax approach, we choose U by maximizing $\|M - \frac{1}{m}\mathbf{1}_m\mathbf{1}_m^\top M\|_F^2$, the Frobenius distance from M to the matrix $\frac{1}{m}\mathbf{1}_m\mathbf{1}_m^\top M$.

To understand this better, note that this distance is zero if and only if $H_m M = \mathbf{0}$, where $H_m = I_m - \frac{1}{m}\mathbf{1}_m\mathbf{1}_m^\top$ is the centering matrix. In other words, each column of M is constant. By maximizing the distance over U we obtain M such that in each column there are a few big entries and the remaining entries are negligible.

The varimax rotation is implemented in the base package in R.

```
# Example data: unrotated factor loadings
loadings <- matrix(c(0.7, 0.3, 0.2, -0.5, 0.8, 0.4), nrow = 3)

# Apply varimax rotation
rotated <- varimax(loadings)

# View rotated loadings
print(rotated$loadings)
```

³ A stochastic matrix is a nonnegative matrix whose rows all sum to 1.

7.3.3 Identifiability via Non-Gaussianity*

An alternative and more robust approach to establish the identifiability of the matrix W in factor analysis and independent component analysis (ICA) relies on the assumption of non-Gaussianity of the underlying components.


One way to understand the issue of non-identifiability in the Gaussian setting is to observe that, if $Z \sim N_r(0, I_r)$, then for any orthogonal matrix $U \in \mathcal{O}(r)$, the transformed vector UZ has the same distribution as Z . In particular, the entries of UZ remain independent, regardless of the choice of U . This invariance to orthogonal transformations is a special property of Gaussian distributions, which does not generally hold for non-Gaussian random vectors.

To formalize this, we cite a result by Kac that highlights the uniqueness of the Gaussian distribution in maintaining independence under orthogonal transformations:

Theorem 7.3.1 (Kac). *If $Z = (Z_1, Z_2)$ is a random vector such that $Z_1 \perp\!\!\!\perp Z_2$ and $(UZ)_1 \perp\!\!\!\perp (UZ)_2$ for some orthogonal matrix $U \in \mathcal{O}(2)$ with all entries non-zero, then Z follows a Gaussian distribution.*

Example 7.3.2. *For a simpler illustration of this result consider two independent variables (Z_1, Z_2) distributed uniformly on $[0, 1]$. Their joint*

If $U \in \mathcal{O}(2)$ has a zero, it must be one of the eight matrices of the form $\begin{bmatrix} \pm 1 & 0 \\ 0 & \pm 1 \end{bmatrix}, \begin{bmatrix} 0 & \pm 1 \\ \pm 1 & 0 \end{bmatrix}$.



samples are depicted on the top in Figure 7.1. Rotating these data by 45 degrees, results in the bottom picture. We can think of it as a sample of the vector (Y_1, Y_2) , which is a rotation of (Z_1, Z_2) . From the picture it should be clear that Y_1 and Y_2 cannot be independent as the support of Y_1 depends on Y_2 and vice versa.

Theorem 7.3.1 can be generalized to higher dimensions in a suitable sense, illustrating that Gaussian distributions are uniquely characterized by their invariance under orthogonal transformations that preserve independence.

In the context of ICA, this property of Gaussian distributions is exploited to establish identifiability of the mixing matrix:

Theorem 7.3.3 (Comon). *If $\mathbb{E}Z = 0$, $\text{var}(Z) = I_r$, the components of Z are independent, and at most one of them is Gaussian, then the mixing matrix W can be identified up to permutations of its columns and sign changes.*

The assumption of non-Gaussianity also allows for other estimation techniques. In the Gaussian distributions the whole information about the distribution is contained in the mean vector μ and the covariance matrix Σ . In general, higher order moments can also be exploited. We briefly explain this in the next section.

7.4 Fitting FA and ICA

Several methods exist for fitting the Factor Analysis (FA) and Independent Component Analysis (ICA) models. One of the most common approaches for FA is maximum likelihood estimation, where we aim to find the parameters Λ and Ψ that maximize the likelihood of the observed data given the factor analysis model. For ICA, maximum likelihood is often replaced by the method of moments, which leverages higher-order statistics.

7.4.1 Maximum Likelihood Estimation for FA

Maximum likelihood estimation in FA aims to find parameters W (the factor loading matrix) and Ψ (the diagonal covariance of unique variances) that best explain the observed covariance structure in (7.2).

Jöreskog's method simplifies the maximum likelihood estimation by providing an analytic update for the loading matrix W and iteratively optimizing over Ψ . The key idea is to iteratively: (1) Update W using an analytic formula. (2) Numerically optimize Ψ .

Theorem 7.4.1. *Let Ψ be fixed, and let $\tilde{S}_n = \Psi^{-1/2} S_n \Psi^{-1/2}$. Using the spectral decomposition, write*

$$\tilde{S}_n = U \tilde{\Lambda} U^\top.$$

Then, the values of W that maximize the likelihood function are of the form

$$\widehat{W} = U_r \Theta Q \quad \text{for any } Q \in \mathcal{O}(m),$$

where U_r is obtained from the first r columns of U (corresponding to the largest eigenvalues) and Θ is a diagonal matrix with i -th entry equal to $\sqrt{\max\{0, \tilde{\Lambda}_{ii} - 1\}}$.

Proof Sketch. We have

$$\Sigma = \Psi + WW^\top = \Psi^{1/2}(I_m + \tilde{W}\tilde{W}^\top)\Psi^{1/2}.$$

If Ψ is fixed then maximizing the log-likelihood function $\log \det \Sigma^{-1} - \text{tr}(S_n \Sigma^{-1})$ is equivalent to maximizing

$$\log \det(I_m + \tilde{W}\tilde{W}^\top)^{-1} - \text{tr}(\tilde{S}_n(I_m + \tilde{W}\tilde{W}^\top)^{-1}).$$

The rest of the proof is essentially the same as the proof of Theorem 4.5.1 with the additional restriction that $\sigma^2 = 1$. \square

We now describe this method in detail:

Initialization: Start with an initial guess for $\Psi^{(0)}$, typically the diagonal elements of the sample covariance matrix $\hat{\Sigma}$.

Update W : For a fixed Ψ , update the loading matrix W analytically by solving:

$$W = \Sigma V^{-1/2} U_q,$$

where $V = \Psi + WW^\top$ is the fitted covariance matrix, and U_q consists of the eigenvectors corresponding to the r largest eigenvalues of $V^{-1/2} \Sigma V^{-1/2}$.

7.4.2 Method of Moments for FA and ICA

Another approach to fitting FA and ICA models is the method of moments, which leverages statistical moments of the data.

In general the method of moments refers to estimating procedures where the population moments can be explicitly expressed as functions of the model parameters $m(\theta)$. Then we try to solve $\hat{m} = m(\theta)$, where \hat{m} refers to the corresponding vector of sample moments. Of course, in general there is no guarantee that such equation has a solution or that this solution is unique so typically additional theoretical analysis is required. When there is no solution, a decent alternative is given by minimizing $\|\hat{m} - m(\theta)\|^2$.

In FA, the covariance structure of the observed data \mathbf{X} is modeled as in (7.2). The method of moments matches the sample covariance matrix S_n to its theoretical counterpart Σ . The estimation involves solving:

$$\text{minimize} \quad \|S_n - (\Lambda \Lambda^\top + \Psi)\|_F^2.$$

⁴ This is referred to as the Generalized Method of Moments.

However, this approach should not be preferred over the MLE presented in the previous section.

The method of moments approach is particularly effective for ICA, where the assumption of non-Gaussianity enables us to exploit higher-order moments such as skewness and kurtosis; see Section 3.4 for some definitions. The assumption of non-Gaussianity ensures that W is identifiable up to scaling and permutation.

One example of this algorithm is the JADE algorithm that takes a look at the fourth cumulant of the data. Since $X = WZ$ and W is invertible, we have $Z = AX$ with $A = W^{-1}$. By independence of the components of Z the $m \times m \times m \times m$ table M that contains all 4-th order cumulants of Z is diagonal. Moreover, by multilinearity of cumulants, its (i, j, k, l) -th (assuming i, j, k, l are not equal) entry satisfies

$$0 = \text{cum}(Z_i, Z_j, Z_k, Z_l) = \sum_{i', j', k', l'=1}^m A_{ii'} A_{jj'} A_{kk'} A_{ll'} \text{cum}(X_{i'}, X_{j'}, X_{k'}, X_{l'}).$$

This gives quartic equations in the entries of A that we solve by replacing the array of cumulants of X by its sample estimator. This is essentially what the JADE algorithm does. The following code presents some simple illustration.

```
library(JADE)

# Generate synthetic data
set.seed(123) # For reproducibility
n <- 1000 # Number of samples

# Create two independent components
Z1 <- runif(n, -1, 1) # Uniform distribution
Z2 <- rt(n, df = 5) # t-distribution with 5 degrees of freedom
Z <- cbind(Z1/sd(Z1), Z2/sd(Z2)) # Combine into a matrix

# Create a mixing matrix
W <- matrix(c(2, -1, 1, 2), nrow = 2)

# Generate mixed signals
X <- Z %*% t(W) # Mixed signals: observed data

# Apply the JADE algorithm
ica_result <- JADE(X, n.comp = 2)

# Check the estimated W (mixing matrix)
print(ica_result$A)
```

Note that in the code below we have

$$W = \begin{bmatrix} 2 & 1 \\ -1 & 2 \end{bmatrix} \quad \text{and} \quad \hat{W} = \begin{bmatrix} 0.74 & 2.10 \\ 2.11 & -0.76 \end{bmatrix}.$$

The matrices do not look that similar. However, W can be identified only up to permuting columns and swapping their signs so we

should compare W with permuted \hat{W} where the two columns are swapped.

7.5 Choosing the Number of Latent Components

A critical step in factor analysis is determining the number of latent components r , which represent the dimensionality of the latent space. Selecting r appropriately is essential to balance model complexity and explainability while avoiding underfitting or overfitting. Several methods exist to guide this choice, including eigenvalue-based criteria, information-theoretic approaches, parallel analysis, and likelihood ratio tests.

The main idea is to exploit the nice representation of the covariance matrix Σ in the FA model; c.f (7.2)

Eigenvalue-Based Criteria. In factor analysis, the covariance structure of the observed data \mathbf{X} is modeled as (7.2). Since Σ is not known, we need to rely on its estimator – for example the sample covariance matrix S_n . The eigenvalues of S_n reflect the total variance explained by each dimension. The goal is to retain the first r components that explain the majority of the variance.

Kaiser's Criterion suggests retaining all components with eigenvalues greater than 1. This criterion assumes that components explaining less variance than a single variable (unit variance) are not meaningful:

Retain components where $\lambda_i > 1$,

where λ_i are the eigenvalues of S_n .

A scree plot is another heuristic tool where the eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_m$ are plotted in descending order. The "elbow point," where the slope of the curve flattens, indicates the optimal number of components r . Mathematically:

$$r = \arg \min_k \frac{\lambda_k}{\lambda_{k-1}} \text{ for } k \geq 2,$$

where the largest change in slope suggests the cutoff.

Information-Theoretic Criteria. Information-theoretic approaches such as the Akaike Information Criterion (AIC) and Bayesian Information Criterion (BIC) are commonly used to balance model fit and complexity. The AIC is defined as:

$$\text{AIC} = -2 \cdot \log \mathcal{L} + 2 \cdot \text{df},$$

where \mathcal{L} is the likelihood of the model and $\text{df} = mr - \frac{r(r-1)}{2}$ is the number of free parameters. The BIC penalizes model complexity more heavily and is given by:

$$\text{BIC} = -2 \cdot \log \mathcal{L} + \log(n) \cdot \text{df},$$

where n is the sample size. The optimal number of factors is the value of r that minimizes either AIC or BIC.

7.6 Illustrative application

Here is an example using real data to illustrate how Factor Analysis can be used in practice. We use the `psych` package and the `bfi` dataset from the `psych` package, which contains responses to the Big Five Personality Test. The dataset consists of 2,800 observations (survey responses). It includes 25 personality-related questions that measure the Big Five Personality Traits: Openness (O), Conscientiousness (C), Extraversion (E), Agreeableness (A), Neuroticism (N). Responses are recorded on a Likert scale (1 to 6), where higher values indicate stronger agreement with the statement. Additionally, the dataset contains demographic information such as gender, education, age. We are going to discard these demographic data and try to show that the survey responses do support the existence of five latent personality traits.

```
# Load necessary packages
library(psych)

# Load the Big Five personality dataset from the psych package
data(bfi)

# Select a subset of personality trait variables (excluding demographic
# variables)
bfi_subset <- bfi[, 1:25] # Selecting only the 25 personality questions

# Remove missing values
bfi_subset <- na.omit(bfi_subset)

# Determine the number of factors using parallel analysis
fa_parallel <- fa.parallel(bfi_subset, fa = "fa")

# Extract the suggested number of factors, with a fallback to a default
# (e.g., 5)
num_factors <- ifelse(!is.na(fa_parallel$ncomp), fa_parallel$ncomp, 5)

# Print chosen number of factors for debugging
print(paste("Number of factors chosen:", num_factors))

# Fit a Factor Analysis model with a safe fallback for nfactors
if (!is.na(num_factors) && num_factors > 0) {
  fa_model <- fa(bfi_subset, nfactors = num_factors, rotate = "varimax")
  print(fa_model$loadings) # Print factor loadings
  fa.diagram(fa_model) # Visualize the factor structure
} else {
  # Print("Error: No valid number of factors found. Try specifying nfactors manually.")
}
```

In the code we try to automatically get the number of latent factors and the answer is... five. When we learn loadings using the varimax

rotation we get that the latent traits correspond exactly to what we would hope; see Figure 7.2.

7.7 Exercises

Exercise 7.7.1 (Basic Properties of Factor Analysis). Consider the Factor Analysis model in (7.1). Show carefully that the marginal distribution of X is Gaussian and rederive its mean and covariance matrix.

Exercise 7.7.2. If X is a Gaussian random vector whose covariance matrix can be written in the form $\Sigma = WW^\top + \Psi$, show that there exist factors Z, ϵ such that (7.1) holds.

Exercise 7.7.3. If $WW^\top = \tilde{W}\tilde{W}^\top$, show that $\tilde{W} = WQ$ for an orthogonal Q .

Exercise 7.7.4 (Identifiability in FA). Show that if (W, Ψ) satisfies the FA model, then for any orthogonal matrix $U \in \mathcal{O}(r)$, the pair (WU, Ψ) defines the same observed distribution for X . What does this imply for parameter estimation?

Exercise 7.7.5 (Maximum Likelihood Estimation in FA). The log-likelihood function in FA is given by:

$$\ell(W, \Psi) = -\frac{n}{2} \left[\log \det(WW^\top + \Psi) + \text{tr}(S_n(WW^\top + \Psi)^{-1}) \right].$$

Explain why maximizing the likelihood is difficult and why iterative methods such as EM are required.

Exercise 7.7.6 (Choosing the Number of Factors). Implement one of the methods to determine the number of factors in FA. Test it using a simulated dataset.

Exercise 7.7.7 (Comparison with PCA). Consider a dataset where all observed variables are generated from a latent factor model. Compare FA and PCA in terms of:

- The estimated loadings
- The explained variance
- The ability to recover the true latent factors

Which method is more appropriate and why?

Exercise 7.7.8 (Generalization to Non-Gaussian Factors). In standard FA, the latent factors Z are assumed to be Gaussian. Suppose instead that Z follows a multivariate Laplace distribution. What impact would this have on estimation and inference?




Figure 7.2:

Exercise 7.7.9 (ICA and FA Differences). *ICA is based on a similar decomposition to FA:*

$$X = WZ.$$

However, ICA assumes that the components of Z are independent but not necessarily Gaussian. Explain why ICA requires non-Gaussianity for identifiability.

Exercise 7.7.10 (Deriving the ICA Objective Function). *A common ICA method minimizes mutual information. Show that for a random variable Z with independent components, minimizing the mutual information $I(Z_1, Z_2, \dots, Z_m)$ is equivalent to maximizing the sum of marginal entropies $\sum_i H(Z_i)$ under a fixed covariance constraint.*

Exercise 7.7.11 (Fourth-Order Cumulants in ICA). *The JADE algorithm for ICA uses fourth-order cumulants to estimate the independent components. Derive the fourth-order cumulant tensor for a 2D random vector and show how it can be used to separate sources in ICA.*

Exercise 7.7.12 (FA vs ICA in High Dimensions). *In high-dimensional data where $m \gg n$, discuss the advantages and disadvantages of FA and ICA. How does the curse of dimensionality affect estimation in these models?*

Graphical Models

In this chapter we only cover basics of graphical models and provide a short introduction to Gaussian graphical models. If you are interested in more details of this beautiful theory, see ¹.

8.1 Introduction

Graphical models are an important tool for representing and analyzing the relationships between variables in multivariate data. They combine ideas from graph theory and probability to describe how variables depend on each other. These models provide a clear and structured way to understand complex systems, making them very useful in modern statistics and data science.

Let $X = (X_1, \dots, X_m)$ be a random vector. A graphical model for X consists of a graph $G = (V, E)$, where $V := \{1, \dots, m\}$ is a set of nodes representing the random variables, and E is a set of edges. If there is an edge between two nodes i and j , we write $ij \in E$. In graphical models, if $ij \notin E$, X_i and X_j are conditionally independent given some other variables. This way of showing how variables depend on each other is helpful for tasks like prediction, testing hypotheses, and choosing the right model.

Graphical models are used in many areas of science and industry. In genetics and biology, they help uncover how genes interact and how biological processes work. In finance and economics, they help understand how different financial factors affect each other and how risks spread through a system. Graphical models also form the basis of many machine learning methods².

There are three main types of graphical models. Undirected graphical models, also called Markov random fields, represent relationships where the direction does not matter. Directed graphical models, or Bayesian networks, show relationships with a specific direction, often representing cause-and-effect. Mixed graphical models combine both directed and undirected edges to describe more complex relationships. Each type of model is suited to different kinds of prob-

¹ Steffen L. Lauritzen. *Graphical Models*, volume 17 of *Oxford Statistical Science Series*. Clarendon Press, Oxford, 1996.
ISBN 978-0-19-852219-5

² The computational aspects of graphical models are treated more in detail in STA414. Here we focus on modelling and statistical aspects.

lems.

In this chapter, we will focus on undirected graphical models because they are easier to understand and widely used. Directed and mixed models will be mentioned briefly, but we leave a deeper discussion of these models for more advanced studies.

8.2 Appendix: Conditional independence

Independence and conditional independence are fundamental distributional restrictions that are both statistically interpretable and allow for scaling down very high dimensional problems. For that reason they have also a considerable impact on machine learning.

Suppose we have a random vector $X \in \mathbb{R}^m$ with density $f(\mathbf{x})$ ³. Let $X = (X_A, X_B)$ be an arbitrary split with $X_A \in \mathbb{R}^k$ and $X_B \in \mathbb{R}^{m-k}$.

The marginal density f_B of X_B is

$$f_B(\mathbf{x}_B) := \int_{\mathbb{R}^k} f_X(\mathbf{x}_A, \mathbf{x}_B) d\mathbf{x}_A$$

and we have the analogous definition of the marginal density f_A of X_A . We say that X_A is independent of X_B if the joint density factorizes $f(\mathbf{x}) = f_A(\mathbf{x}_A)f_B(\mathbf{x}_B)$ for all $\mathbf{x} \in \mathbb{R}^m$; we write $X_A \perp\!\!\!\perp X_B$. It can be shown that $X_A \perp\!\!\!\perp X_B$ if and only if $X_i \perp\!\!\!\perp X_j$ for all $i \in A, j \in B$.

Let $f_{A|B}$ be the conditional density of $X_A | X_B$, that is,

$$f_{A|B}(\mathbf{x}_A | \mathbf{x}_B) := \frac{f_X(\mathbf{x}_A, \mathbf{x}_B)}{f_B(\mathbf{x}_B)}.$$

Note that if $X_A \perp\!\!\!\perp X_B$ then, equivalently,

$$f_{A|B}(\mathbf{x}_A | \mathbf{x}_B) = f_A(\mathbf{x}_A).$$

This characterization of independence gives also its convenient interpretation: If we want to predict (or sample etc) X_A then the information about X_B is irrelevant.

Now consider a split $X = (X_A, X_B, X_C)$. We say that X_A is conditionally independent of X_B given X_C , which we denote by $X_A \perp\!\!\!\perp X_B | X_C$, if

$$f_{A \cup B | C}(\mathbf{x}_A, \mathbf{x}_B | \mathbf{x}_C) = f_{A|C}(\mathbf{x}_A | \mathbf{x}_C)f_{B|C}(\mathbf{x}_B | \mathbf{x}_C)$$

or equivalently

$$f_{A|B \cup C}(\mathbf{x}_A | \mathbf{x}_B, \mathbf{x}_C) = \frac{f_{A \cup B | C}(\mathbf{x}_A, \mathbf{x}_B | \mathbf{x}_C)}{f_{B|C}(\mathbf{x}_B | \mathbf{x}_C)} = f_{A|C}(\mathbf{x}_A | \mathbf{x}_C).$$

In other words, if we want to predict (sample, etc) X_A after observing the value of X_C then an additional observation of X_B should not change our predictions.

³ The same definitions hold for discrete variables. In this case f is the probability mass function and integrals below should be replaced with sums.

8.2.1 Testing independence

Suppose we observe some bivariate data $\text{Data} = (X_1, Y_1), \dots, (X_n, Y_n)$ i.i.d. from bivariate distribution $\mathbb{P}_{X,Y}$. How would you test if the components of the underlying vector are independent? We consider testing problem: $H_0 : X \perp\!\!\!\perp Y, H_1 : X \not\perp\!\!\!\perp Y$. There are many tests of independence. For example, we know that if $X_1 \perp\!\!\!\perp X_2$ then $\text{corr}(X_1, X_2) = 0$. We could then compute the sample correlation and try to draw conclusions from this. Many popular independence tests are based on the sample correlation or on its nonparametric version the Kendall's τ . Consider the following code.

```
> set.seed(1); n <- 200; rho <- 0.2; Z <- runif(n);
> X <- runif(n)^2+sqrt(rho)*Z; Y <- runif(n)+sqrt(rho)*Z
> cor.test(X, Y, method = "pearson")$p.value
[1] 0.03417231
> cor.test(X, Y, method = "kendall")$p.value
[1] 0.01100592
```

Here we sample a bivariate vector with some non-negligible dependence structure. The vector is highly non-Gaussian so the standard test based on the Pearson correlation (using Fisher's z-transform) is generally not valid, and we see that the p-value is rather high when we use this test. The nonparametric test based on the Kendall's tau is better behaved.

8.2.2 Some cautionary examples

I want to emphasize however that, using such tests is in general a very bad idea. Remember, $\text{corr}(X_1, X_2) = 0$ is **not** equivalent to independence! One important cautionary example is the dataset considered by Bowman& Azzalini (1997) on aircraft wing span and speed data.

```
> library(sm); set.seed(1);
> X <- aircraft$Span
> Y <- aircraft$Speed
> cor.test(X,Y)$p.value
[1] 0.7816014
> dcor.test(X,Y,R=1000)$p.value
[1] 0.000999001
```

The correlation based tests all retain the null hypothesis simply because the sample correlation is -0.01 and so, very close to zero. But these two variables are very highly dependent, just dependence is non-linear. Here the distance correlation test is a nonparametric test of dependence that works under general assumptions, its p-value is nearly zero, confirming dependence in the data.

I have another cautionary example regarding conditional independence. This one is often referred to as the Simpson's paradox.

Consider the admission figures of the grad school at UC Berkeley in 1973: 8442 (44%) men, 4321 (35%) women admitted. A group of statisticians started looking into these data trying to understand this potential disparity. Is there a dependence between the random variable representing the gender and another binary random variable representing The conclusion becomes strikingly different if we condition on a third variable, the department: Try to think what is

Department	Men: Applicants	Men: Admitted	Women: Applicants	Women: Admitted
A	825	62%	108	82%
B	560	63%	25	68%
C	325	37%	593	34%
D	417	33%	375	35%
E	191	28%	393	24%
F	373	6%	341	7%

a possible explanation for this seemingly paradoxical situation? As the original study puts it: "Measuring bias is harder than is usually assumed, and the evidence is sometimes contrary to expectation."⁴.

Figure 8.1 gives yet another illustration for what the Simpson paradox is. Here we observe 2D dataset with points corresponding to two groups (regions). In each group the effect is clearly positive but if we analyse the data jointly the trend is negative.

8.3 Graphs and Conditional Independence

A graph $G = (V, E)$ consists of a set of nodes $V = \{1, \dots, m\}$ and edges E . Nodes represent random variables, and edges represent direct dependencies or associations between the variables. A simple example with $m = 5$ is given in Figure 8.3. This will be the running example in this chapter.

Graphs can be represented in multiple ways. An **adjacency matrix** is a square matrix A where $A_{ij} = 1$ if there is an edge between nodes i and j , and $A_{ij} = 0$ otherwise. For example, for the graph in Figure 8.3, the adjacency matrix is

$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 \end{bmatrix}.$$

Another way to represent a graph is by listing for each node the set of all its neighbors. In our running example: $\mathcal{N}_1 = \{2, 3, 4\}$, $\mathcal{N}_2 = \{1, 3\}$, $\mathcal{N}_3 = \{1, 2, 5\}$, $\mathcal{N}_4 = \{1, 5\}$, and $\mathcal{N}_5 = \{3, 4\}$.

⁴ Bickel et al, *Sex Bias in Graduate Admissions: Data From Berkeley*, Science, 1975




Figure 8.1: Simpson's paradox in action.




Figure 8.2: A simple graph with five nodes and six edges.

A **clique** is a subset of nodes in an undirected graph where every pair of nodes is connected by an edge. A maximal clique is a clique that is maximal with respect to inclusion (no other clique strictly containing it).

A path is a sequence of nodes connected by edges. A cycle is a path that starts and ends at the same node.

Separation in graphs is a crucial concept for us. We say that two subsets $A, B \subseteq V$ are **separated** from each other by a subset $C \subseteq V$ if every path between any node $i \in A$ and $j \in B$ necessarily contains a node in C ; we write $A \perp_G B|C$. For example, in Figure 8.3, $A = \{2\}$ is separated from $B = \{4, 5\}$ by $C = \{1, 3\}$. However, $C' = \{1\}$ does not separate A from B .

In the running example, the maximal cliques are $\{1, 2, 3\}, \{1, 4\}, \{3, 5\}, \{4, 5\}$. On the other hand, $\{1, 3\}$ is a clique but it is not maximal.

8.3.1 Graph factorization

An important defining concept in graphical models is that of **factorization with respect to G** . Let G be an undirected graph with nodes $V = \{1, \dots, m\}$ and maximal cliques C_1, \dots, C_k . For any $C \subseteq V$, \mathbf{x}_C denotes a subvector of \mathbf{x} , $\mathbf{x}_C = (x_i)_{i \in C}$.

We say that density⁵ $f(x)$ of X factorizes according to G if

$$f(\mathbf{x}) = \prod_{C \in \mathcal{C}} \phi_C(\mathbf{x}_C) \quad \text{for all } \mathbf{x} \in \mathcal{X}, \quad (8.1)$$

where $\phi_C(\mathbf{x}_C) \geq 0$.⁶ For the graph in Figure 8.3,

$$f(x_1, x_2, x_3, x_4, x_5) = \phi_{123}(x_1, x_2, x_3)\phi_{14}(x_1, x_4)\phi_{35}(x_3, x_5)\phi_{45}(x_4, x_5).$$

We define the graphical model $M(G)$ as the set of all distributions on \mathcal{X} that factorize with respect to G .

⁵ When we say “density” we mean the density function for continuous distributions or the probability mass function for discrete distributions.

⁶ This is a notion of simplicity, which states that the distribution can be glued in a specific way from simpler pieces.

8.3.2 Conditional Independence and Graph Separation

Conditional independence and separation are the core ideas linking graphs to probabilistic models. In undirected graphs, conditional independence is determined by the concept of graphical separation called the **Global Markov Property**: If a set of nodes C separates A from B in the graph, then $X_A \perp\!\!\!\perp X_B | X_C$. For example, in the graph in Figure 8.3, $X_2 \perp\!\!\!\perp \{X_4, X_5\} | \{X_1, X_2\}$.

The most fundamental result of graphical models that links the Global Markov property with factorization is the Hammersley-Clifford theorem.

← Exercise 8.6.1

Theorem 8.3.1 (Hammersley-Clifford theorem). *Let $f > 0$ be a density function for $X = (X_1, \dots, X_m)$. Then the following are equivalent:*

(F) f factorizes according to $G = (V, E)$.

(G) $X_A \perp\!\!\!\perp X_B | X_C$ whenever C separates A and B in G .

(P) $X_i \perp\!\!\!\perp X_j | X_{V \setminus \{i,j\}}$ for all $ij \notin E$.⁷

By the Hammersley-Clifford theorem, graphical models could be equivalently defined by specific conditional independences.⁸

In what follows, we discuss two most important versions of undirected graphical models. One for Gaussian distributions and the other for finite discrete distributions.

⁷ (P) is called the Pairwise Markov Property.

⁸ For simplicity, in this lecture, we work only with positive distributions for which the Hammersley-Clifford theorem applies.

8.4 Gaussian Graphical Models

Gaussian Graphical Models (GGMs) describe the conditional independence structure of multivariate normal distributions. They are widely used in statistics, machine learning, and computational biology due to their interpretability and mathematical tractability.

8.4.1 Definition and Properties

We already learned about the multivariate normal distributions, and many of its properties will become essential here in this chapter. In particular, recall Proposition 2.3.2, which states that $K_{ij} = 0$ if and only if $X_i \perp\!\!\!\perp X_j | X_{V \setminus \{i,j\}}$. By the Hammersley-Clifford theorem⁹, this is equivalent to density factorization with respect to the graph that represents the support of the inverse covariance matrix $K = \Sigma^{-1}$, namely the graph G given by

$$(i, j) \notin E \iff K_{ij} = 0.$$

Moreover, by the Global Markov Property, we get a way of reading from the graph other more complicated conditional independences. Slightly abusing notation, we write

$$M(G) = \{\Sigma \in \mathbb{S}^m : \Sigma_{ij}^{-1} = 0 \text{ if } ij \notin E\}.$$

In other words, we associate the Gaussian graphical model over G with the family of covariance matrices for distributions in $M(G)$. The graphical model makes no restrictions on the mean vector.

8.4.2 Estimation with fixed G

Given n i.i.d. samples $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^m$ consider the sample covariance matrix S_n as defined in (1.4). Throughout we assume that the sample comes from $N(\mu, \Sigma)$, $\Sigma \in M(G)$. A natural way to estimate parameters in a Gaussian graphical model $M(G)$ is by maximizing the log-likelihood function. This results in a consistent and asymptotically normal estimator by standard results.

⁹ Note that the Gaussian density is always positive so this theorem applies.

Note that the mean μ is unrestricted in this model and so, by the same argument as in Section 2.4.1, the maximum likelihood estimator of μ is equal to the sample average; $\hat{\mu} = \bar{x}_n$. Thus, expressed in terms of $K = \Sigma^{-1}$, the log-likelihood function takes a simpler form (as in (2.6)):

$$\ell(\bar{x}, K) = \text{const} + \frac{n}{2}(\log \det(K) - \text{tr}(S_n K)).$$

Denote $f(K) = -\log \det(K) + \text{tr}(S_n K)$ so that maximizing $\ell(\bar{x}, K)$ is equivalent to minimizing $f(K)$. We noted already in Section 2.4.1 that f is a strictly convex function of K . Define the linear subspace

$$\mathcal{V}(G) = \{K \in \mathbb{S}^m : K_{ij} = 0 \text{ if } ij \notin E\}.$$

With this notation optimizing the log-likelihood for a Gaussian graphical model results in being equivalent to minimizing $f(K)$ a convex optimization problem; $\min f(K)$ subject to $K \in \mathcal{V}(G) \cap \mathbb{S}_+^m$.¹⁰ This problem does not have a closed form solution in general¹¹ so we solve it using numerical methods.

Some of the most efficient methods use the dual formulation of the maximum likelihood problem.

Proposition 8.4.1. *The maximum likelihood estimator $\hat{\Sigma}$, if it exists, is the unique maximizer of $\log \det \Sigma$ subject to $\Sigma_{ii} = (S_n)_{ii}$ for all $i \in V$ and $\Sigma_{ij} = (S_n)_{ij}$ for all $ij \in E$.*

The idea of the following proof sketch applies to a wide range of related problems in multivariate statistics. Thus, it is beneficial to follow it closely.

Proof Sketch. We can extend f to whole \mathbb{S}^m by setting $f(K) = +\infty$ if $K \notin \mathbb{S}_+^m$. Minimizing this extended f over \mathbb{S}^m is equivalent to minimizing f over \mathbb{S}_+^m . Consider the orthogonal complement¹²

$$\mathcal{W}(G) := \mathcal{V}(G)^\perp = \{K \in \mathbb{S}^m : K_{ii} = 0 \text{ for all } i \text{ and } K_{ij} = 0 \text{ for } ij \in E\}.$$

Observe that

$$\sup_{\Lambda \in \mathcal{W}(G)} f(K) + \text{tr}(\Lambda K) = \begin{cases} f(K) & \text{if } K \in \mathcal{V}(G), \\ +\infty & \text{otherwise.} \end{cases}$$

Indeed, if $K \in \mathcal{V}(G)$ then $\text{tr}(\Lambda K) = 0$ and so $f(K) + \text{tr}(\Lambda K) = f(K)$. On the other hand, if $K \notin \mathcal{V}(G)$, there exists $\Lambda \in \mathcal{W}(G)$ such that $\text{tr}(\Lambda K) > 0$. Taking $t\Lambda$ with $t \rightarrow \infty$ shows that $\sup_{\Lambda \in \mathcal{W}(G)} f(K) + \text{tr}(\Lambda K) = +\infty$. This all implies that

$$\inf_{K \in \mathbb{S}^m} f(K) = \inf_{K \in \mathbb{S}^m} \sup_{\Lambda \in \mathcal{W}(G)} f(K) + \text{tr}(\Lambda K).$$

By the minimax theorem from convex analysis¹³, which we are not

← Exercise 8.6.2

¹⁰ Intersection of two convex sets is convex. Moreover, \mathbb{S}_+^m is convex directly by definition.

¹¹ It does if and only if G is **decomposable**. See Lauritzen 1996 for details.

¹² If a finite dimensional linear subspace is defined by vanishing of some of the coordinates then its complement is defined by vanishing of the remaining coordinates.

¹³ This is the only informal step of this proof sketch.

going to prove here, we have

$$\inf_{K \in \mathbb{S}^m} \sup_{\Lambda \in \mathcal{W}(G)} f(K) + \text{tr}(\Lambda K) = \sup_{\Lambda \in \mathcal{W}(G)} \inf_{K \in \mathbb{S}^m} f(K) + \text{tr}(\Lambda K).$$

The latter expression is easy to handle since

$$\inf_{K \in \mathbb{S}^m} f(K) + \text{tr}(\Lambda K) = \inf_{K \in \mathbb{S}^m} -\log \det(K) + \text{tr}((S_n + \Lambda)K) = -\log \det((S_n + \Lambda)^{-1}) + m,$$

where the last equality comes from the fact that minimizing the middle expression is equivalent to maximizing the (unrestricted) Gaussian log-likelihood with sample covariance matrix $S_n + \Lambda$. We get that

$$\inf_{K \in \mathbb{S}^m} f(K) = \sup_{\Lambda \in \mathcal{W}(G)} \log \det(S_n + \Lambda) + m,$$

which gives the desired conclusion. \square

8.4.3 Matrix completion

Proposition 8.4.1 implies that the MLE $\hat{\Sigma}$ has the same diagonal as the sample covariance matrix S_n . The same applies to all entries corresponding to the edges of G . Thus, it remains to fill in the remaining entries so that the determinant is maximally possible.

Denote $\Sigma_{\setminus i, \setminus i} \in \mathbb{S}_+^{m-1}$ the submatrix of Σ obtained by removing the i -th row/columns. Let $\Sigma_{\setminus i, i} \in \mathbb{R}^{m-1}$ be the i -th column of Σ with the i -th entry removed. Note that by fixing the i -th row/column, we can write

$$\det(\Sigma) = \det(\Sigma_{\setminus i, \setminus i})(\Sigma_{ii} - \Sigma_{\setminus i, i}\Sigma_{\setminus i, \setminus i}^{-1}\Sigma_{\setminus i, i}).$$

Suppose we want to optimize the determinant with respect to $y := \Sigma_{\setminus i, i}$ only. Since $\det(\Sigma_{\setminus i, \setminus i}) > 0$, equivalently, we minimize $y^\top \Sigma_{\setminus i, \setminus i}^{-1} y$. Since $\Sigma_{\setminus i, \setminus i}$ is positive definite¹⁴, the global optimum of this quadratic form is $y = \mathbf{0}$. But in our case $y = \Sigma_{\setminus i, i}$ must satisfy $\Sigma_{ji} = (S_n)_{ji}$ for $j - i \in E$. Nevertheless, this is a simple quadratic problem with linear constraints for which there are very efficient algorithms and implementations.

We can now imagine an iterative procedure, where we update Σ row by row. This is an example of block coordinate descent. Although we are not going to prove this formally, since the original function $\log \det \Sigma$ is strictly concave, this is guaranteed to converge to the global optimum.

8.4.4 Challenges in High-Dimensional Settings

Let us reiterate what we already mentioned earlier in these notes. Estimating the covariance is hard! Estimating the inverse covariance matrix is even harder. Consider a trivial scenario of a well-behaved

In the example of Figure 8.3, if S is the sample covariance matrix then the MLE is of the form

$$\hat{\Sigma} = \begin{bmatrix} S_{11} & S_{12} & S_{13} & S_{14} & * \\ S_{12} & S_{22} & S_{23} & * & * \\ S_{13} & S_{23} & S_{33} & * & S_{35} \\ S_{14} & * & * & S_{44} & S_{45} \\ * & * & S_{35} & S_{45} & S_{55} \end{bmatrix}.$$

¹⁴ Every principal submatrix of a positive definite matrix is positive definite.

5×5 covariance matrix in the model $M(G)$ over the graph in Figure 8.3 as in the code below.

```
library(MASS)
Sigma <- solve(matrix(c(1, 0.3, 0.4, 0.5, 0,
                      0.3, 1, 0.2, 0, 0,
                      0.4, 0.2, 1, 0, 0.4,
                      0.5, 0, 0, 1, 0.5,
                      0, 0, 0.4, 0.5, 1), nrow = 5))
set.seed(123)
X10 <- mvtnorm(10,c(0,0,0,0,0),Sigma);
S10 <- cov(X10)
X100 <- mvtnorm(100,c(0,0,0,0,0),Sigma)
S100 <- cov(X100)
X1000 <- mvtnorm(1000,c(0,0,0,0,0),Sigma)
S1000 <- cov(X1000)
round(solve(Sigma),2)
round(solve(S10),2)
round(solve(S100),2)
round(solve(S1000),2)
```

The true inverse covariance matrix is

$$K = \begin{bmatrix} 1 & 0.3 & 0.4 & 0.5 & 0 \\ 0.3 & 1 & 0.2 & 0 & 0 \\ 0.4 & 0.2 & 1 & 0 & 0.4 \\ 0.5 & 0 & 0 & 1 & 0.5 \\ 0 & 0 & 0.4 & 0.5 & 1 \end{bmatrix}$$

and

$$\Sigma = \begin{bmatrix} 3.62 & -0.63 & -2.29 & -3.02 & 2.43 \\ -0.63 & 1.16 & 0.13 & 0.45 & -0.28 \\ -2.29 & 0.13 & 2.79 & 2.27 & -2.25 \\ -3.02 & 0.45 & 2.27 & 3.95 & -2.89 \\ 2.43 & -0.28 & -2.25 & -2.89 & 3.34 \end{bmatrix}$$

The following output shows that only for very large sample sizes the inverse of the sample covariance matrix becomes a decent estimation of the inverse covariance matrix.

```
> round(solve(Sigma),2)
 [,1] [,2] [,3] [,4] [,5]
[1,] 1.0 0.3 0.4 0.5 0.0
[2,] 0.3 1.0 0.2 0.0 0.0
[3,] 0.4 0.2 1.0 0.0 0.4
[4,] 0.5 0.0 0.0 1.0 0.5
[5,] 0.0 0.0 0.4 0.5 1.0
> round(solve(S10),2)
 [,1] [,2] [,3] [,4] [,5]
[1,] 1.29 0.12 0.30 0.28 -0.52
[2,] 0.12 6.28 2.21 -3.70 -3.63
[3,] 0.30 2.21 1.62 -1.44 -0.63
[4,] 0.28 -3.70 -1.44 3.63 3.28
[5,] -0.52 -3.63 -0.63 3.28 5.06
> round(solve(S100),2)
 [,1] [,2] [,3] [,4] [,5]
[1,] 1.07 0.31 0.57 0.53 0.08
[2,] 0.31 1.01 0.19 0.02 0.10
[3,] 0.57 0.19 1.28 0.10 0.53
[4,] 0.53 0.02 0.10 1.14 0.59
[5,] 0.08 0.10 0.53 0.59 1.06
> round(solve(S1000),2)
 [,1] [,2] [,3] [,4] [,5]
[1,] 1.05 0.30 0.44 0.54 0.04
[2,] 0.30 1.00 0.16 0.02 -0.05
[3,] 0.44 0.16 1.02 0.02 0.42
[4,] 0.54 0.02 0.02 1.03 0.53
[5,] 0.04 -0.05 0.42 0.53 1.02
```

For this reason, whenever the problem has some structure, it is important to exploit it. In this chapter we will show how to use the fact that the true inverse covariance matrix has some zeros. We will go back to this example in the next section to illustrate how graphical

model constraints can help.

The problem with estimating the inverse covariance matrix becomes particularly profound in higher-dimensional scenarios. We noted already in Section 4.7 that if the underlying dimension is not negligible with respect to n , the classical asymptotic results may not be very accurate. For this reason, in modern applications, we use estimators that rely on some form of regularization. We will discuss in Section 8.4.7 a form of regularization that allows us to learn the underlying graph G .

8.4.5 Appendix: Likelihood Ratio Test

Hypothesis Testing in multivariate statistics offer comes with many issues related with high-dimensionality of the problems. Likelihood Ratio Test provides a rather general approach that can be used in a variety of situations and which, under some regularity conditions comes with asymptotic statistical guarantees.

I assume that you are familiar with the likelihood ratio procedure from univariate statistics. Consider a parametric model with parameter space Θ . Given $\Theta_0 \subseteq \Theta$ we consider the null hypothesis $H_0 : \theta \in \Theta_0$. Given the log-likelihood function $\ell_n(\theta) = \sum_i \log f(\mathbf{x}_i)$. Define $\ell_1^* = \sup_{\theta \in \Theta} \ell_n(\theta)$ and $\ell_0^* = \sup_{\theta \in \Theta_0} \ell_n(\theta)$. The likelihood ratio statistic is then defined as

$$\Lambda_n = 2(\ell_1^* - \ell_0^*).$$

Note that Λ_n is always nonnegative. In general, one tends to favor H_1 when the LR statistic is low, and H_0 when it is high. More formally, the Likelihood Ratio Test of size α for testing H_0 against H_1 has as its rejection region

$$R = \{\mathbf{X} : \Lambda_n < c\},$$

where c is determined so that

$$\sup_{\theta \in \Theta_0} \mathbb{P}_{\theta}(\mathbf{X} \in R) = \alpha.$$

For the hypotheses we are interested in, the distribution of Λ_n does not depend on the particular value of $\theta \in \Theta_0$, so the above supremum is unnecessary. The LRT has the following very important property.

Theorem 8.4.2. *If $\Theta_0 \subseteq \Theta$ with $\dim(\Theta) = q$ and $\dim(\Theta_0) = r$ then, under suitable regularity conditions, for each $\theta \in \Theta_0$, Λ_n has an asymptotic χ^2_{q-r} distribution as $n \rightarrow \infty$.*

Example 8.4.3. *Consider the Gaussian distribution $N_m(\mu, \Sigma)$ and $H_0 : \mu = \mu_0$ for some fixed vector μ_0 . As this constrained defined m additional*

restrictions, we get that, under the null, Λ_n will have the asymptotic χ_m^2 distribution.

8.4.6 Structure learning (aka model selection)

In practice, the graph G is rarely known. In this case, the focus is on learning/selecting a suitable graph based on the observed data. In this section we briefly discuss some classic approaches to model selection.

Likelihood Ratio Test. A classical approach to determine the structure of G is to test a candidate graph G_0 (the null model $M(G_0)$) against its supergraph G (the alternative model $M(G)$) using the likelihood ratio test.

Suppose $X \sim N_m(\mu, \Sigma)$. Let $\mathbf{x}_1, \dots, \mathbf{x}_n$ be a random sample from this distribution. Let $G_0 = (V, E_0)$ be a graph of interest and $G = (V, E)$ be any graph containing it ($E_0 \subseteq E$). Based on the observed data we would like to test

$$H_0 : \Sigma \in M(G_0) \quad \text{against} \quad H_1 : \Sigma \in M(G).$$

Let $\ell(\bar{\mathbf{x}}_n, K)$ be as defined in (2.6). Let \hat{K} be the MLE of the concentration matrix under the model $M(G)$ and let \hat{K}_0 be the the MLE under the submodel $M(G_0)$. The likelihood ratio statistic is defined as:

$$\Lambda_n := 2(\ell(\bar{\mathbf{x}}_n, \hat{K}) - \ell(\bar{\mathbf{x}}_n, \hat{K}_0)).$$

Proposition 8.4.4. If $G_0 = (V, E_0)$ is a subgraph of $G = (V, E)$, and $v = |E \setminus E_0|$, then under the null hypothesis that $M(G_0)$ is the true model:

$$\Lambda_n \xrightarrow{d} \chi_v^2.$$

This result allows us to test whether the additional edges in G significantly improve the model fit, by comparing Λ_n to the critical value of the χ_v^2 distribution:

← Exercise 8.6.3

1. Compute Λ_n based on the data.
2. Compute p-value as $P = \mathbb{P}(\chi_v^2 \geq \Lambda_n)$.
3. If P is small, say less then 0.05, reject the null hypothesis.

We illustrate the process using a small simulated dataset. Suppose we observe data from a multivariate normal distribution, and we wish to select a Gaussian graphical model that explains the dependencies among the variables. In this code, the true covariance matrix generating the data lies in the model defined by the graph in Figure 8.3. We will test this model against the unrestricted Gaussian model defined by the complete graph.

← Exercise 8.6.5

```

library(ggm)

# Generate synthetic data for illustration
set.seed(123)
n <- 100 # Sample size
p <- 5 # Number of variables

# Simulate data from a multivariate normal distribution
Sigma <- solve(matrix(c(1, 0.3, 0.4, 0.5, 0,
                        0.3, 1, 0.2, 0, 0,
                        0.4, 0.2, 1, 0, 0.4,
                        0.5, 0, 0, 1, 0.5,
                        0, 0, 0.4, 0.5, 1), nrow = 5))
data <- MASS::mvrnorm(n = n, mu = rep(0, p), Sigma = Sigma)

# Define two competing graphs (G0: subgraph, G: full graph)
G0 <- matrix(c(0, 1, 1, 1, 0,
               1, 0, 1, 0, 0,
               1, 1, 0, 0, 1,
               1, 0, 0, 0, 1,
               0, 0, 1, 1, 0), nrow = 5, byrow = TRUE)
G <- matrix(1, nrow = 5, ncol = 5) # Full graph (fully connected)
diag(G) <- 0

# Assign variable names
var_names <- c("X1", "X2", "X3", "X4", "X5")
colnames(data) <- colnames(G0) <- colnames(G) <- rownames(G) <- rownames(G0)
↪ <- var_names
S <- cov(data)

# Fit models under G0 and G
Shat0 <- fitConGraph(G0, S, nrow(data))$Shat
Shat <- fitConGraph(G, S, nrow(data))$Shat
Khat0 <- solve(Shat0)
Khat <- solve(Shat)

# Likelihood Ratio Test
logLik_G0 <- (n/2)*(log(det(Khat0)) - sum(Khat0*S))
logLik_G <- (n/2)*(log(det(Khat)) - sum(Khat*S))
LRT_stat <- 2 * (logLik_G - logLik_G0)
df <- (sum(G) - sum(G0))/2 # Difference in number of edges
p_value <- pchisq(LRT_stat, df, lower.tail = FALSE)

cat("Likelihood Ratio Test Statistic:", LRT_stat, "\n")
cat("Degrees of Freedom:", df, "\n")
cat("p-value:", p_value, "\n")

```

The corresponding p-value¹⁵ of the likelihood ratio test is 0.294, this is not super conclusive but we have no enough evidence to reject the null hypothesis.

¹⁵ Recall that the p-value is the probability (under the null) that Λ_n is greater than what was observed in the data. Very small values provide evidence against H_0 .

Remark 8.4.5. We note in passing that the estimator of the inverse covari-

ance matrix under the true graph in Figure 8.3 is

$$\begin{bmatrix} 1.04 & 0.34 & 0.40 & 0.48 & 0 \\ 0.34 & 1.02 & 0.17 & 0 & 0 \\ 0.40 & 0.17 & 0.99 & 0 & 0.47 \\ 0.48 & 0 & 0 & 0.99 & 0.51 \\ 0 & 0 & 0.47 & 0.51 & 1.19 \end{bmatrix}.$$

This is a slightly more accurate estimator of the real inverse covariance matrix than the one given by the inverse of the sample covariance matrix in Section 8.4.4.

Akaike's Information Criterion (AIC) and Bayesian Information Criterion (BIC). An alternative approach to model selection is based on information criteria, which balance the goodness of fit and model complexity. The dimension of the model $M(G)$ is simply the number of nodes m plus the number of edges E ; $\dim M(G) = m + |E|$. AIC is defined as:

$$\text{AIC} = -2\ell(\bar{\mathbf{x}}_n, \hat{K}) + 2 \cdot \dim M(G)$$

and BIC is defined as:

$$\text{BIC} = -2\ell(\bar{\mathbf{x}}_n, \hat{K}) + \log(n) \cdot \dim M(G).$$

The model with the smallest AIC/BIC value is selected. Unlike hypothesis testing, information criteria provide a way to compare non-nested models directly.

The number of possible undirected graphs on m nodes is $2^{\binom{m}{2}}$ and so it grows exponentially, making exhaustive search computationally infeasible for even moderate m . A stepwise procedure offers a practical alternative:

1. Start with the complete graph or a simpler initial model.
2. Iteratively add or remove edges:
 - Evaluate all models obtained by adding or removing one edge.
 - Select the model with the smallest AIC/BIC.
3. Stop when no improvement in AIC/BIC is observed.

← Exercise 8.6.4

This greedy approach does not guarantee the globally optimal model, but it is computationally efficient.

The package `gRim` offers an implementation of such a stepwise procedure in R. Consider the following code.

```
library(gRim)
library(qgraph)

sat.model <- cmod(~.^., data=data);
```

Note that the “aic” criterion in the code refers in fact to the BIC as $k = \log(n)$.

```

test.model <-
  ↵ stepwise(sat.model,details=1,criterion="aic",type="unrestricted",k=log(
    n))
# plot(test.carc) or using the qgraph package
Sigma.hat <- solve(test.model$fitinfo$K)
qgraph::qgraph(Sigma.hat,graph="pcor")

```

We used here the package `qgraph` only to plot the underlying partial correlation graph. We will mention this package again when discussing high-dimensional scenarios.

8.4.7 Graphical LASSO

In high-dimensional scenarios, stepwise procedures described above are impractical and we often use other approaches based on regularization techniques.

The graphical LASSO estimates sparse precision matrices by solving the optimization problem:

$$\hat{K} = \arg \max_{K \in S_+^m} \{ \log \det(K) - \text{tr}(S_n K) - \lambda \|K\|_{1,\text{off}} \}, \quad (8.2)$$

where $\lambda > 0$ is a fix penalty parameter and

$$\|K\|_{1,\text{off}} := \sum_{i \neq j} |K_{ij}|.$$

In other words the graphical LASSO estimator is obtained by penalizing the log-likelihood function with a term that penalizes the off-diagonal entries to be too large effectively pushing some of them to be zero. This is in the same spirit as the LASSO approach in linear regression, hence the name.

An efficient algorithm to find \hat{K} follows exactly the same lines as the block coordinate descent algorithm for maximizing the log-likelihood function. First we formulate a variational form of $\|K\|_{1,\text{off}}$. Noting that $|x| = \sup_{-1 \leq y \leq 1} xy$ we get

$$\lambda \|K\|_{1,\text{off}} = \lambda \sum_{i \neq j} \sup_{-1 \leq \Lambda_{ij} \leq 1} \Lambda_{ij} K_{ij} = \sup_{\Lambda} \text{tr}(\Lambda K),$$

where the supremum is taken over all $\Lambda \in S^m$ with zeros on the diagonal and off-diagonal entries in $[-\lambda, \lambda]$.

This allows to rewrite the main function in (8.2) as

$$\sup_{\Lambda} \left\{ \log \det(K) - \text{tr}((S_n + \Lambda)K) \right\}.$$

Proceeding as in the proof of Proposition 8.4.1, we get that the dual problem is to maximize $\log \det \Sigma$ subject to $\Sigma_{ii} = (S_n)_{ii}$ for $i = 1, \dots, m$ and $|\Sigma_{ij} - (S_n)_{ij}| \leq \lambda$.

We can again solve this problem using the same block coordinate descent approach.




Figure 8.3: The selected model is close to the true graph. Only the edge 2 – 3 is missing.

1. Set $\Sigma = S_n$.
2. For each $i = 1, \dots, m$, set $\Sigma_{\setminus i, i}$ and $\Sigma_{i, \setminus i}$ to be equal to \hat{y} , which is the minimizer of $y^\top (\Sigma_{\setminus i, \setminus i})^{-1} y$ subject to $y_j = (S_n)_{ij}$ for all j such that $ij \in E$.
3. Check a convergence criterion (pick one you like in advance). If it is not satisfied, go back to Step 2. Otherwise stop.

Here is an example of applying the graphical LASSO in R using the `glasso` package:

```
library(glasso)

# Simulated data
set.seed(123)
n <- 100
p <- 10
X <- matrix(rnorm(n * p), nrow = n, ncol = p)

# Sample covariance matrix
S <- cov(X)

# Apply graphical LASSO
lambda <- 0.1
fit <- glasso(S, rho = lambda)

# Extract precision matrix
precision_matrix <- fit$wi
print(precision_matrix)
```

Remark 8.4.6. Many packages, like `glasso`, penalize in (8.2) not only the off-diagonal entries of K but also the diagonal entries. This however introduces an unnecessary bias.

8.5 Log-linear Graphical Models

In this lecture course, we have largely ignored discrete variables. We try to partly fix this in this section by discussing discrete graphical models.

8.5.1 Introduction

Log-linear graphical models describe the dependence structure of discrete random variables using a graph-based factorization of the joint probability distribution. Discrete data arise naturally in many real-world settings, making their analysis essential in applied statistics and data science. Here are a few explicit examples of scenarios where finite discrete data are relevant:

- **Genetics:** In genetics, categorical variables often represent genotypes (e.g., AA, Aa, aa) or the presence or absence of certain genetic mutations. For example, a study investigating the relationship between multiple genetic variants and a disease status might use a log-linear model to analyze these interactions.
- **Market Basket Analysis:** In retail analytics, log-linear models are applied to understand the co-occurrence patterns of items in shopping baskets. Each variable represents whether a customer bought a specific item (e.g., milk, bread, eggs), and the observed data are counts of item combinations.
- **Survey Data:** Social science surveys frequently collect categorical responses to questions, such as political preferences (e.g., liberal, conservative) or education levels (e.g., high school, college, graduate school). Log-linear models can analyze the relationships between these variables, adjusting for potential confounding factors.

8.5.2 Notation

Let $X = (X_1, \dots, X_m)$ be a vector of discrete random variables, where each X_i takes values in a finite set \mathcal{X}_i . Without loss of generality, label the states as $\mathcal{X}_i = \{0, 1, \dots, r_i - 1\}$ for some $r_i \geq 2$. The joint distribution of X is represented by a probability mass function $f(\mathbf{x}) = P(X = \mathbf{x})$ for $\mathbf{x} = (x_1, \dots, x_m) \in \mathcal{X}_1 \times \dots \times \mathcal{X}_m =: \mathcal{X}$.

Suppose now we observe a random sample $\mathbf{x}_1, \dots, \mathbf{x}_n$ from this distribution. The likelihood is

$$\prod_{i=1}^n f(\mathbf{x}_i) = \prod_{\mathbf{x} \in \mathcal{X}} f(\mathbf{x})^{n(\mathbf{x})},$$

where $n(\mathbf{x})$ denotes the number of times the state \mathbf{x} was observed in the sample. The collection of these counts $\mathbf{n} = (n(\mathbf{x}))$ is called the contingency table. The contingency table is an $r_1 \times \dots \times r_m$ array, where the (x_1, \dots, x_m) -th entry $n(x_1, \dots, x_m)$ represents the number of times the configuration (x_1, \dots, x_m) is observed in the data. The total sample size n satisfies:

$$n = \sum_{\mathbf{x} \in \mathcal{X}} n(\mathbf{x}).$$

If $m = 2$ and X_1, X_2 are binary variables taking values in $\{0, 1\}$, the contingency table might look like this:

		$X_2 = 0$	$X_2 = 1$
$X_1 = 0$	$n(0, 0)$	$n(0, 1)$	
	$n(1, 0)$	$n(1, 1)$	

where $n(0, 0)$ records the number of observations where $(X_1, X_2) = (0, 0)$, $n(0, 1)$ records the number where $(X_1, X_2) = (0, 1)$, and so on.

Example 8.5.1. To make this concrete, consider a contingency table summarizing the relationship between smoking status (*Smoker*, *Non-Smoker*) and disease status (*Disease*, *No Disease*). The observed data might look like this:

	Disease	No Disease
Smoker	50	200
Non-Smoker	30	300

This table indicates that 50 smokers were observed to have the disease, while 30 non-smokers were observed to have the disease, and so on. The total sample size is $n = 580$.

For $m > 2$, the contingency table becomes a multidimensional array. For example, if $m = 3$ and each X_i is binary, the contingency table is a $2 \times 2 \times 2$ cube, with entries $n(x_1, x_2, x_3)$ representing the frequencies of each configuration (x_1, x_2, x_3) .

Example 8.5.2. Suppose a dataset summarizes the co-occurrence of purchases of *Milk*, *Bread*, and *Eggs*. A contingency table is a $2 \times 2 \times 2$ array where the entry $(0, 0, 0)$ gives the number of clients that bought none of the three products, $(0, 1, 1)$ entry gives the number of clients that bought both *Bread* and *Eggs* etc

no Eggs:	no Milk	no Bread		Bread		Eggs:	no Milk	no Bread		Bread	
		10	15	25	30			Milk	35	10	
	Milk	20	5								

The joint probability of the observed contingency table is

$$\mathbb{P}(n(\mathbf{x}) : \mathbf{x} \in \mathcal{X}) = \frac{n!}{\prod_{\mathbf{x} \in \mathcal{X}} n(\mathbf{x})!} \prod_{\mathbf{x} \in \mathcal{X}} f(\mathbf{x})^{n(\mathbf{x})},$$

which differs from the likelihood by a multinomial coefficient which is just a constant.

If we do not restrict the probabilities in any way (except requiring that they are non-negative and sum to unity), then it is easily shown that the maximum likelihood estimates are given by $\hat{f}(\mathbf{x}) = n(\mathbf{x})/n$ for $\mathbf{x} \in \mathcal{X}$. The unrestricted model is known as the saturated model. In most substantive contexts, it is of interest to restrict the probabilities further to obtain parsimony and to identify or exploit structural information.

We need a little more notation. The expected cell counts are written $m(\mathbf{x}) = nf(\mathbf{x})$ for $\mathbf{x} \in \mathcal{X}$, and the fitted values as $\hat{m}(\mathbf{x}) = n\hat{f}(\mathbf{x})$. We need to work with marginal tables and to do this must first define marginal cells. For a subset of variables $C \subseteq V$, we define marginal counts as:

$$n(\mathbf{x}_C) = \sum_{\mathbf{y} \in \mathcal{X}: \mathbf{y}_C = \mathbf{x}_C} n(\mathbf{x}).$$

We use the analogous notation for $f(\mathbf{x}_C)$, $m(\mathbf{x}_C)$ etc.

Recall $V = \{1, \dots, m\}$.

8.5.3 Discrete Graphical Models through Factorization

Let \mathcal{C} be the set of (maximal) cliques of the graph G . For each $C \in \mathcal{C}$, consider the set $\mathcal{X}_C = \prod_{i \in C} \mathcal{X}_i$, and let $\theta_C : \mathcal{X}_C \rightarrow \mathbb{R}$.¹⁶ The joint probability mass function $f(\mathbf{x}) = P(X = \mathbf{x})$ of a log-linear model corresponding to G is expressed as:

$$f(\mathbf{x}; \theta) = \frac{1}{Z(\theta)} \exp \left(\sum_{C \in \mathcal{C}} \theta_C(\mathbf{x}_C) \right), \quad (8.3)$$

where $Z(\theta) = \sum_{\mathbf{x} \in \mathcal{X}} \exp \left(\sum_{C \in \mathcal{C}} \theta_C(\mathbf{x}_C) \right)$ is the normalizing constant and $\theta = (\theta_C)_{C \in \mathcal{C}}$ represents the collection of clique-specific parameters.

Example 8.5.3. Consider $m = 3$, where X_1, X_2, X_3 are binary variables taking values in $\{0, 1\}$. Suppose the underlying graph G is a chain $\bullet^1 - \bullet^2 - \bullet^3$. The cliques of G are $\mathcal{C} = \{\{1, 2\}, \{2, 3\}\}$. The joint probability mass function can be written as:

$$f(x_1, x_2, x_3; \theta) = \frac{1}{Z(\theta)} \exp (\theta_{12}(x_1, x_2) + \theta_{23}(x_2, x_3)),$$

where $\theta_{12}(x_1, x_2)$ captures the interaction between X_1 and X_2 , and $\theta_{23}(x_2, x_3)$ captures the interaction between X_2 and X_3 . Both θ_{12} and θ_{23} are encoded as vectors in \mathbb{R}^4 :

$$\theta_{12} = \begin{bmatrix} \theta_{12}(0, 0) \\ \theta_{12}(0, 1) \\ \theta_{12}(1, 0) \\ \theta_{12}(1, 1) \end{bmatrix}, \quad \theta_{23} = \begin{bmatrix} \theta_{23}(0, 0) \\ \theta_{23}(0, 1) \\ \theta_{23}(1, 0) \\ \theta_{23}(1, 1) \end{bmatrix}$$

so effectively we have 8 parameters¹⁷.

The normalizing constant $Z(\theta)$ ensures that the probabilities sum to 1:

$$Z(\theta) = \sum_{x_1, x_2, x_3 \in \{0, 1\}} \exp (\theta_{12}(x_1, x_2) + \theta_{23}(x_2, x_3)).$$

The formula in (8.3) is equivalent to f factorizing with respect to G , as defined in Equation (8.1)¹⁸. By the Hammersley-Clifford theorem, the corresponding distributions satisfy the global, local, and pairwise Markov properties.

This equivalence between the graph structure and conditional independence makes log-linear models highly interpretable.

8.5.4 Working with Log-Linear Models in R

To work with log-linear models, the R package **MASS** provides functions such as **loglm**, which can fit log-linear models to contingency tables.

¹⁶ Since $d_C := |\mathcal{X}_C|$ is finite, we can represent θ_C as a vector in \mathbb{R}^{d_C} .

¹⁷ We note in passing that these parameters are of no direct interest. In fact the model has only dimension 5 and so this parametrization leads to lack of identifiability.

¹⁸ Clearly, if (8.3) holds then f factorizes so it remains to show that if $f > 0$ factorizes then it can be always written as (8.3).

For example, suppose we have a contingency table for three binary variables X_1, X_2, X_3 as in Example 8.5.2. We could fit a model with the chain structure $\bullet - \bullet - \bullet^3$. Here is how to fit the model in R:

```
# Load the MASS package
library(MASS)

# Define the contingency table
contingency_table <- array(
  c(10, 20, 15, 5, 25, 35, 30, 10), # Frequencies
  dim = c(2, 2, 2),                  # Dimensions for X1, X2, X3
  dimnames = list(
    X1 = c("0", "1"),
    X2 = c("0", "1"),
    X3 = c("0", "1")
  )
)

# Fit the log-linear model with the chain structure
model <- loglm(~ X1 * X2 + X2 * X3, data = contingency_table)

# Display the results
summary(model)
```

The output looks like this.

```
Formula:
~X1 * X2 + X2 * X3
attr(,"variables")
list(X1, X2, X3)
attr(,"factors")
  X1 X2 X3 X1:X2 X2:X3
X1  1  0  0    1    0
X2  0  1  0    1    1
X3  0  0  1    0    1
attr(,"term.labels")
[1] "X1"    "X2"    "X3"    "X1:X2" "X2:X3"
attr(,"order")
[1] 1 1 1 2 2
attr(,"intercept")
[1] 1
attr(,"response")
[1] 0
attr(,".Environment")
<environment: R_GlobalEnv>

Statistics:
      X^2  df  P(> X^2)
Likelihood Ratio 0.5906840 2 0.7442770
Pearson          0.5844156 2 0.7466134
```

The p-values in this code refer to the test of this model with respect to the saturated model.

8.5.5 The Log-Likelihood Function

Suppose we have observed data points $\mathbf{x}_1, \dots, \mathbf{x}_n$, where each \mathbf{x}_i takes values in the discrete space \mathcal{X} . For a fixed graph G , the goal is to

maximize the log-likelihood function

$$\ell(\theta) = \sum_{i=1}^n \log f(\mathbf{x}_i; \theta) = \sum_{i=1}^n \sum_{C \in \mathcal{C}} \theta_C((\mathbf{x}_i)_C) - n \log Z(\theta),$$

where $f(\mathbf{x}; \theta)$ is the joint probability mass function parameterized by θ , and $Z(\theta)$ is the normalizing constant ensuring $f(\mathbf{x}; \theta)$ sums to 1 over all configurations of \mathbf{x} .

The notation here may seem dense, so let us break it down step by step. In the above expression, the double summation over i and C can be reordered because both ranges are finite:

$$\sum_{i=1}^n \sum_{C \in \mathcal{C}} \theta_C((\mathbf{x}_i)_C) = \sum_{C \in \mathcal{C}} \sum_{i=1}^n \theta_C((\mathbf{x}_i)_C).$$

For a fixed clique $C \in \mathcal{C}$, the term $\sum_{i=1}^n \theta_C((\mathbf{x}_i)_C)$ can be simplified by grouping terms according to $\mathbf{x}_C \in \mathcal{X}_C$

$$\sum_{i=1}^n \theta_C((\mathbf{x}_i)_C) = \sum_{\mathbf{x}_C \in \mathcal{X}_C} n_C(\mathbf{x}_C) \theta_C(\mathbf{x}_C).$$

Substituting this back into the log-likelihood function gives:

$$\ell(\theta) = \sum_{C \in \mathcal{C}} \sum_{\mathbf{x}_C \in \mathcal{X}_C} n_C(\mathbf{x}_C) \theta_C(\mathbf{x}_C) - n \log Z(\theta).$$

8.5.6 Fitting Log-linear Models: Iterative Proportional Scaling (IPS)

The parameter vectors $\theta_C \in \mathbb{R}^{|\mathcal{X}_C|}$ in log-linear models can be estimated using a simple iterative procedure. For simplicity of exposition, we present a dual algorithm, called Iterative Proportional Scaling, that tries to estimate the expectations $m(\mathbf{x})$ for $\mathbf{x} \in \mathcal{X}$.

Let \mathcal{C} be the set of maximal cliques of the graph G . The corresponding marginal tables $n_C(\mathbf{x}_C)$ are a set of sufficient statistics. The maximum likelihood estimate is obtained by equating the sufficient statistics with their expectations $m_C(\mathbf{x}_C)$.

Initially the $m(\mathbf{x})$ are set to some constant, say $m(\mathbf{x}) = 1$ for all $\mathbf{x} \in \mathcal{X}$. One iteration consists of updating for each $C \in \mathcal{C}$

$$m(\mathbf{x}) \leftarrow m(\mathbf{x}) \frac{n_C(\mathbf{x}_C)}{m_C(\mathbf{x}_C)} \quad \forall \mathbf{x} \in \mathcal{X}.$$

Iteration continues until convergence which happens when $m_C(\mathbf{x}_C) = n_C(\mathbf{x}_C)$ for all C and \mathbf{x} . The algorithm is always theoretically convergent with the limiting value being the maximum likelihood estimate under the model $\{\hat{m}(\mathbf{x})\}_{\mathbf{x} \in \mathcal{X}}$, although these may not all have $\hat{m}(\mathbf{x}) > 0$ for all cells \mathbf{x} and thus may not admit a logarithmic expansion.

8.5.7 The Ising Model

The dimensionality of log-linear models does not scale well with the number of variables and their number of values. For this reason, it is important to consider simpler models for discrete data that retain computational feasibility while capturing essential dependencies.

One prominent example is the Ising model.

The Ising model is a special case of a log-linear graphical model that is widely used in statistics, machine learning, and physics. It models binary variables with pairwise interactions.

Consider a random vector $X = (X_1, \dots, X_m)$, where each $X_i \in \{-1, 1\}$. Fix a graph $G = (V, E)$. The probability mass function is given by:

$$f(\mathbf{x}) = \mathbb{P}(X = \mathbf{x}) \propto \exp \left(\sum_{i \in V} \theta_i x_i + \sum_{(i,j) \in E} \theta_{ij} x_i x_j \right),$$

where θ_i are node potentials representing individual variable effects, and θ_{ij} are edge potentials representing pairwise interactions between variables. The normalization constant, or partition function, $Z(\theta)$, is:

$$Z(\theta) = \sum_{x \in \{-1, 1\}^m} \exp \left(\sum_{i \in V} \theta_i x_i + \sum_{(i,j) \in E} \theta_{ij} x_i x_j \right).$$

Applications. The Ising model is highly versatile and is widely applied in different domains:

- **Network Analysis:** Social networks where binary variables X_i represent the presence or absence of a specific behavior or trait for individual i , and θ_{ij} encodes the strength of the interaction between individuals i and j .
- **Spatial Statistics:** Modeling dependencies in spatial data, such as image segmentation, where X_i represents the label of pixel i (e.g., foreground or background) and θ_{ij} captures the similarity between neighboring pixels.
- **Statistical Physics:** Originally proposed to model ferromagnetic materials, where X_i represents the spin of an electron, and θ_{ij} models the coupling between spins.
- **Genetics:** Modeling binary genetic traits with pairwise dependencies between genes.

Statistical Estimation and Inference. Consider a random sample $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$. Maximum likelihood is the canonical method to esti-

mate the parameters θ_i and θ_{ij} . The log-likelihood is

$$\ell(\theta) = \sum_{k=1}^n \left(\sum_{i \in V} \theta_i x_i^{(k)} + \sum_{(i,j) \in E} \theta_{ij} x_i^{(k)} x_j^{(k)} \right) - n \log Z(\theta),$$

The log-likelihood can be maximized using iterative procedures like IPS described above.

The presence of $Z(\theta)$ makes exact inference and parameter estimation computationally challenging for large graphs. Indeed, computing $Z(\theta)$ requires performing a sum over w^m terms, which is often computationally prohibitive. This motivates other estimation methods.

Pseudo-Likelihood Methods: To avoid the computational burden of $Z(\theta)$, one can replace the likelihood with the pseudo-likelihood:

$$\ell_{\text{pseudo}}(\theta) = \sum_{k=1}^n \sum_{i \in V} \log P(X_i = x_i^{(k)} | X_{-i}),$$

where X_{-i} denotes all variables except X_i . This approximation significantly reduces computational cost while providing reasonably accurate parameter estimates.

Regularization: Sparse Ising models can be estimated using penalized pseudo-likelihood, incorporating ℓ_1 -penalties on the edge parameters θ_{ij} .

Below is a simple example of fitting an Ising model using the qgraph package in R.

```
# Install and load necessary packages
install.packages("qgraph")
library(qgraph)

# Generate synthetic binary data
set.seed(123)
n <- 100 # Number of observations
p <- 5 # Number of variables
data <- matrix(sample(c(-1, 1), n * p, replace = TRUE), nrow = n)

# Fit Ising model using pseudo-likelihood
fit <- IsingFit(data, plot = TRUE)

# Display estimated node and edge parameters
fit$thresholds # Node potentials
fit$weights # Edge potentials
```

This code generates synthetic data, fits an Ising model using pseudo-likelihood, and plots the resulting graph. The output includes estimated node and edge parameters, which can be interpreted in the context of the problem domain.

8.6 Exercises

Exercise 8.6.1. Consider all binary distributions that factorize $f(x_1, x_2, x_3) = \phi_{12}(x_1, x_2)\phi_{23}(x_2, x_3)$ for $(x_1, x_2, x_3) \in \{0, 1\}^3$. Prove the Hammerley-Clifford theorem in this very special case.

Exercise 8.6.2. Show that the set \mathbb{S}_+^m is convex. Also, show that intersection of two convex sets must be also convex.

Exercise 8.6.3. Verify by simulations that the likelihood ratio statistics has the reported asymptotic distribution χ_4^2 in the leading example of Section 8.4.6.

Exercise 8.6.4. Suppose we compare a graph G with G_0 obtained from G by removing a single edge. What difference in likelihood is needed so that the AIC/BIC criteria prefer the bigger model?

Exercise 8.6.5. Consider the code on estimating the inverse covariance matrix for the simple graphical model in Figure 8.3, where the true inverse covariance matrix is

$$K = \begin{bmatrix} 1 & 0.3 & 0.4 & 0.5 & 0 \\ 0.3 & 1 & 0.2 & 0 & 0 \\ 0.4 & 0.2 & 1 & 0 & 0.4 \\ 0.5 & 0 & 0 & 1 & 0.5 \\ 0 & 0 & 0.4 & 0.5 & 1 \end{bmatrix}.$$

Using simulations, study the performance of the MLE of K under $M(G)$ compared to the full graph as provided in the code in Section 8.4.4.

